



Technische Hochschule
Ingolstadt

Fakultät Informatik

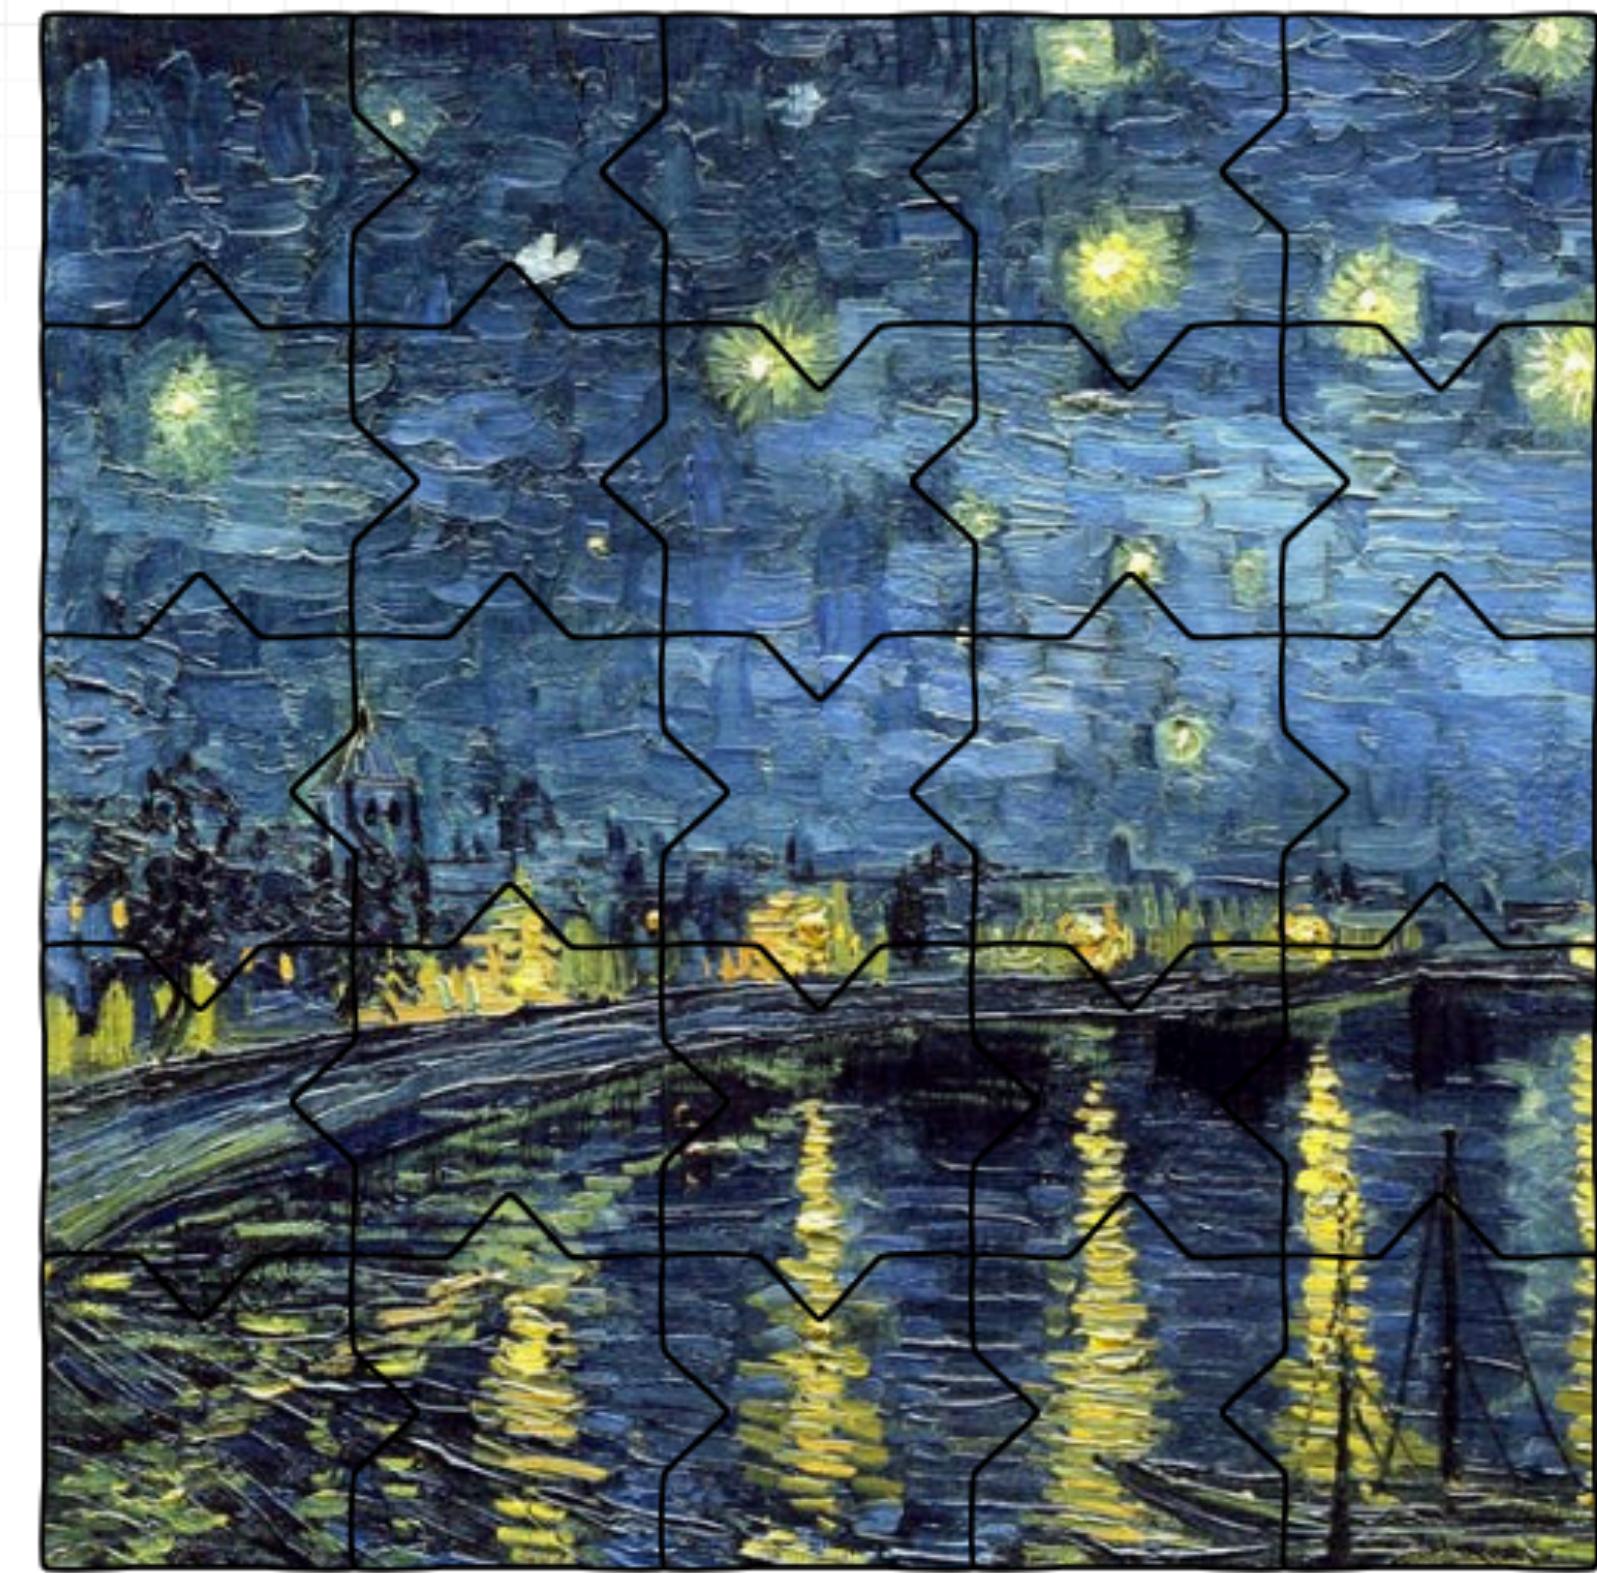
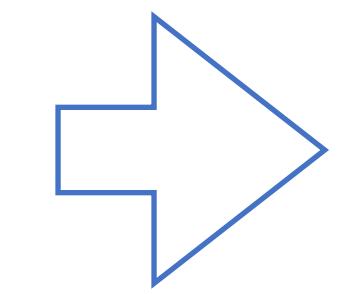
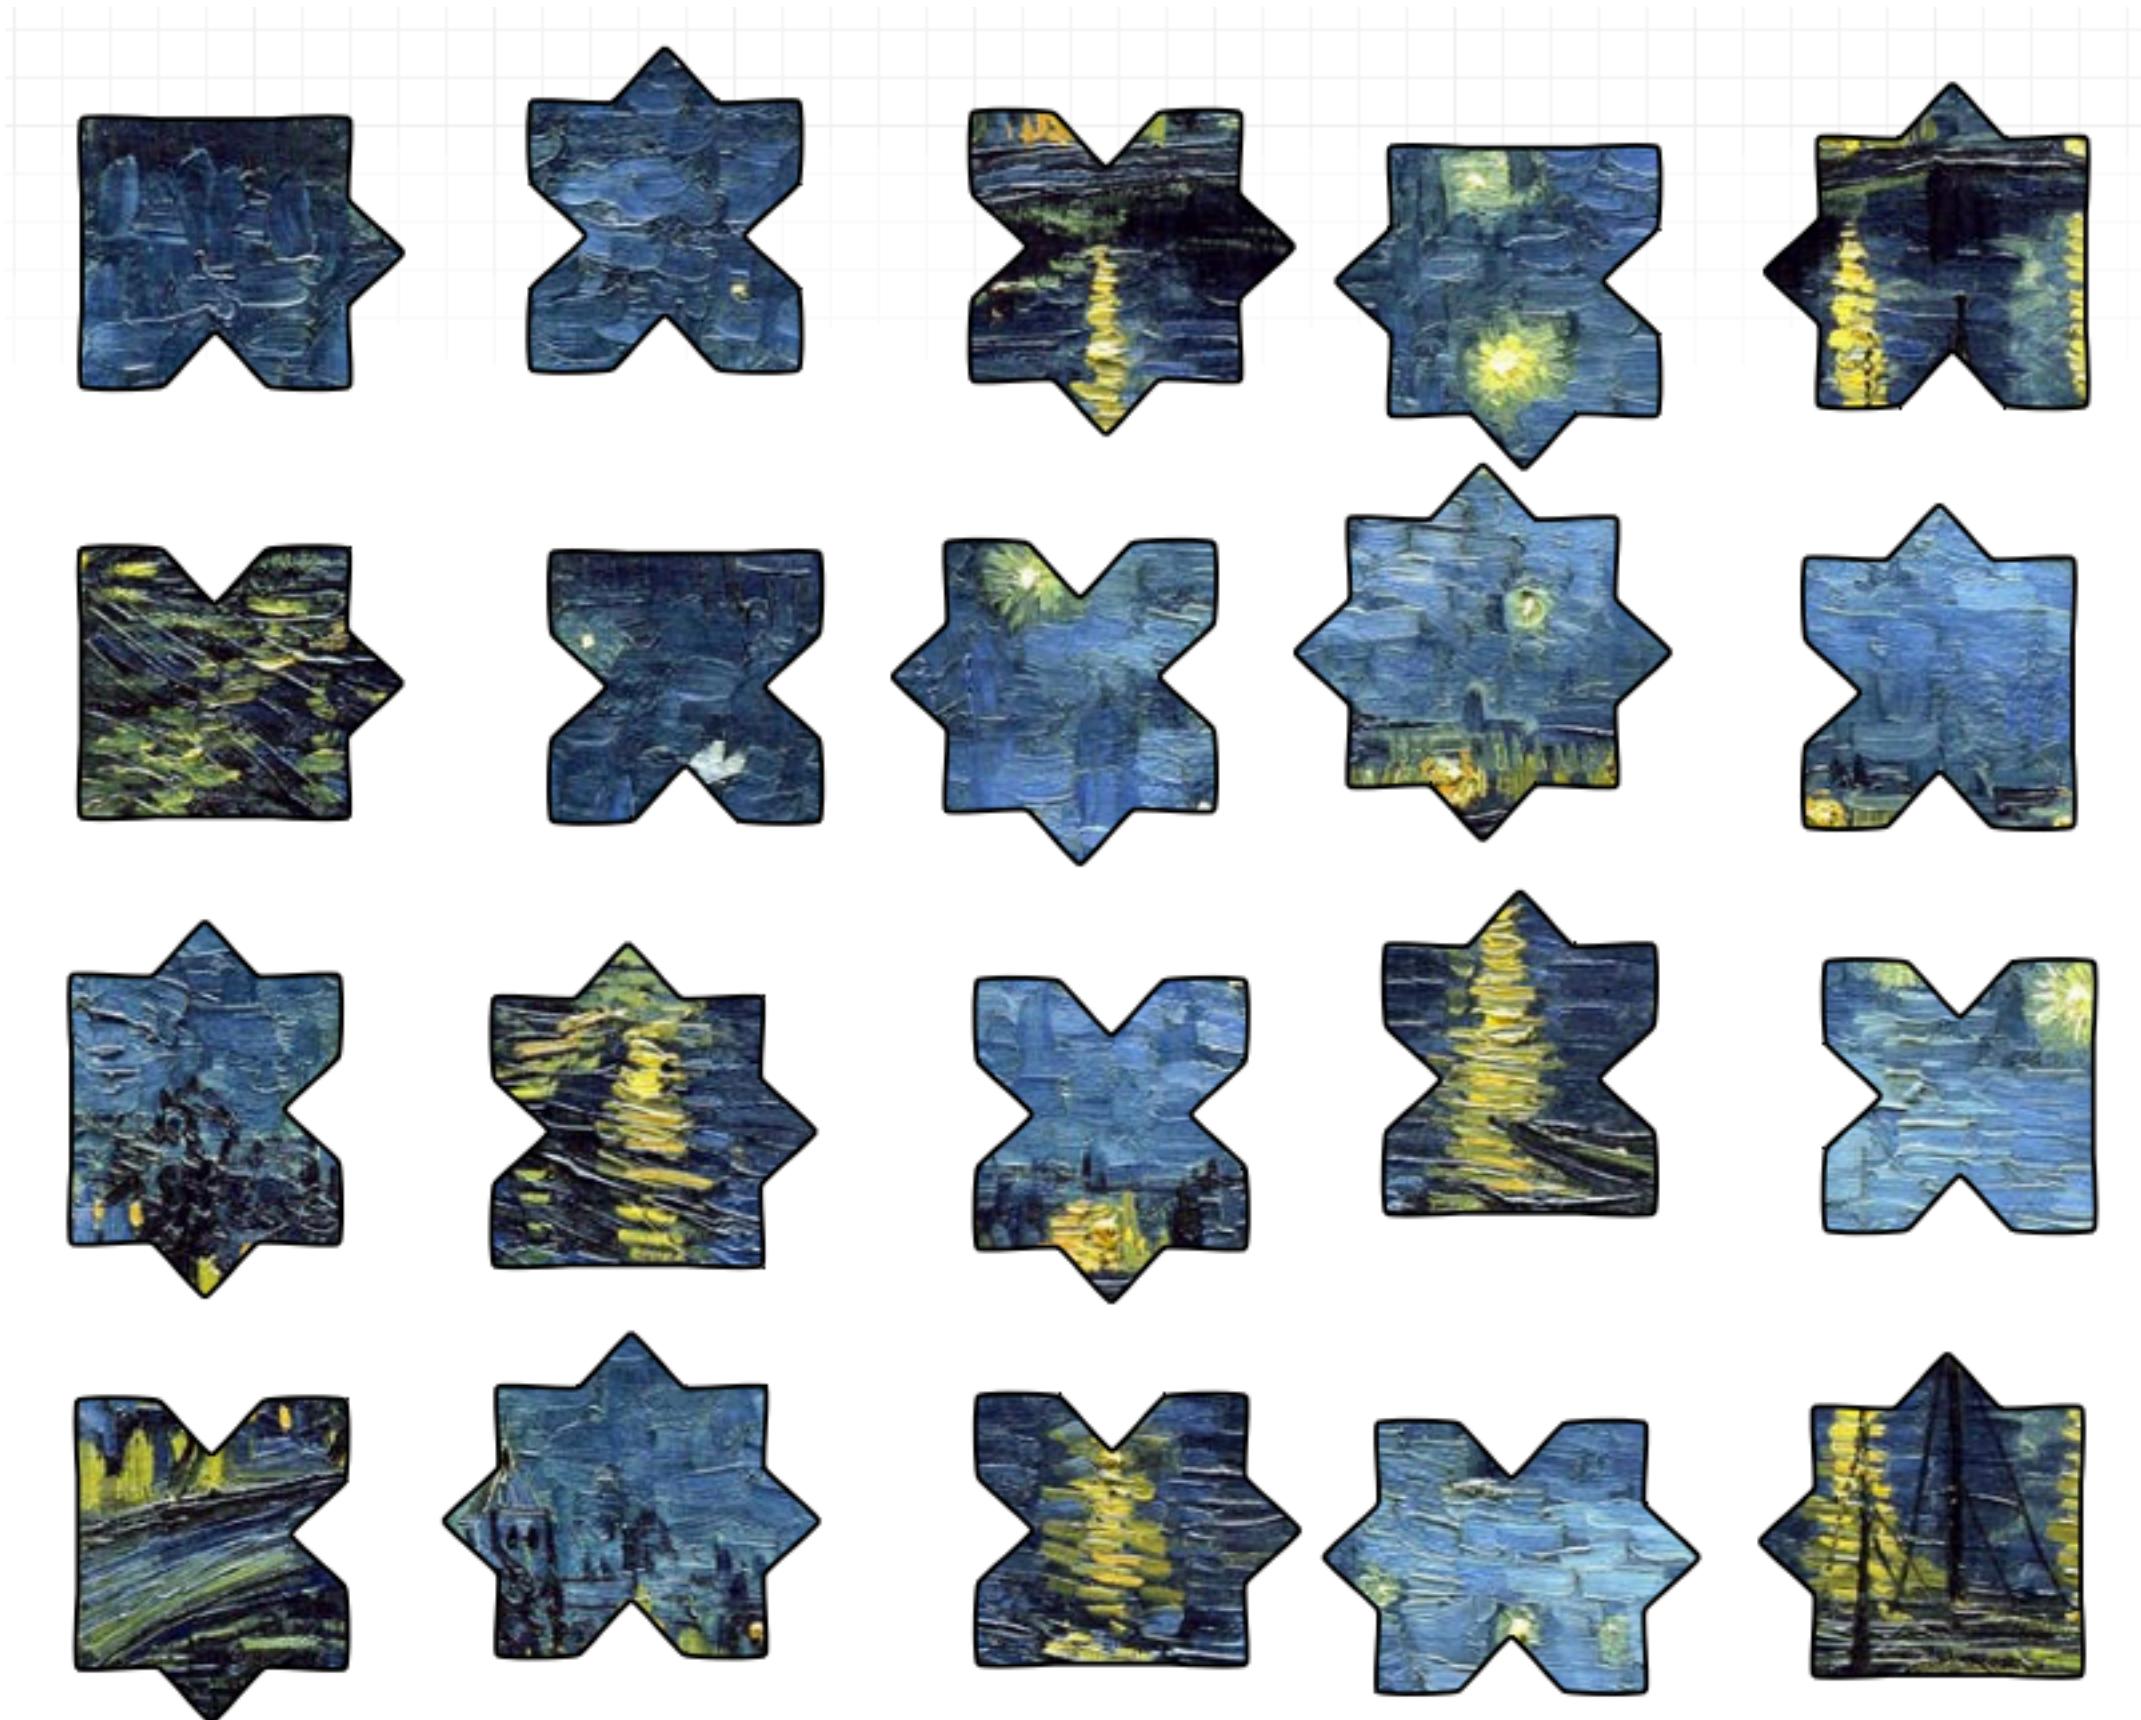
Computer Vision

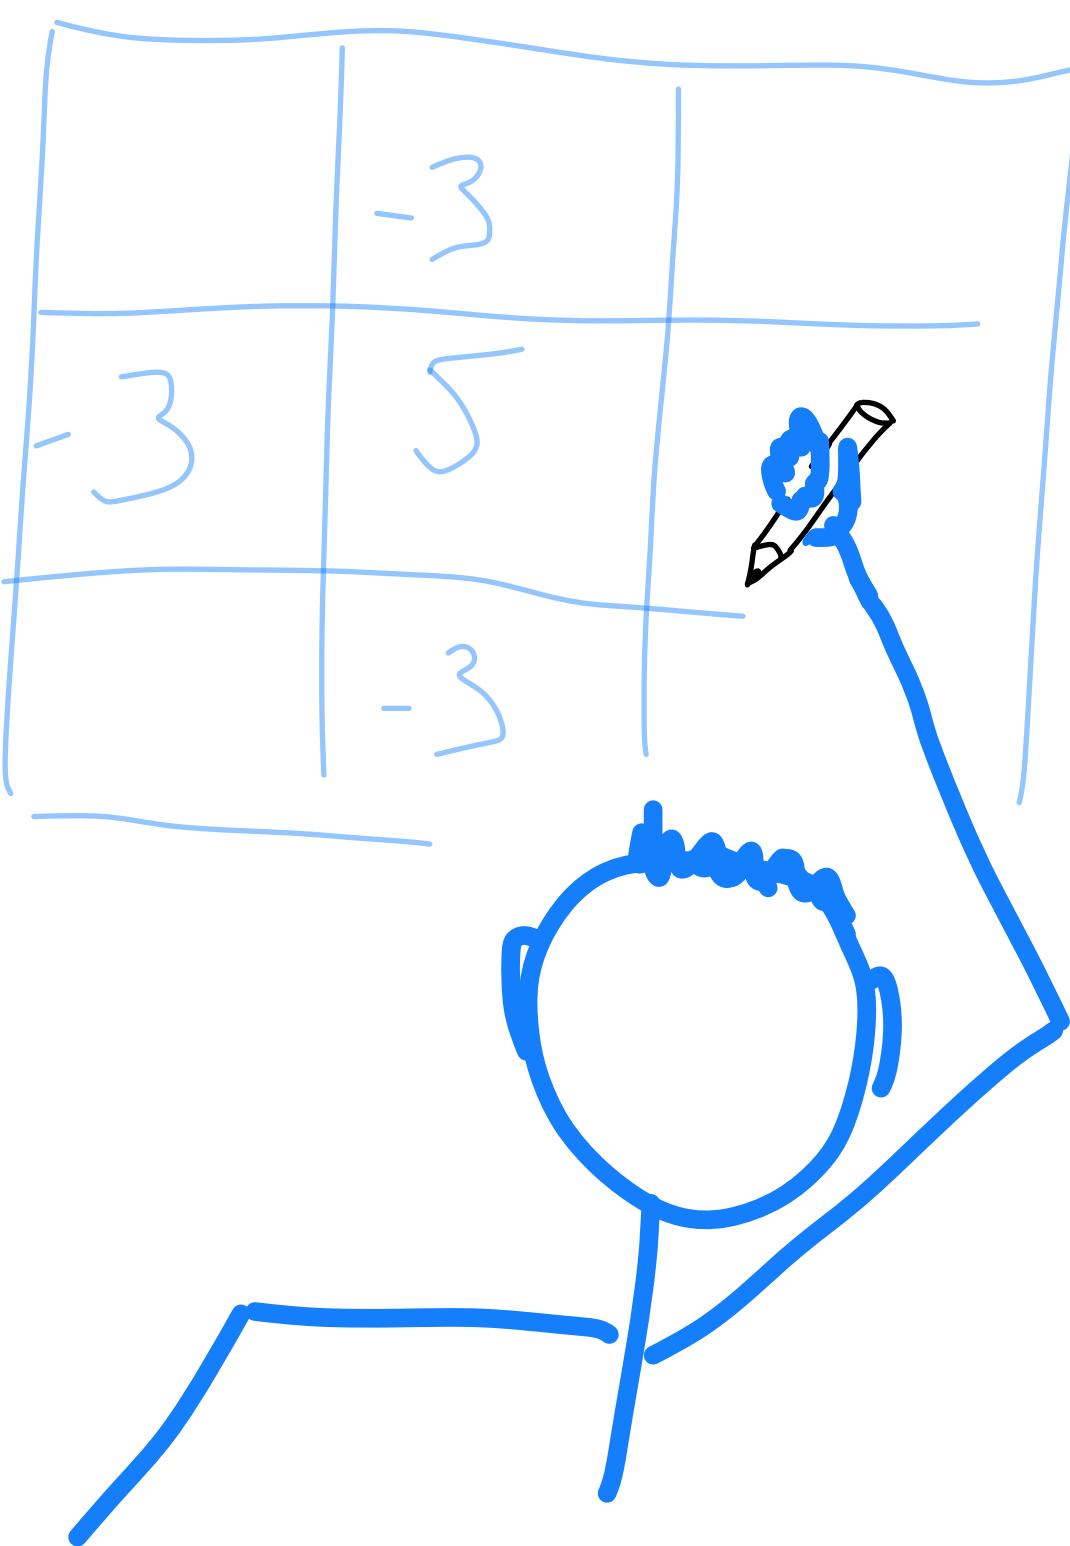
Filtering and Kernels

Prof. Dr. Marc Aubreville, SS 2022

15.03.2022

Motivation: In images, position matters!





Convolutional
kernels

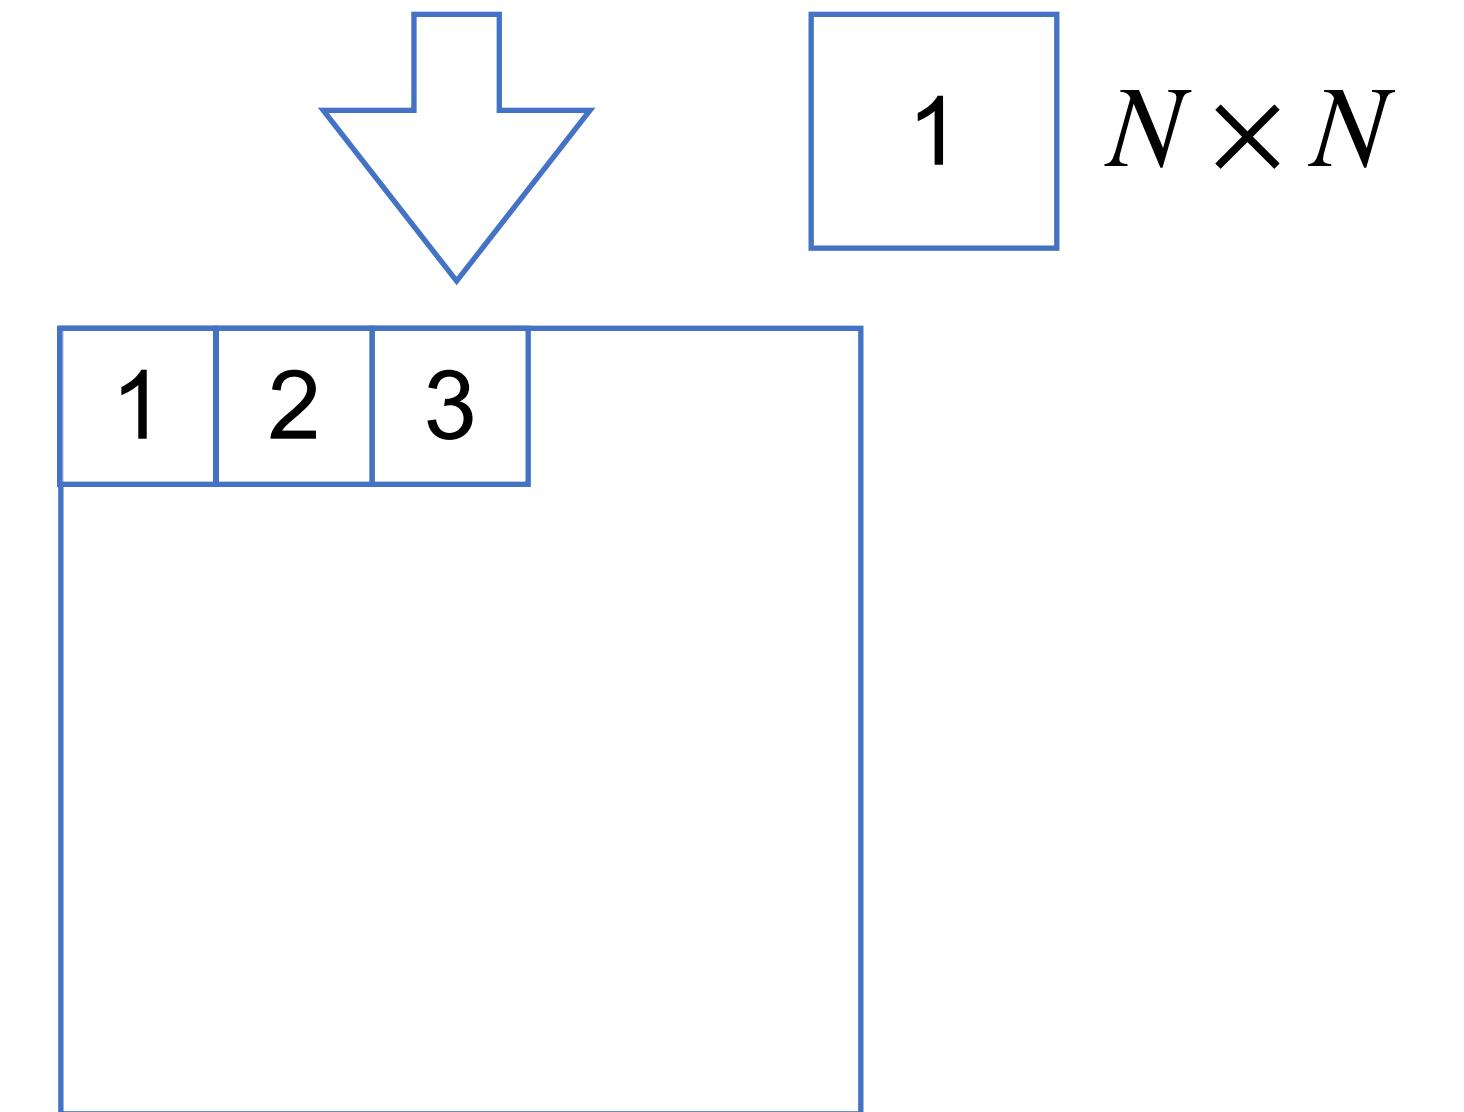
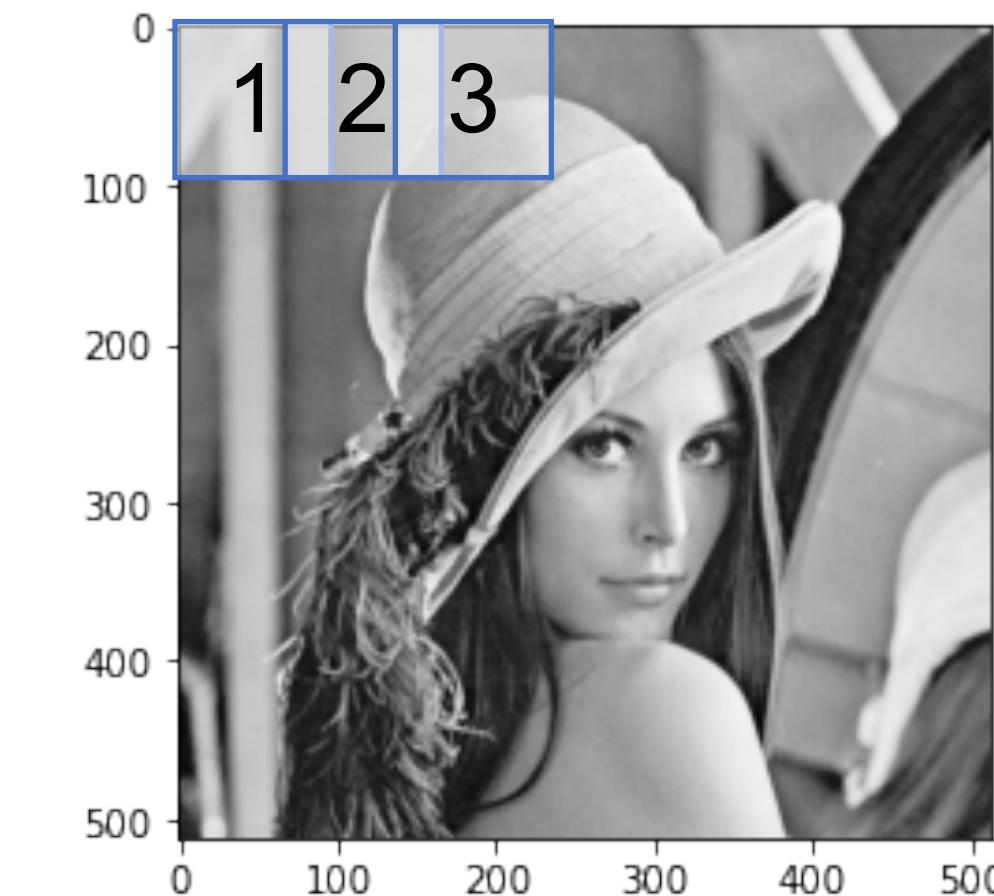


- So far, every operation that we have been doing, has been performed on **independent pixels** only
- The location of these pixels in the image could have been randomly changed without affecting the operations
- We used the values of the pixels from the image to derive characteristics and we have seen some modifications that can be performed on these pixels

In many cases, we are not only interested in single pixel values,
but also in the values of its neighbors!

Motivation: The convolution

- We want to be able to apply an operation that extracts the relationships between neighboring pixels.
- We call the result of this operation a descriptor, because it describes the pixel relationship
- Naively, this operation should always consider a local environment (e.g., N pixels squared)
- And the operation should be applied to all such local environments.





- The convolution of two functions is defined as:

$$s(t) * h(t) = \int_{\mathbb{R}^N} s(x + \tau) - h(\tau) d\tau \quad \text{or, in discrete terms: } s(n) * h(n) = \sum_u s(n + u) - h(u)$$

- We can extend this to a two-dimensional convolution

$$s(x, y) * h(u, v) = \sum_u \sum_v s(x + u, y + v) \cdot h(u, v)$$

Convolution (cont'd)

0	0	0	0
0	85	255	0
<hr/>			
0	255	255	0
0	0	0	0

Image S

1	1	1
1	1	1
1	1	1

Filter kernel h

size: $m \times m$

- It is sensible to normalize the convolution by the size of the filtering element (kernel):

$$s(x, y) * h(u, v) = \frac{1}{m^2} \sum_u \sum_v s(x + u, y + v) \cdot h(u, v)$$

- For convenience, let's shift the equation so the kernel is centered:

$$s_{out}(x, y) = \frac{1}{m^2} \sum_{u=0}^{m-1} \sum_{v=0}^{m-1} s_{in}(x + k - u, y + k - v) \cdot h(u, v)$$

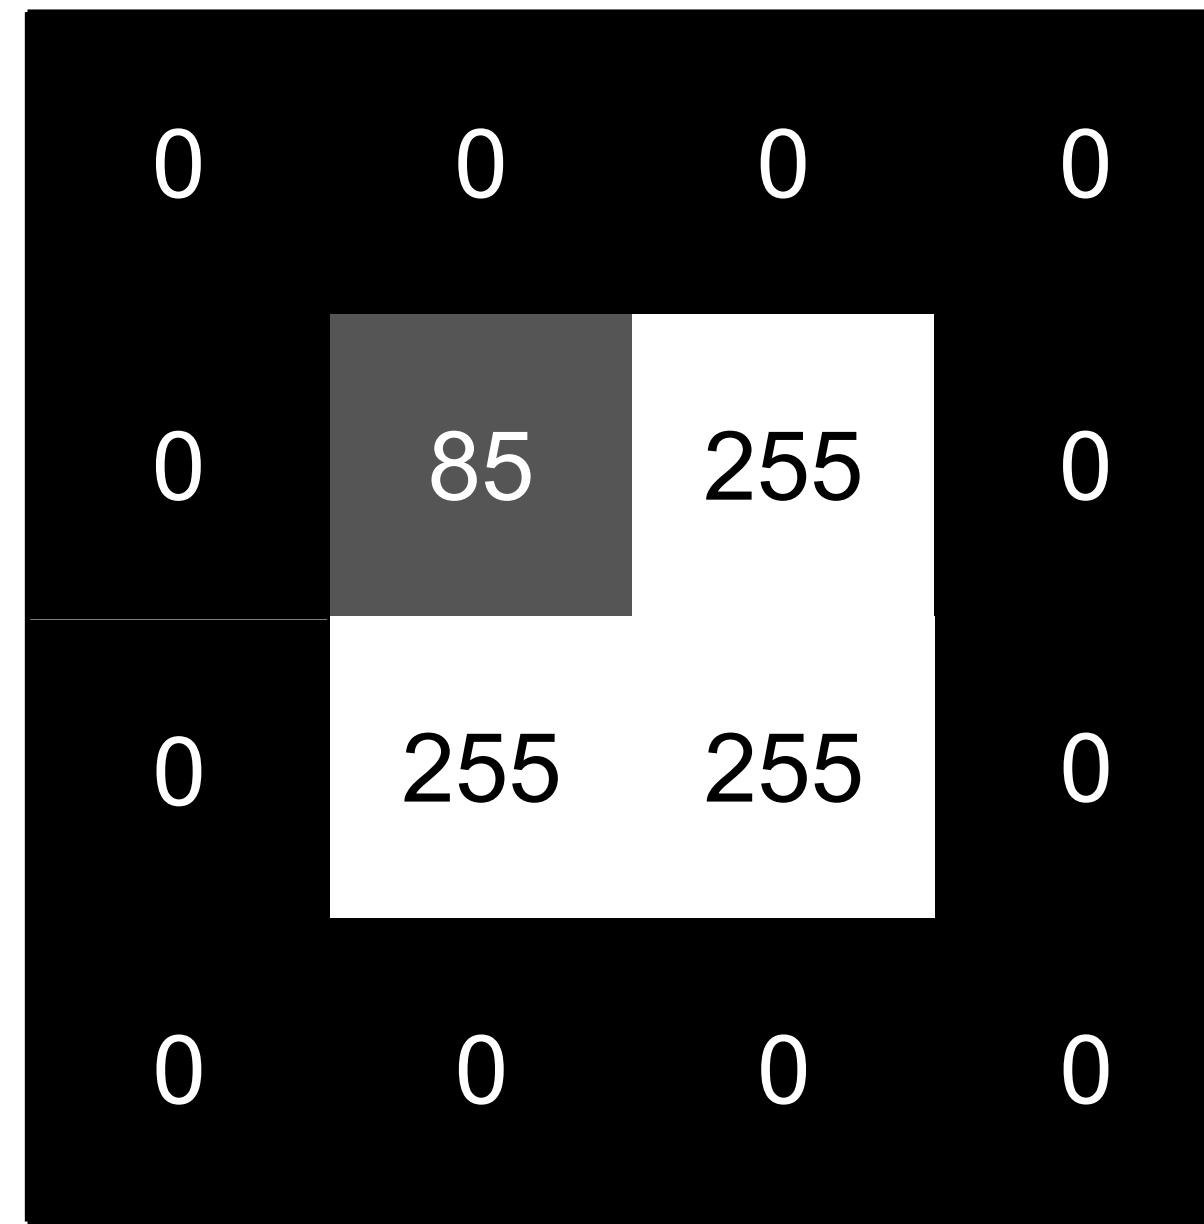
$$k = \frac{m - 1}{2}$$

Note: To have a clear center definition, odd sized kernels are sensible.

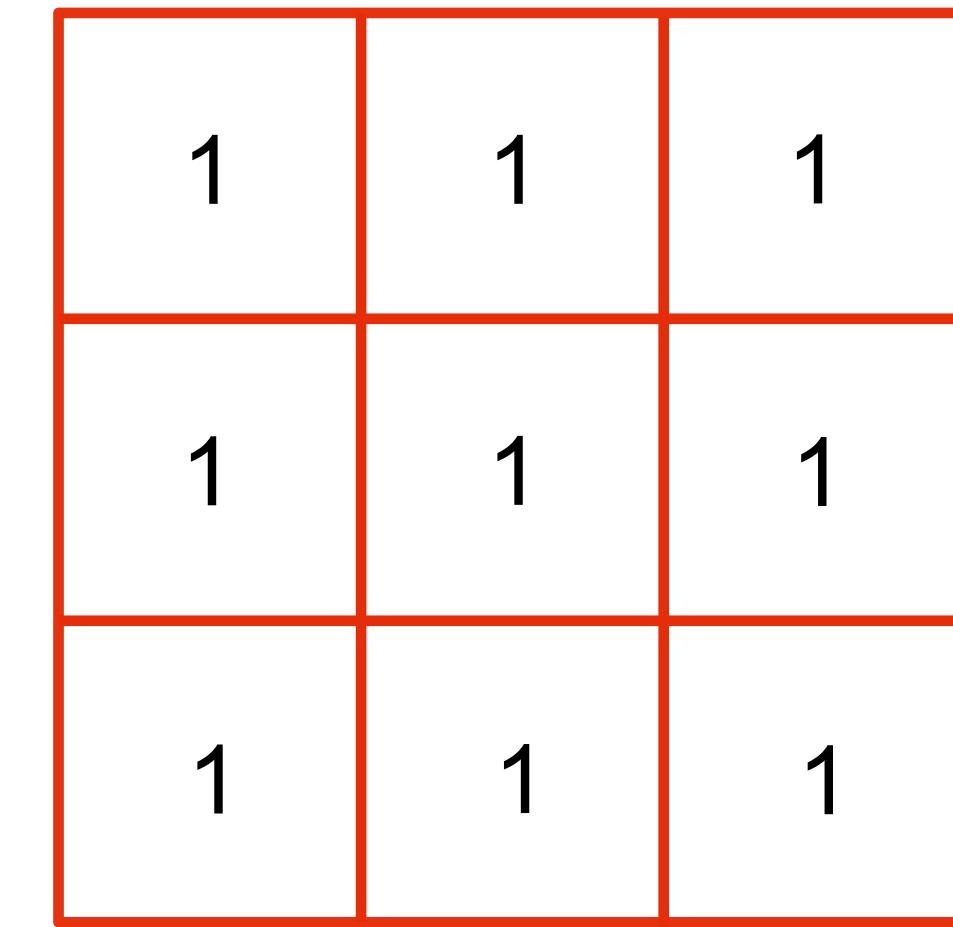
Kernels

The general idea

4x4 Image s_e



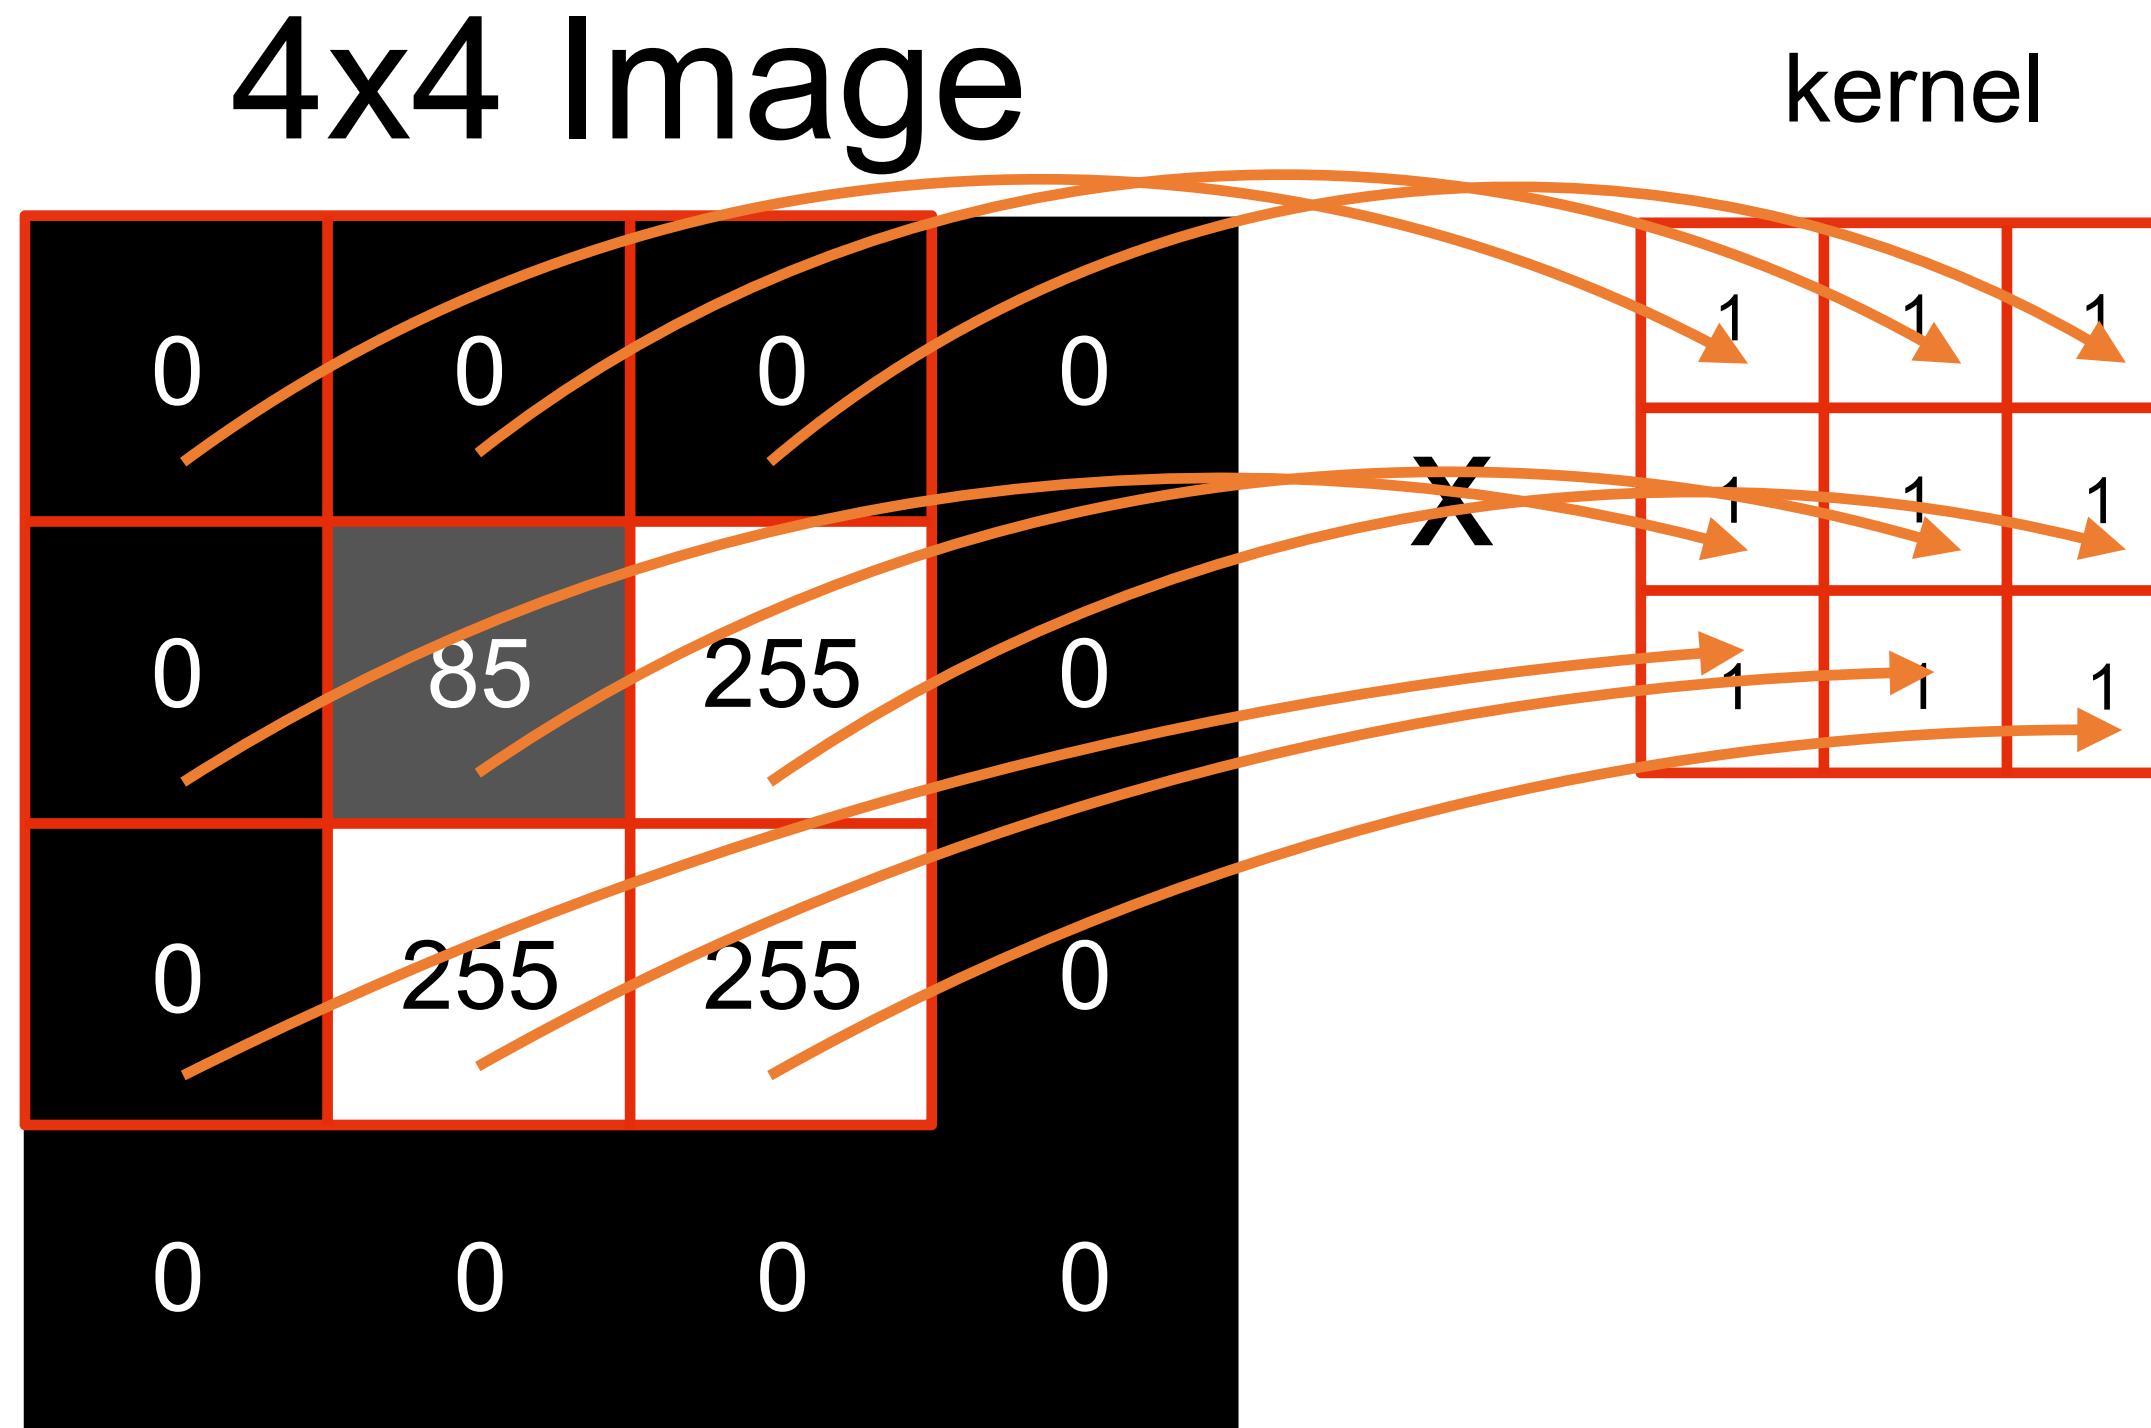
3x3 Kernel $h(u, v)$



$$s_{out}(x, y) = \frac{1}{m^2} \sum_{u=0}^{m-1} \sum_{v=0}^{m-1} s_{in}(x + k - u, y + k - v) \cdot h(u, v)$$

where m is the length of the mask

Convolutional kernels: Example



$$\begin{aligned} & 0 * 1 + 0 * 1 + 0 * 1 \\ & + 0 * 1 + 85 * 1 + 255 * 1 + 0 * 1 + 255 * 1 + 255 * 1 \\ & = 85 + 255 + 255 + 255 \\ & = 850 \end{aligned}$$

$$\frac{850}{9} = 94$$

Convolutional Kernels: Example (2)

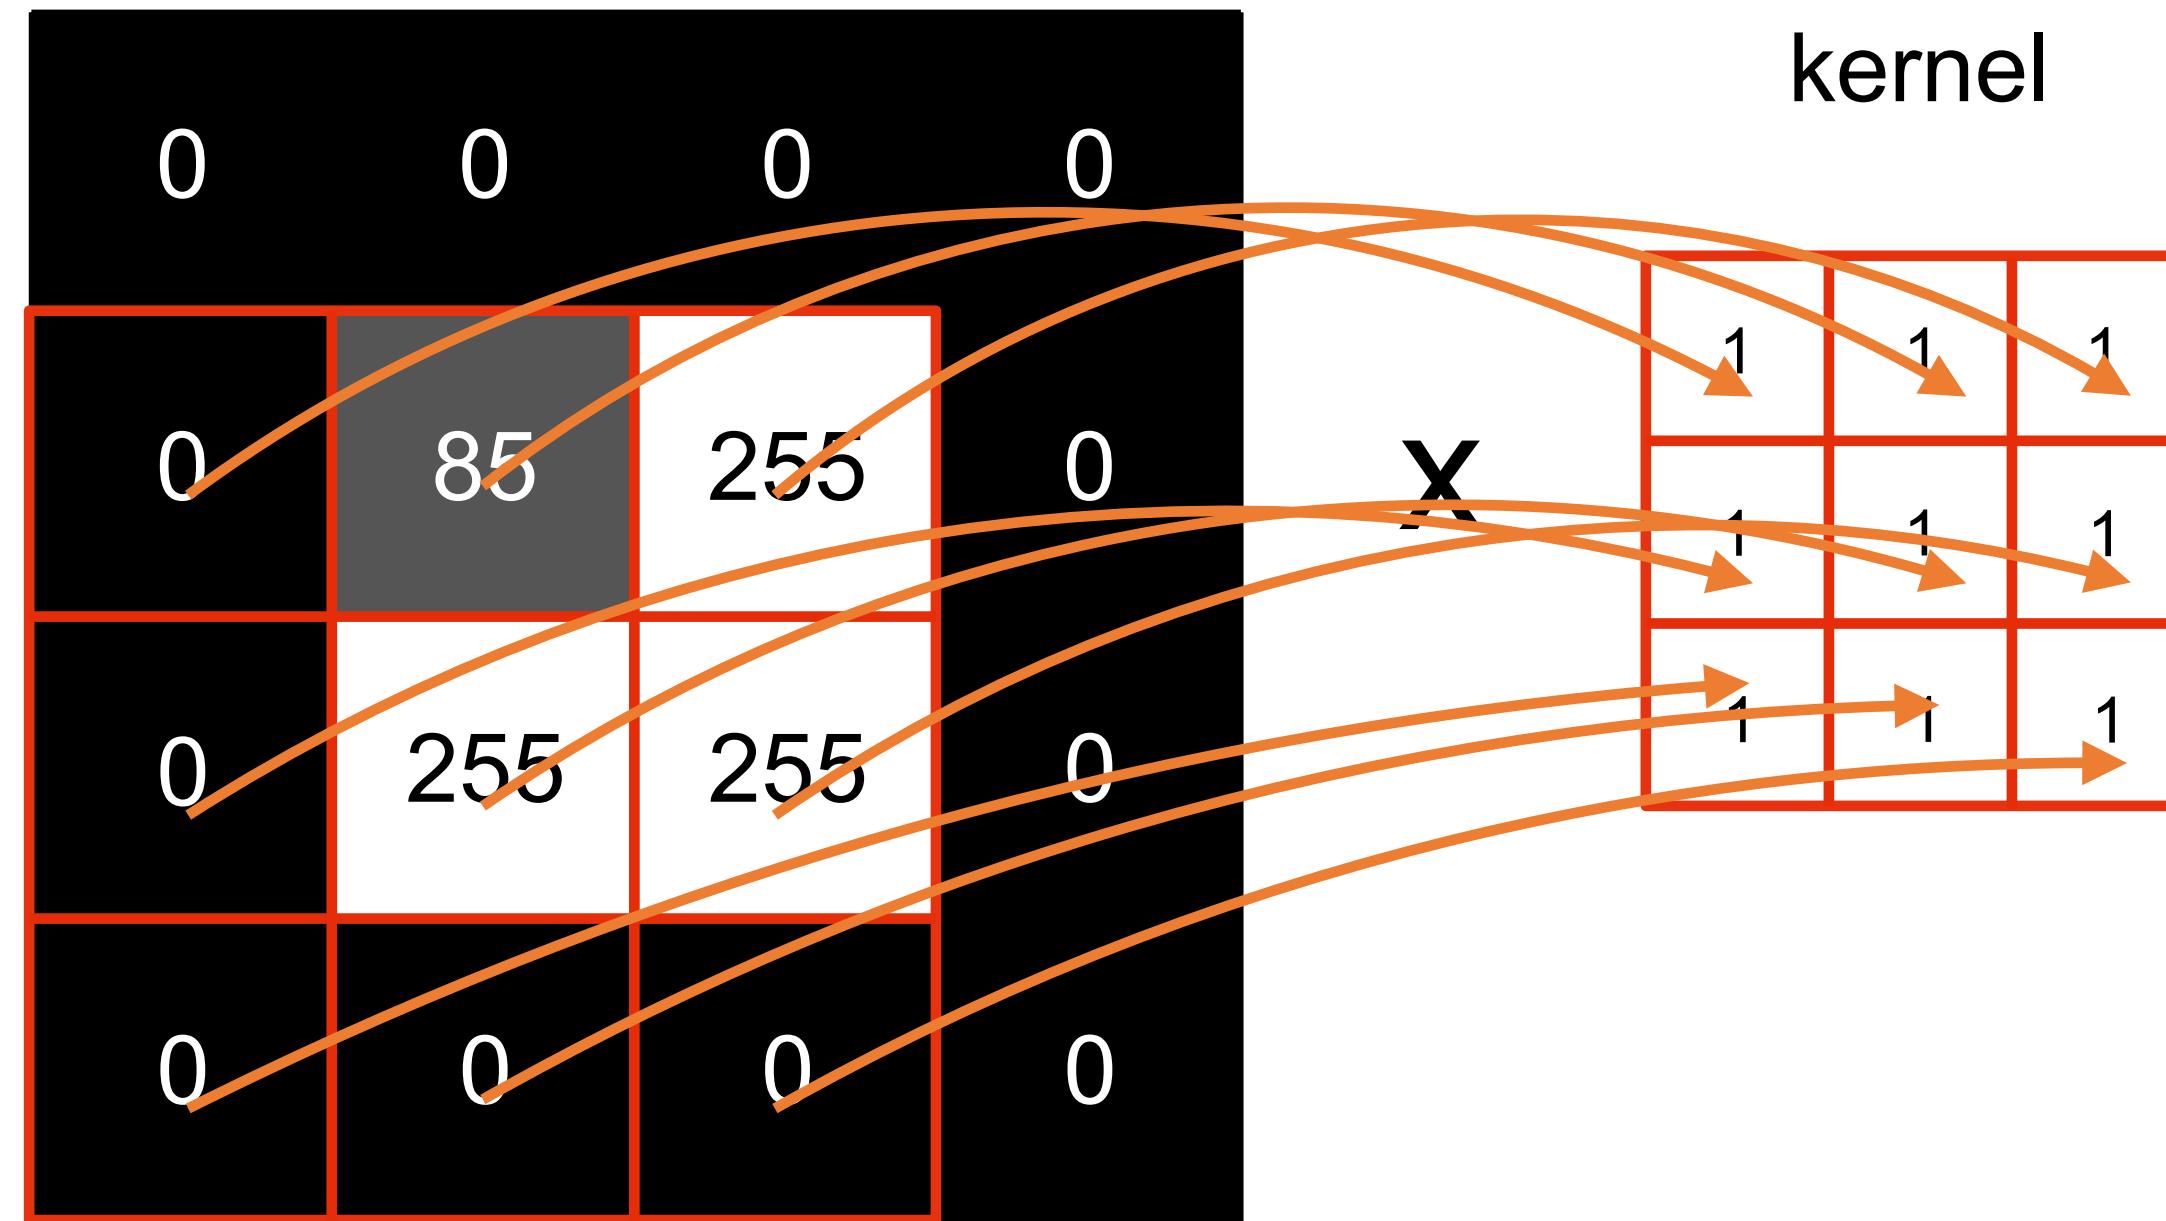


$$\begin{aligned} & 0 * 1 + 0 * 1 + 0 * 1 \\ & + 85 * 1 + 255 * 1 + 0 * 1 + 255 * 1 + 255 * 1 + 0 * 1 \\ & = 85 + 255 + 255 + 255 \\ & = 850 \end{aligned}$$

$$\frac{850}{9} = 94$$

Convolutional Kernels: Example (3)

4x4 Image



$$0 * 1 + 85 * 1 + 255 * 1 + 0 * 1 + 255 * 1 + 255 * 1 + 0 * 1 + 0 * 1 + 0 * 1$$

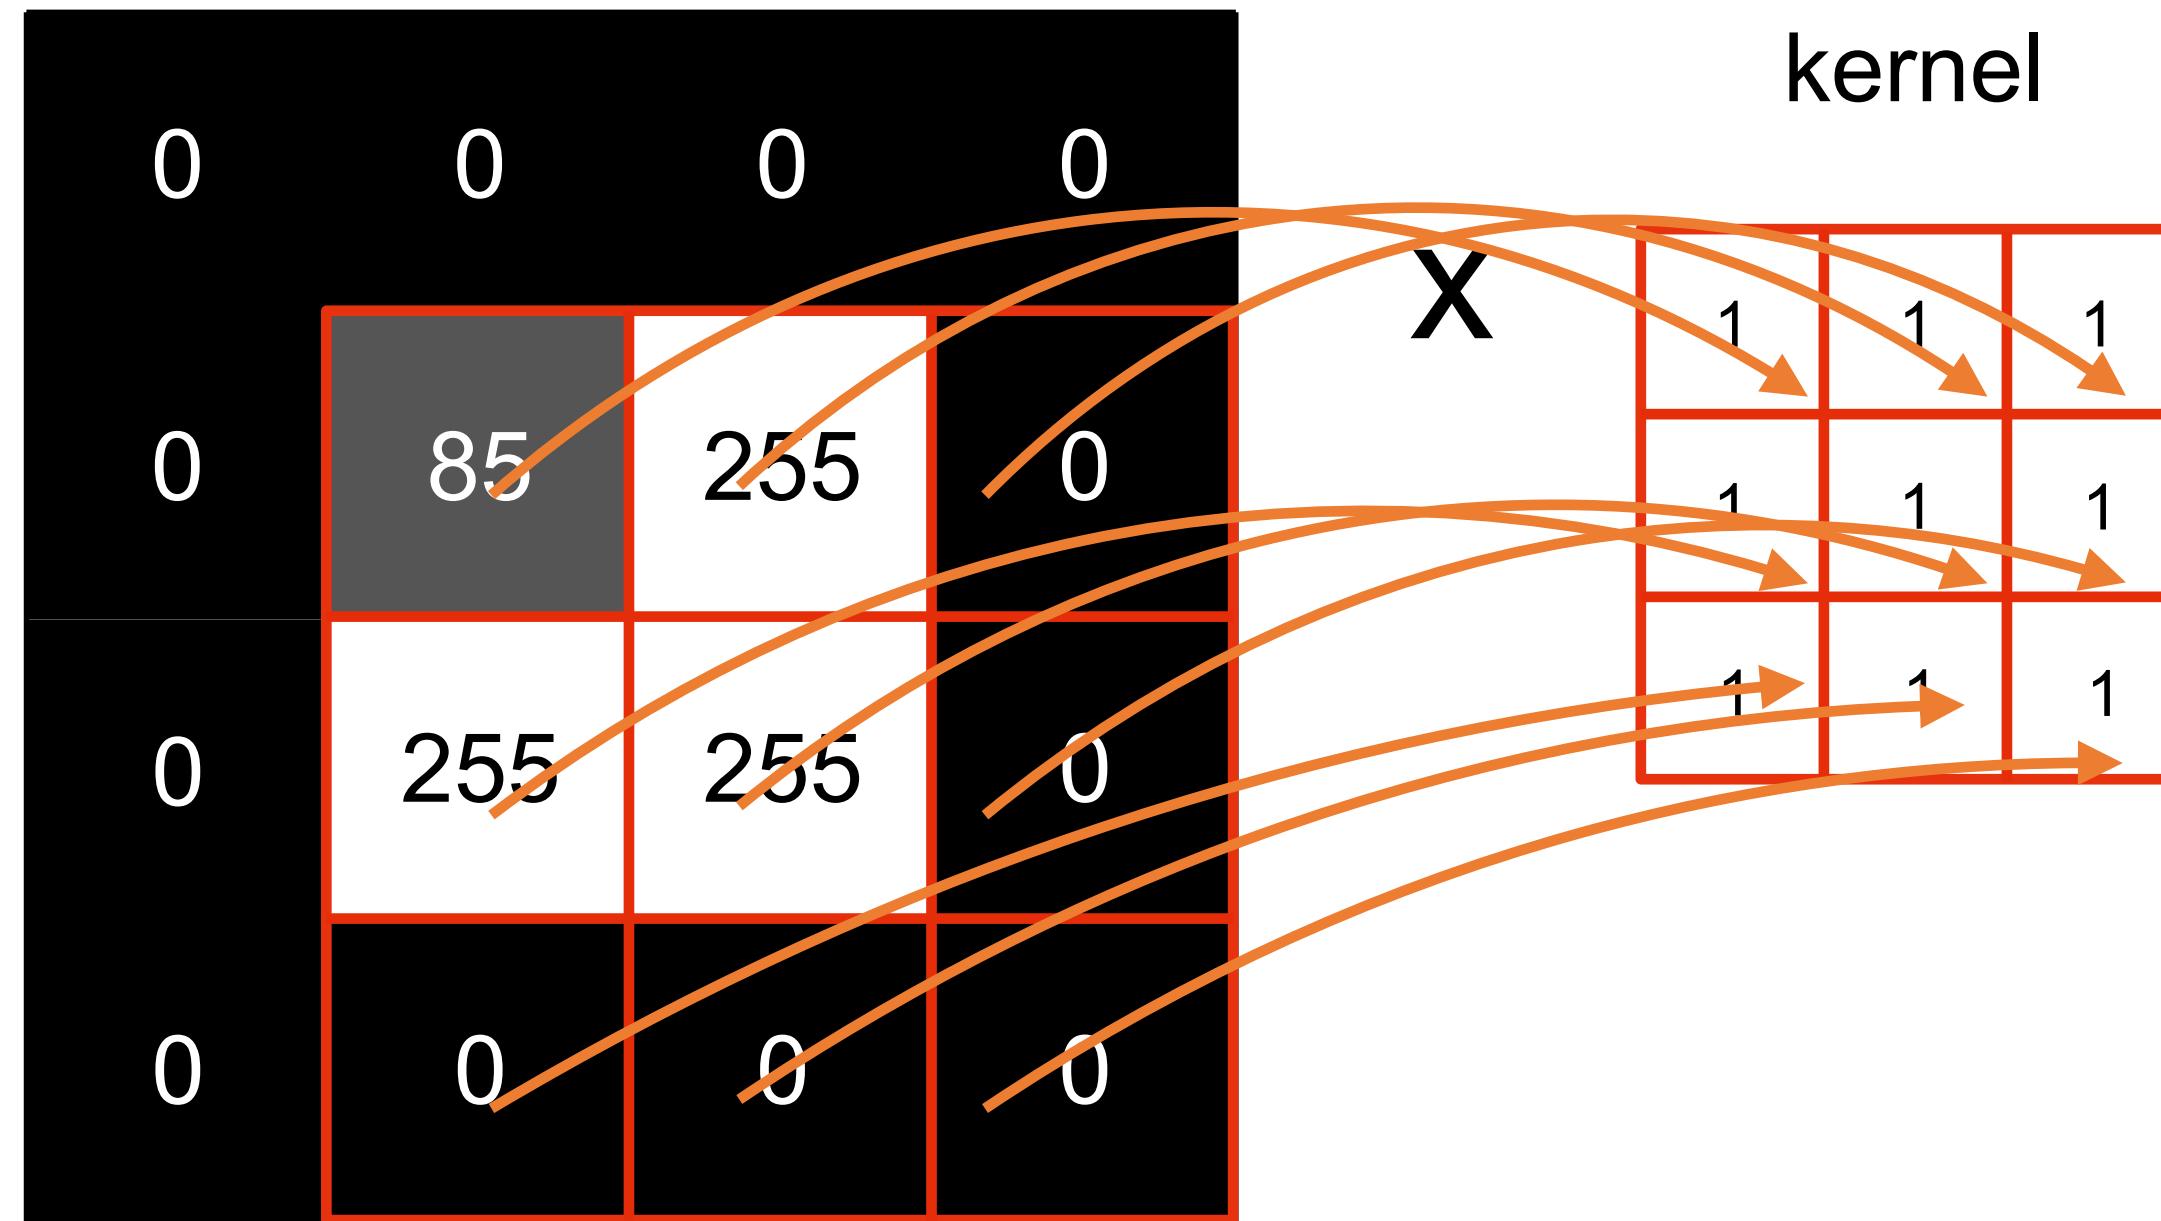
$$= 85 + 255 + 255 + 255$$

$$= 850$$

$$\frac{850}{9} = 94$$

Convolutional Kernels: Example (4)

4x4 Image



$$85 * 1 + 255 * 1 + 0 * 1 + 255 * 1 + 255 * 1 + 0 * 1 + 0 * 1 + 0 * 1 + 0 * 1$$

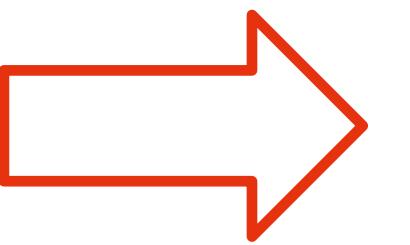
$$= 85 + 255 + 255 + 255$$

$$= 850$$

$$\frac{850}{9} = 94$$

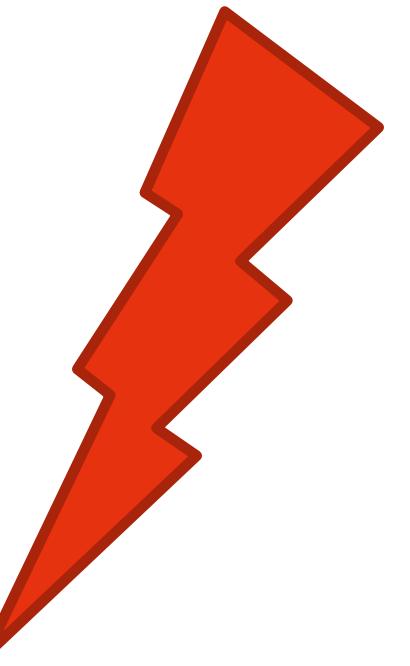
The general idea

0	0	0	0
0	85	255	0
0	255	255	0
0	0	0	0



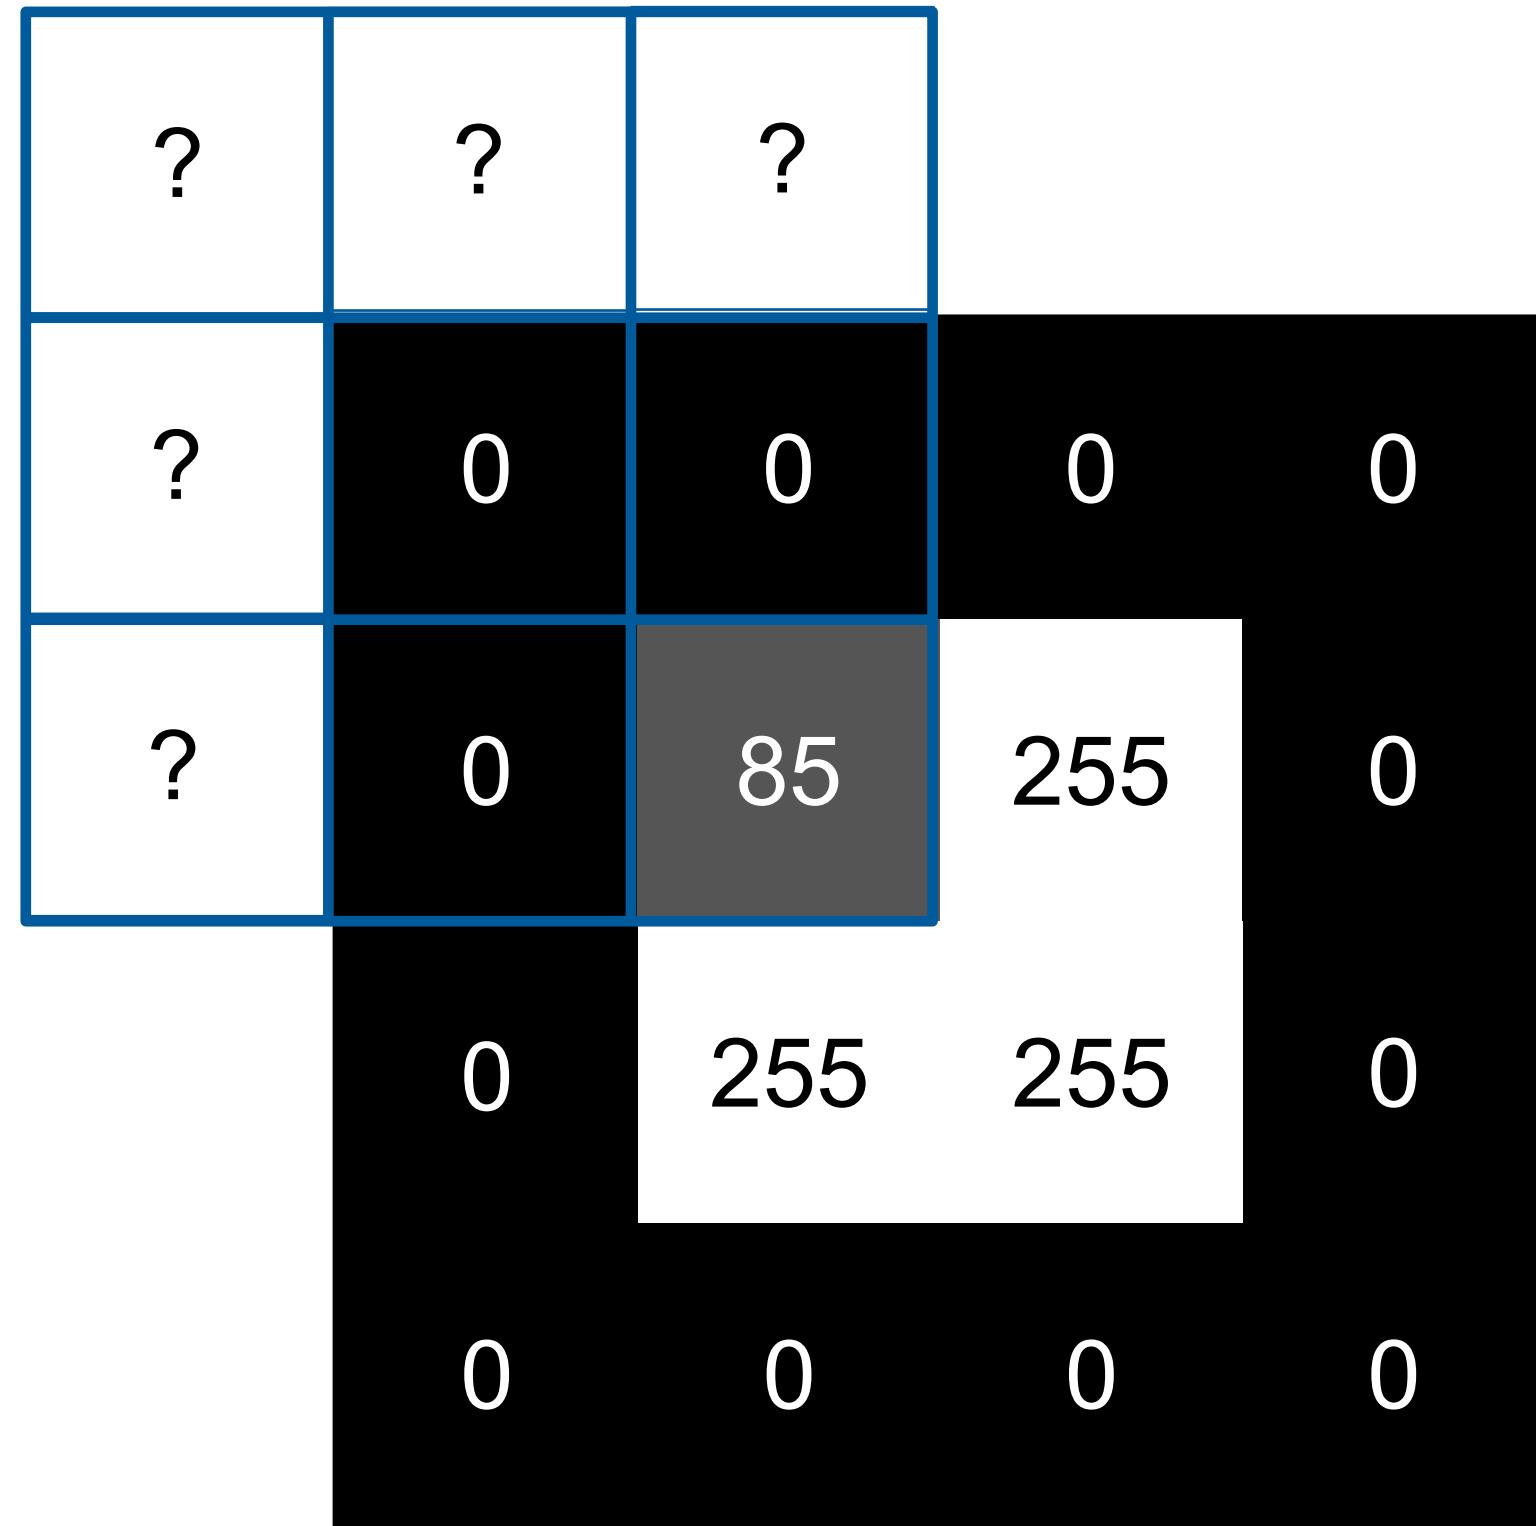
0	0	0	0
0	94	94	0
0	94	94	0
0	0	0	0

We skipped all
the border
pixels

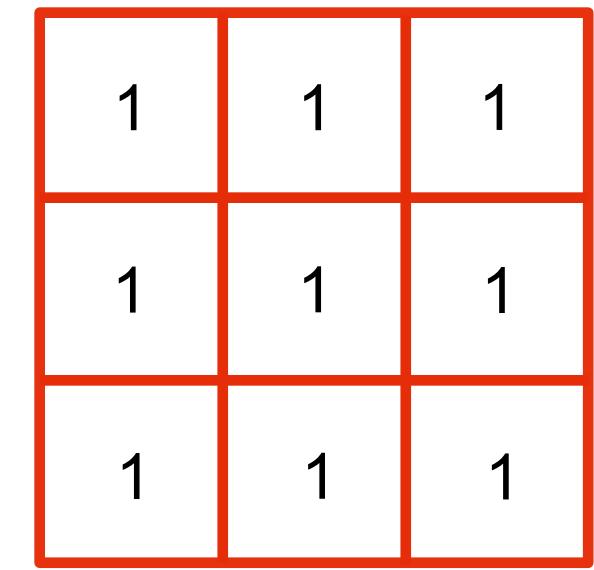


Kernels

What can we do at the borders?

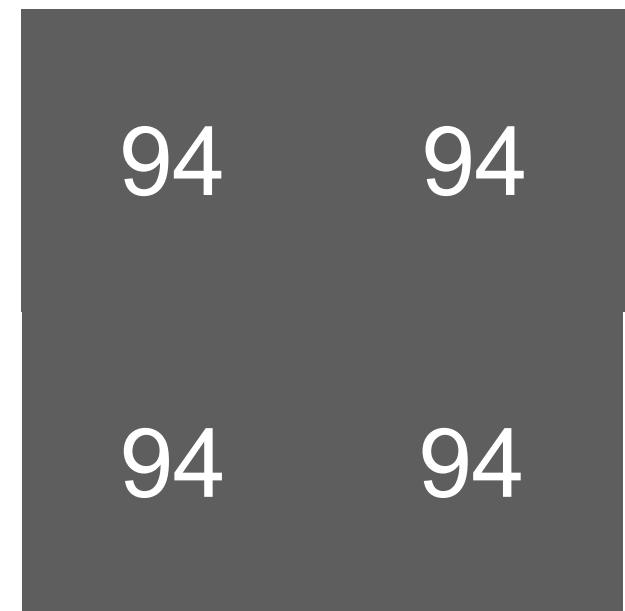


X



Kernels

What can we do at the borders?



1st option:

- Skip the outer pixels
- Image gets smaller

Kernels

What can we do at the borders?

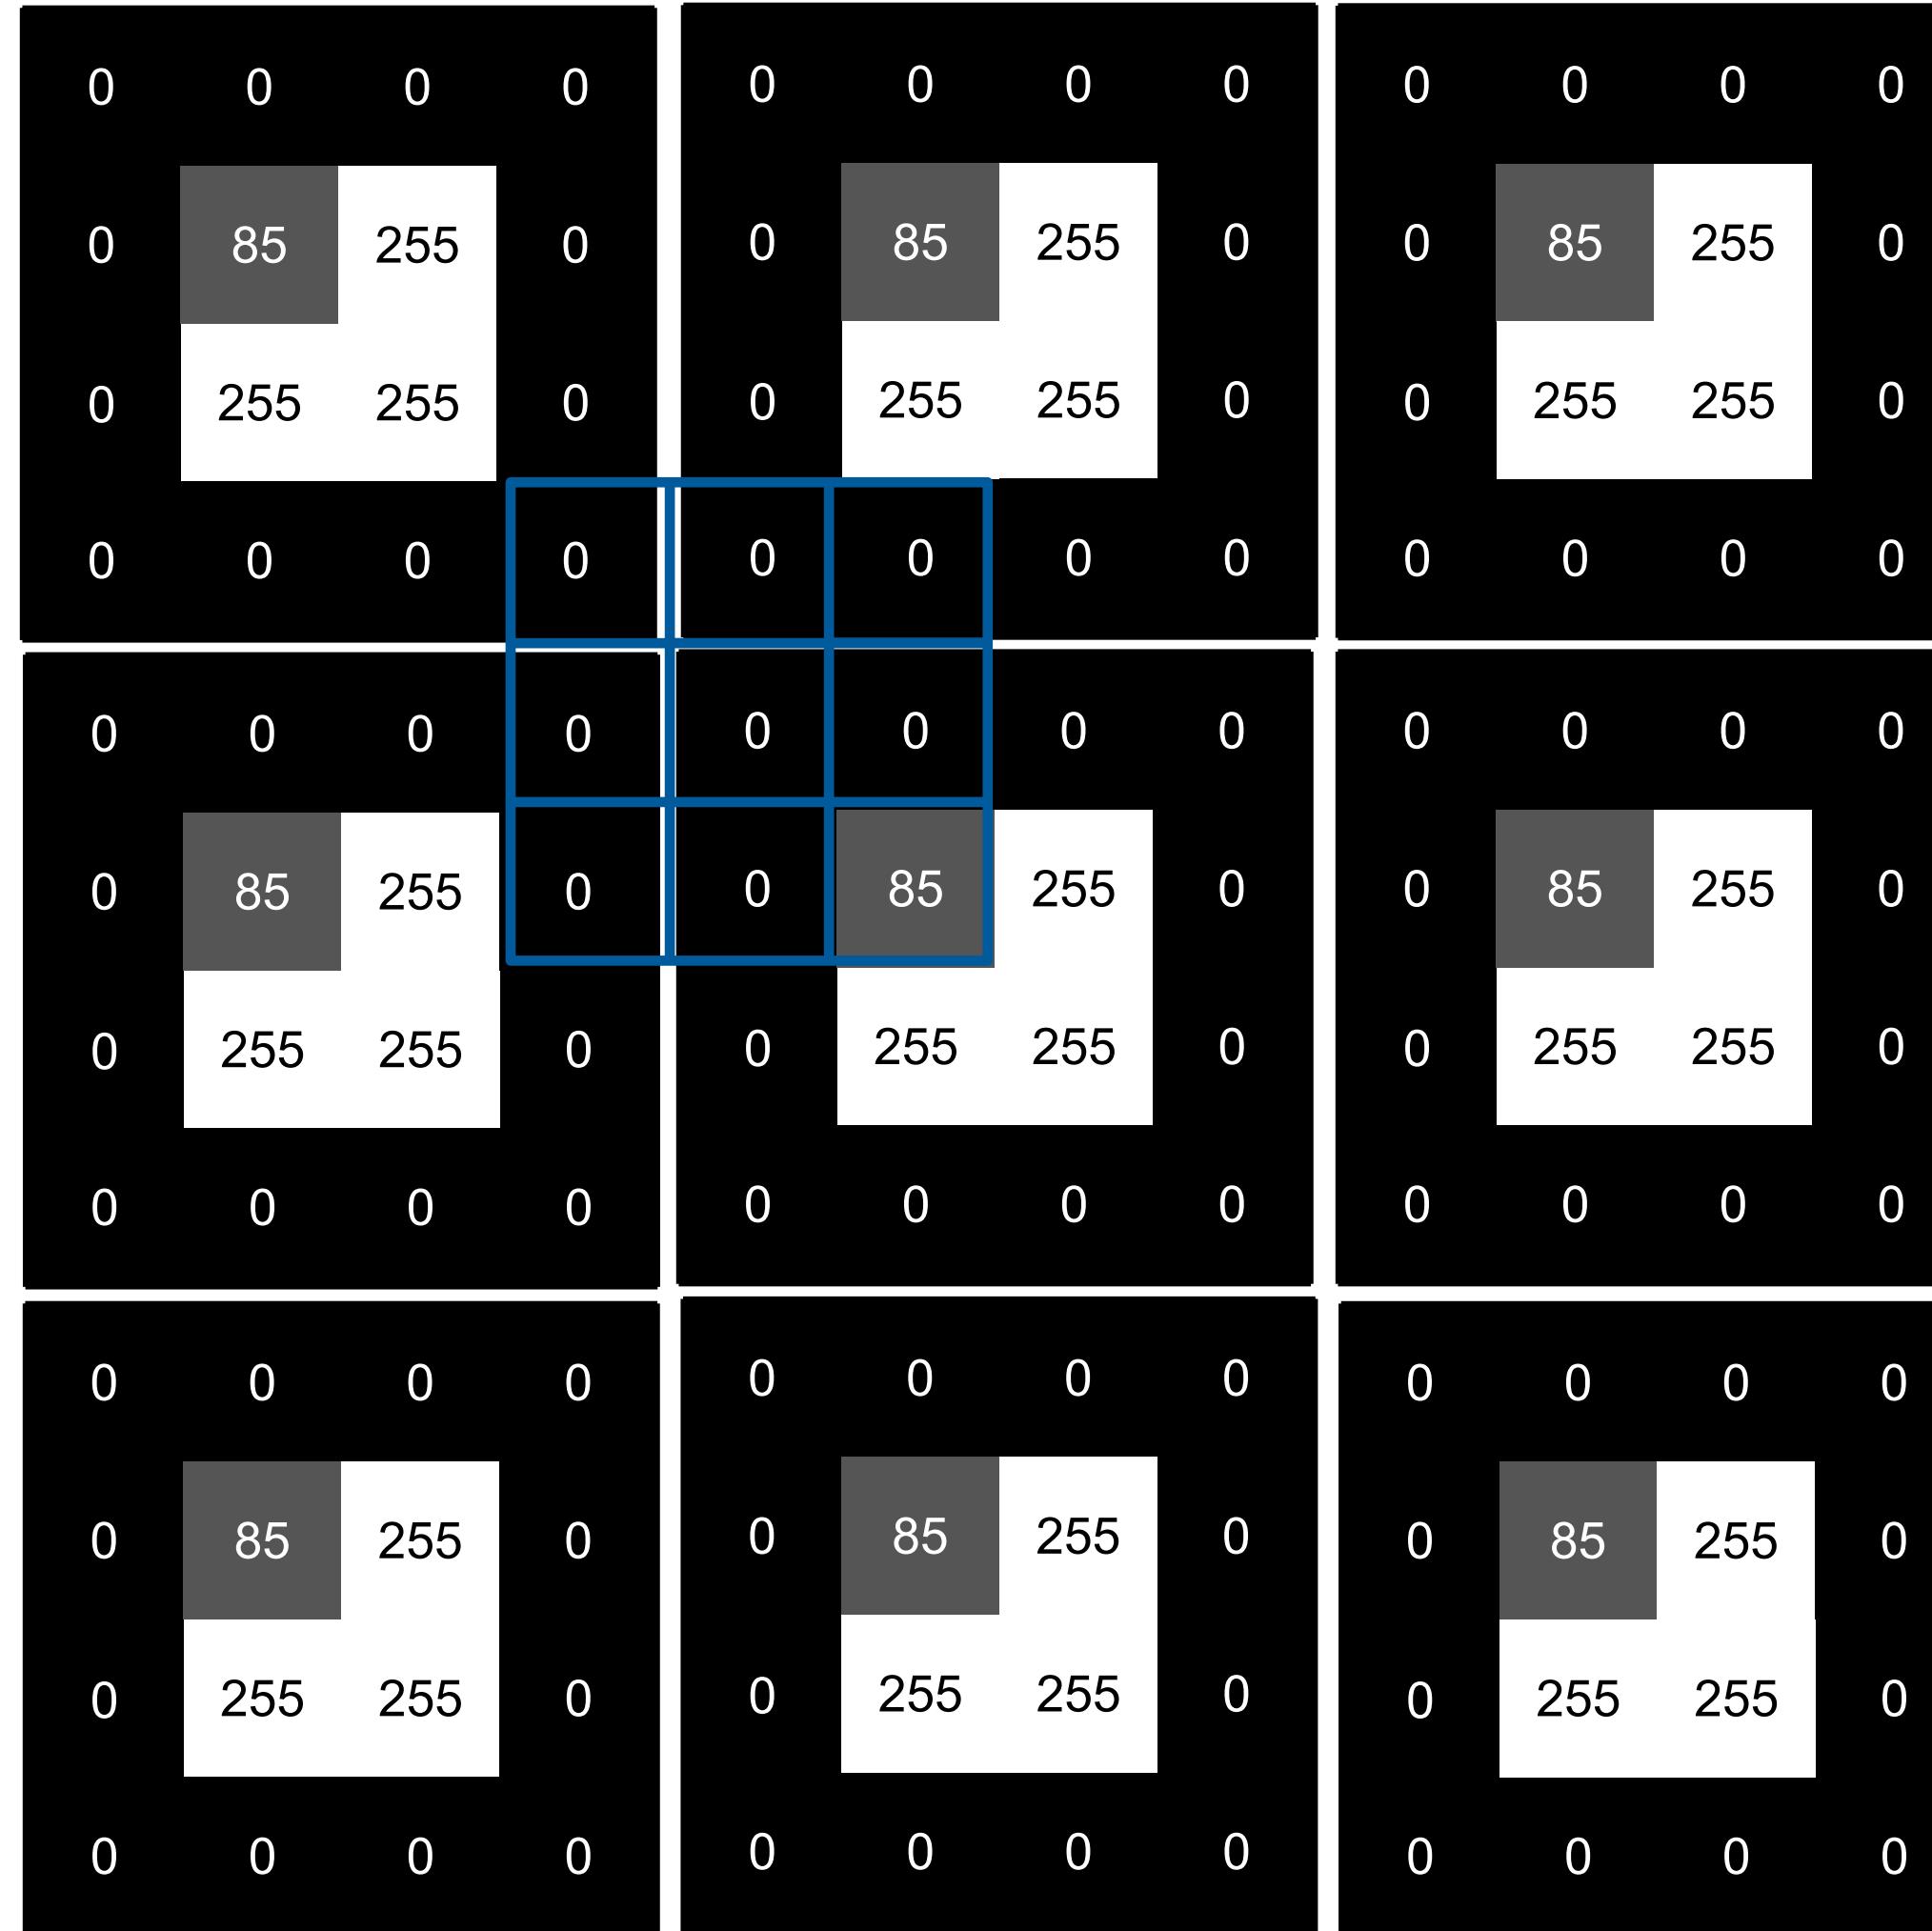
0	0	0	0	0	0
0	0	0	0	0	0
0	0	85	255	0	0
0	0	255	255	0	0
0	0	0	0	0	0
0	0	0	0	0	0

2nd option: Extend:

- Copy the border pixels
- Image keeps same size

Kernels

What can we do at the borders?



3nd option: Concatenate:

→ Add the image next to each other

Kernels

What can we do at the borders?



4th Option: Mirror:

→ Add a mirrored image to each side

Kernels

What can we do at the borders?



4th Option: Mirror:

→ Add a mirrored image to each side

Kernels

What can we do at the borders?

$$\begin{array}{c}
 \begin{array}{ccccc}
 \text{X} & 0 & 0 & 0 & 0 \\
 0 & 0 & 85 & 255 & 0 \\
 0 & 255 & 255 & 0 & 0
 \end{array} \\
 \times \\
 \begin{array}{ccccc}
 1 & 1 & 1 \\
 1 & 1 & 1 \\
 1 & 1 & 1
 \end{array}
 \end{array}$$

The diagram illustrates a convolution operation. On the left is a 5x5 input image (X) with values: top-left (0), top-middle (0), top-right (0), middle-left (0), middle-middle (85), middle-right (255), bottom-left (0), bottom-middle (255), bottom-right (255). The bottom row of the input is zeroed out. To its right is a 3x3 kernel with all elements set to 1. A blue 'X' symbol indicates the center of the kernel. The result of the convolution is shown on the right, where the output value is 85 (the sum of the kernel elements multiplied by the central input value).

5th option: kernel crop

→ Crop the kernel according to the available pixels in the image.

Kernel normalization

- We defined the convolution to include a normalization of the kernel

$$s(x, y) * h(u, v) = \frac{1}{m^2} \sum_u \sum_v s(x + u, y + v) \cdot h(u, v)$$

- This is sensible, since the value range of pixels (0..255) shall be retained in the filtered image.
- Dividing by m^2 actually ensures that we don't exceed the value limit of 255.
- This also makes applications of (e.g.) cropped filters possible (different normalization in the cropped case)

Kernel normalization (cont'd)

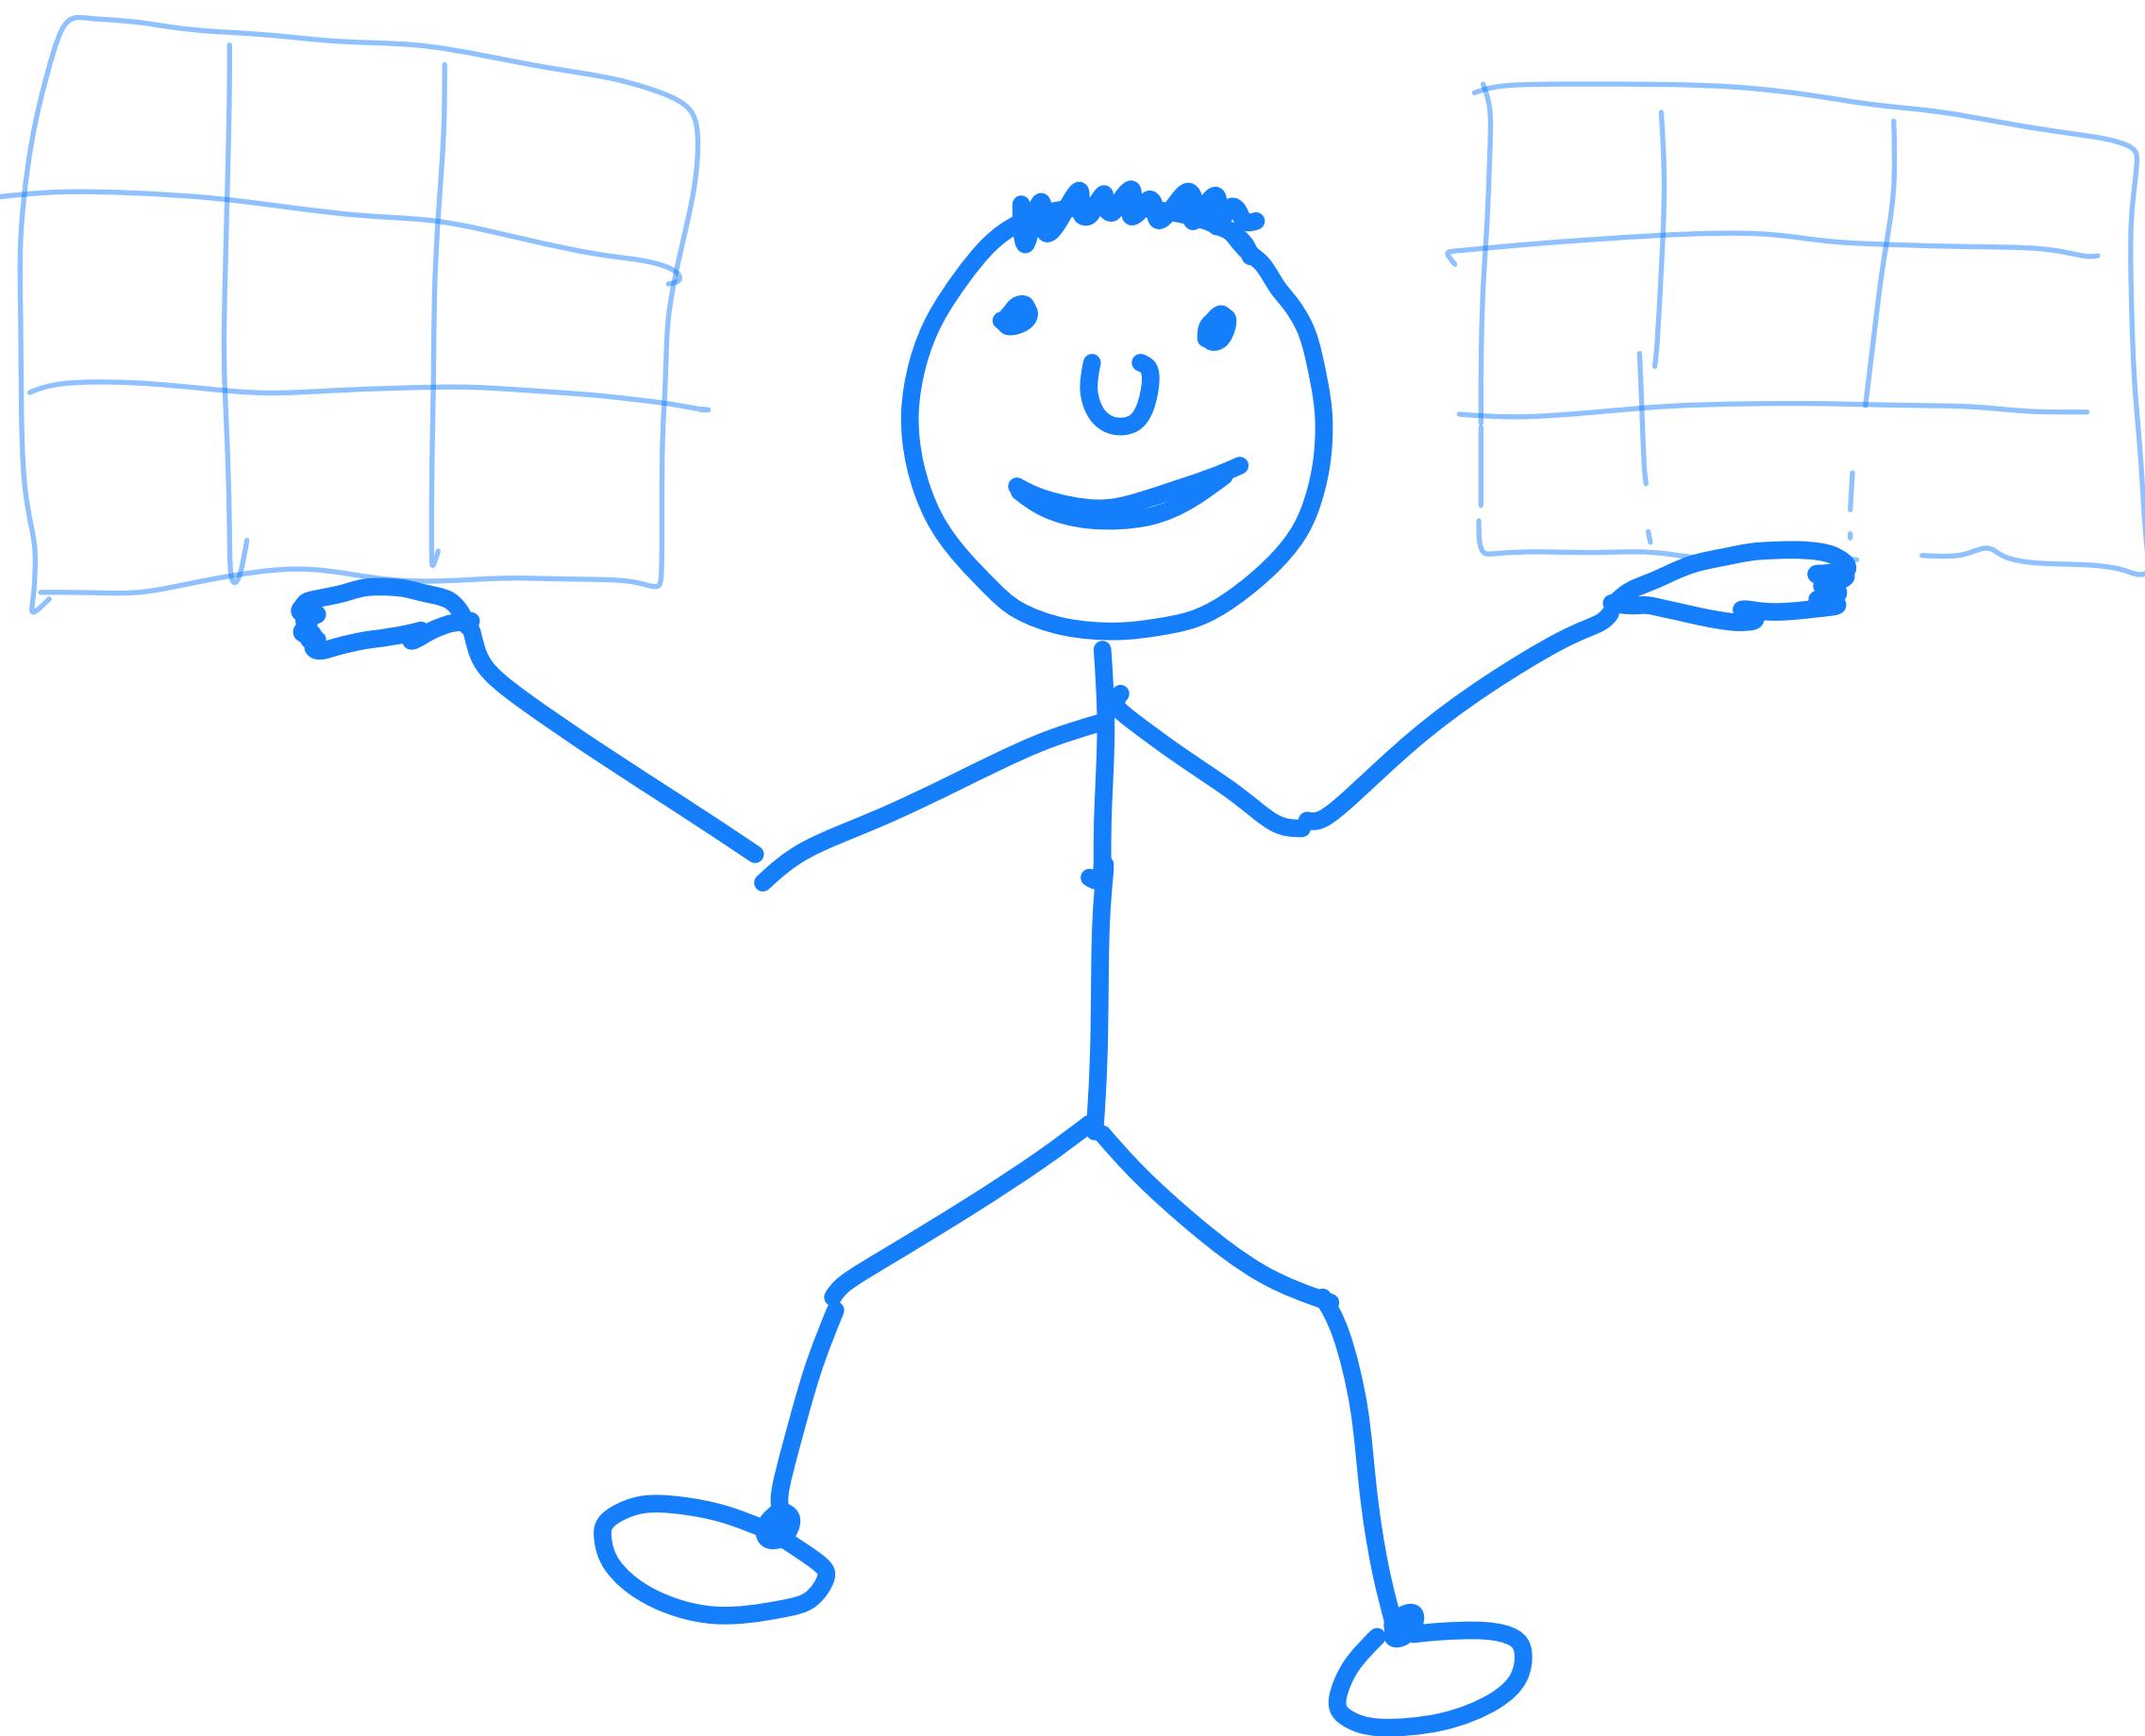
■ But what if the kernel is not only ones?

0	1	0
1	2	1
0	1	0

- Dividing by 9 would lead to a too low value range.
- Instead, we can divide by the sum of the kernel:

$$s_{out}(x, y) = \frac{1}{\sum_u \sum_v h(u, v)} \sum_{u=0}^{m-1} \sum_{v=0}^{m-1} s_{in}(x + k - u, y + k - v) \cdot h(u, v)$$

- For the example on the left, the normalization term would be $\frac{1}{6}$



Kernel

types

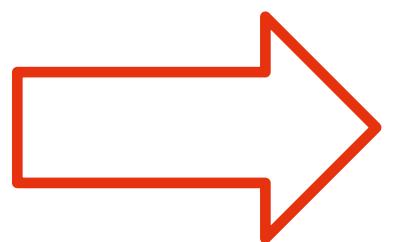
Kernel Types

Identity

0	0	0
0	1	0
0	0	0

Keeps the original image

Normalizing factor: $\frac{1}{1} = 1$



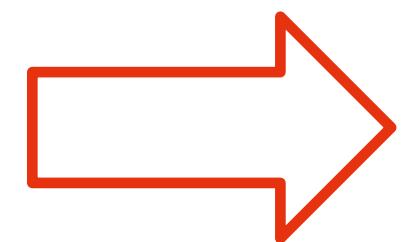
Kernel Types

Box Blur

1	1	1
1	1	1
1	1	1

Takes an **average** of all neighbors with **equal weight**

Normalizing factor: $\frac{1}{9}$



Kernel
15x15

Kernel Types

Gaussian Blur

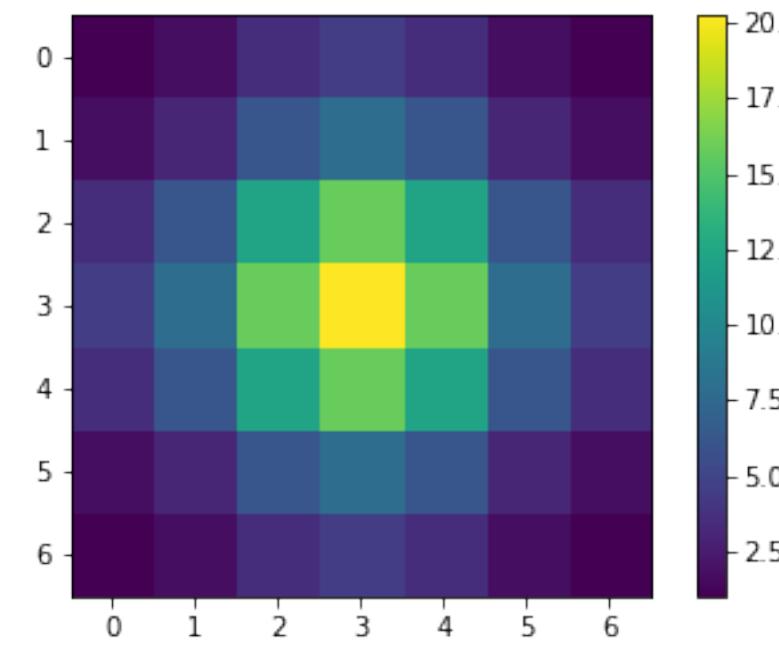
3x3 gaussian blur

1	2	1
2	4	2
1	2	1

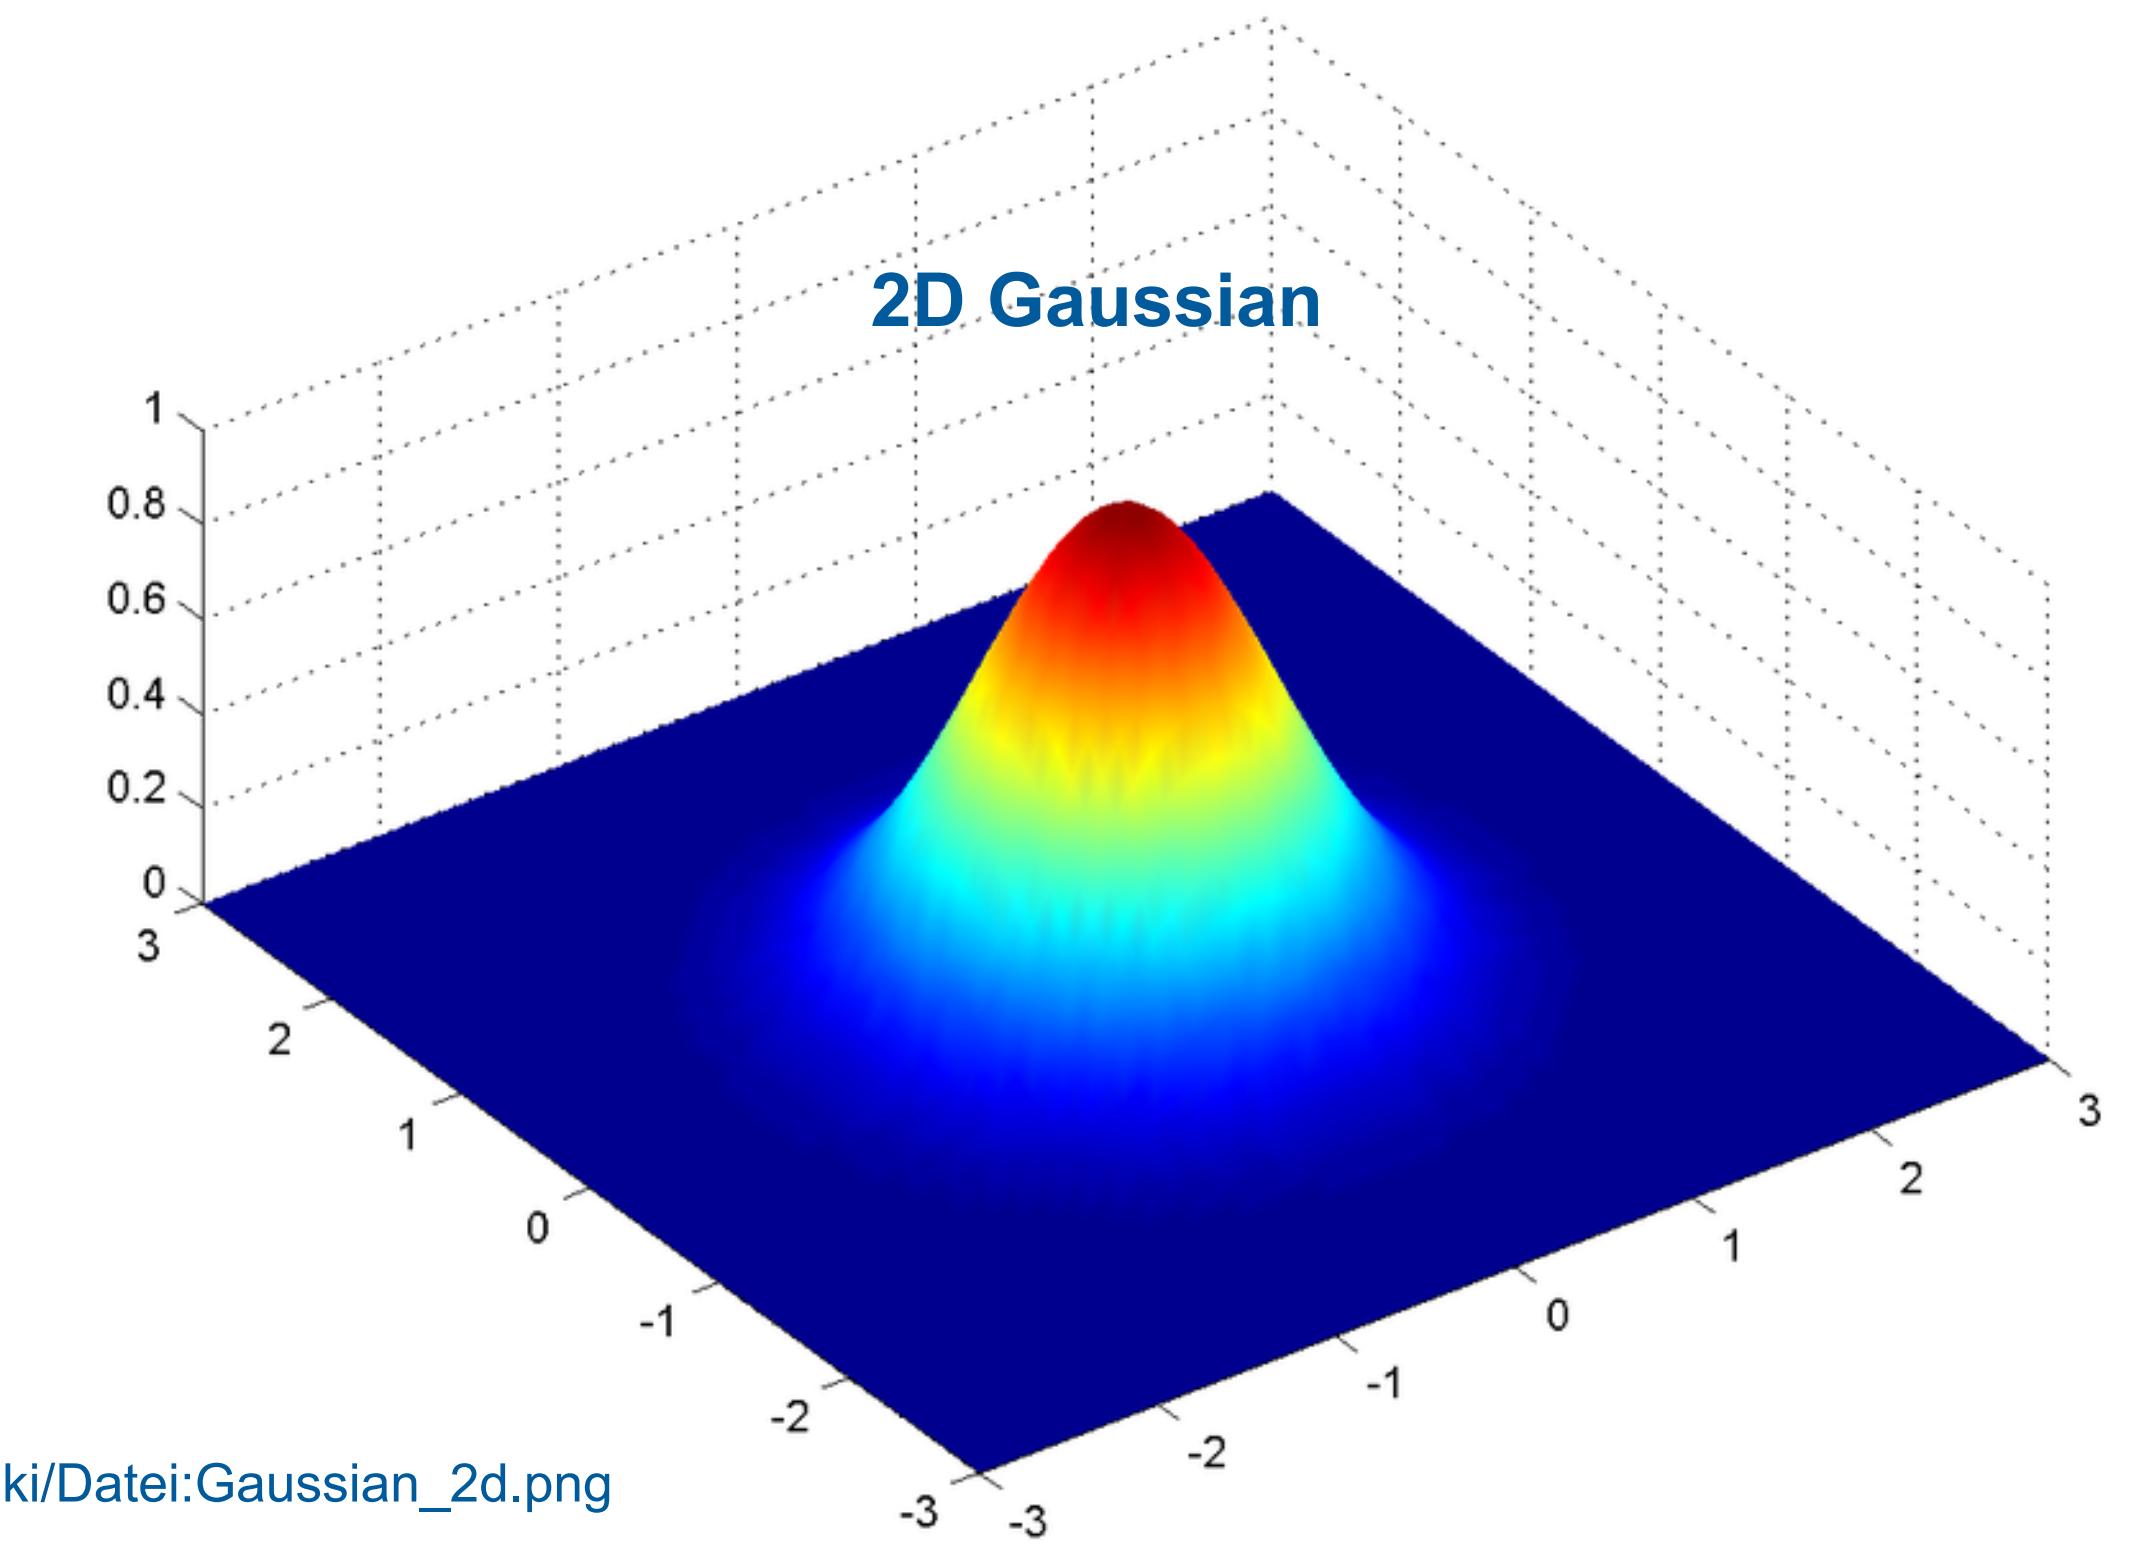
Takes an **average** of all neighbors with **approximated gauss distribution weight**

Normalizing factor: $\frac{1}{16}$

7x7 gaussian blur



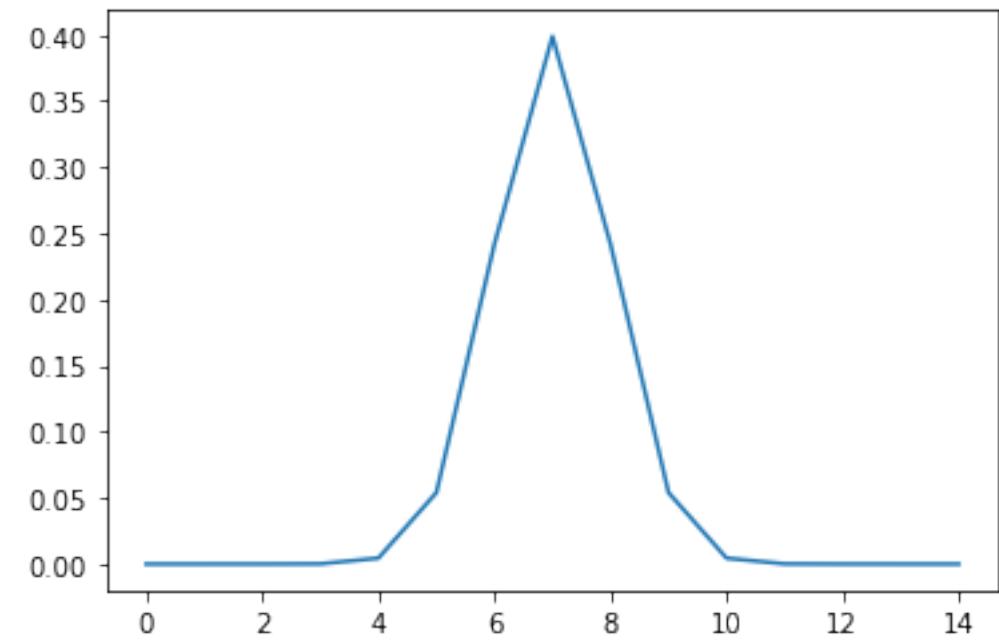
Quelle:
https://de.m.wikipedia.org/wiki/Datei:Gaussian_2d.png



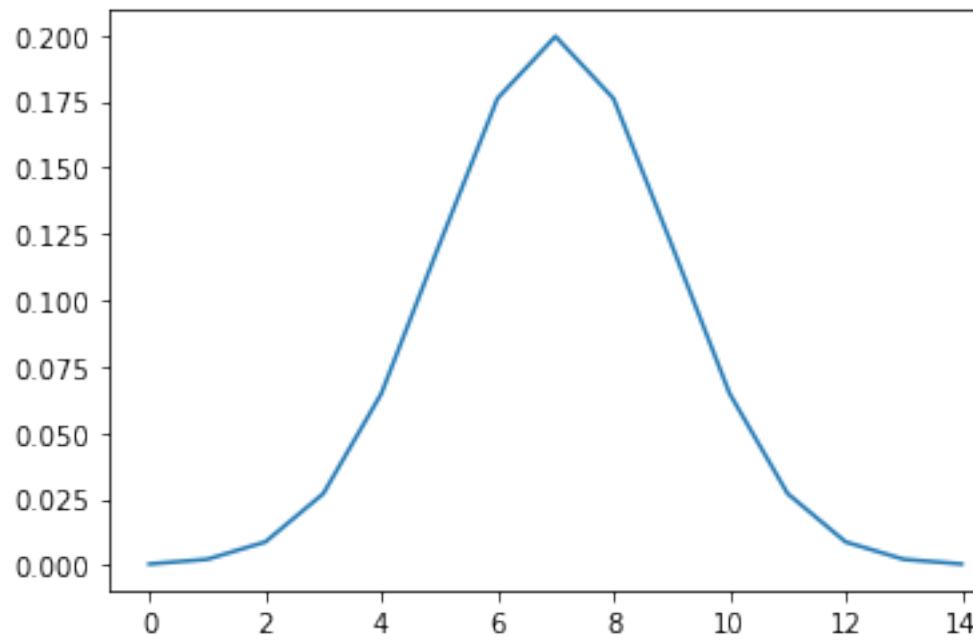
Gaussian blur filters

- Gaussian bell curves have a parameter: σ

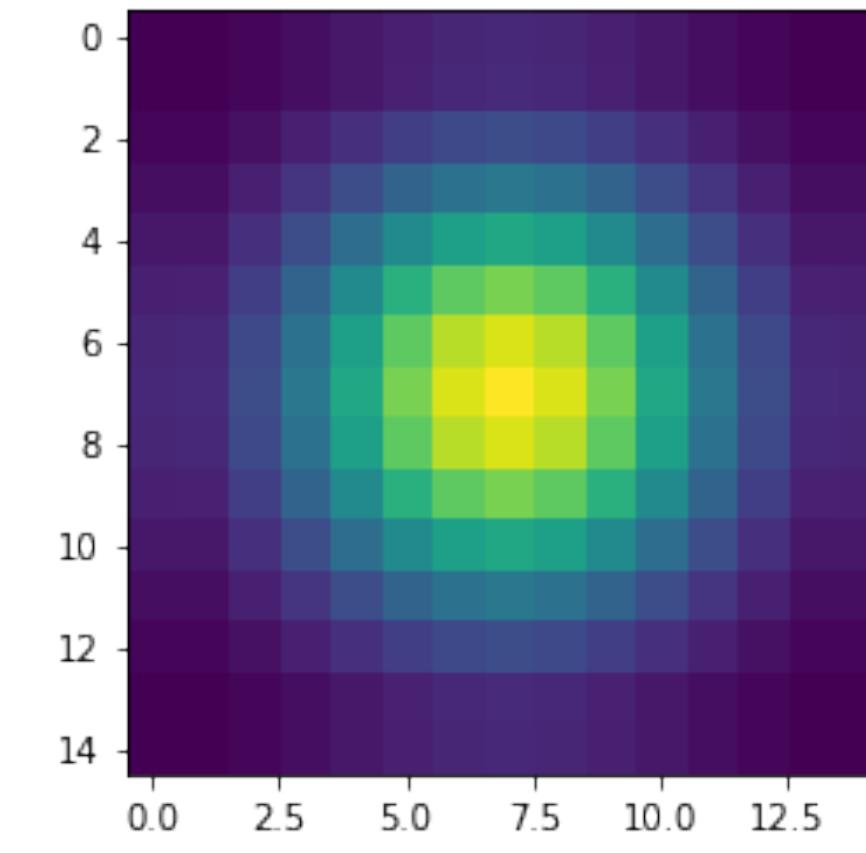
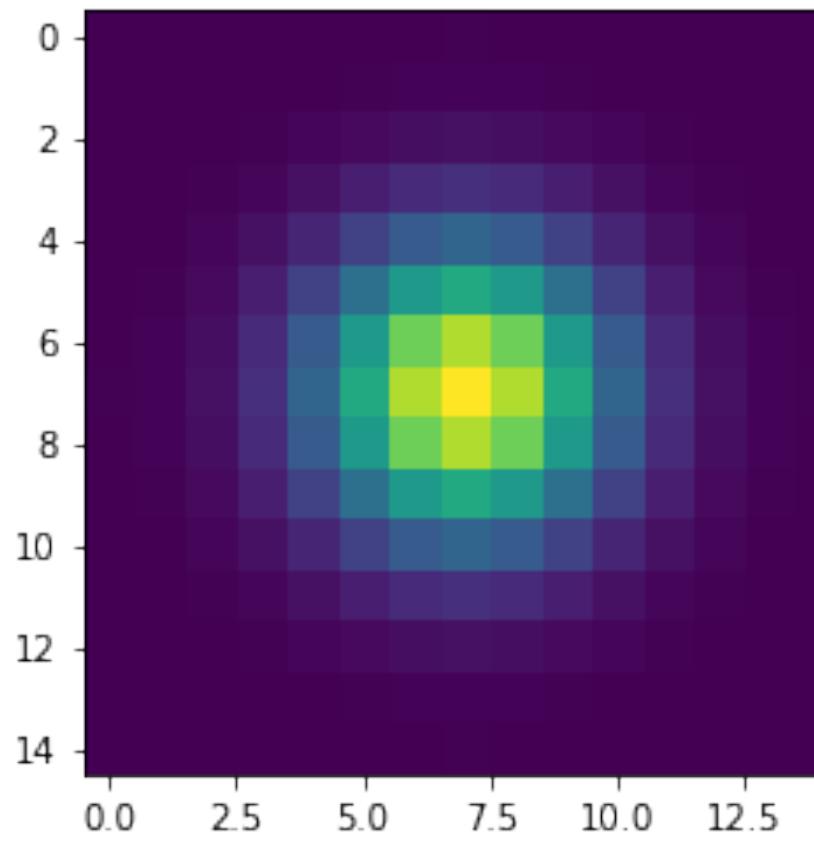
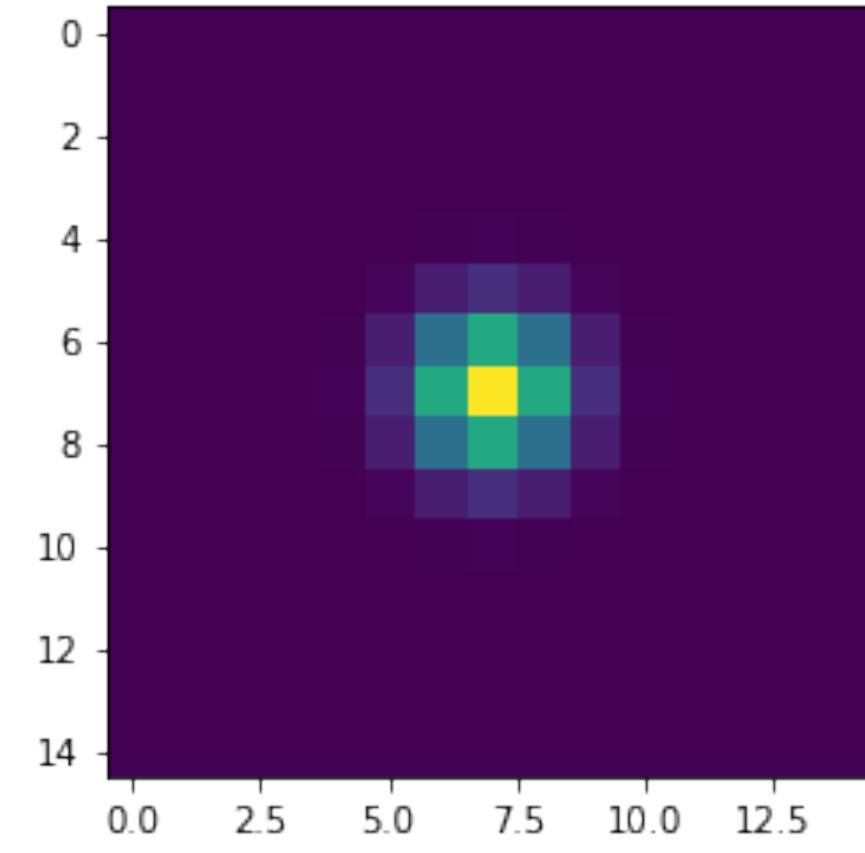
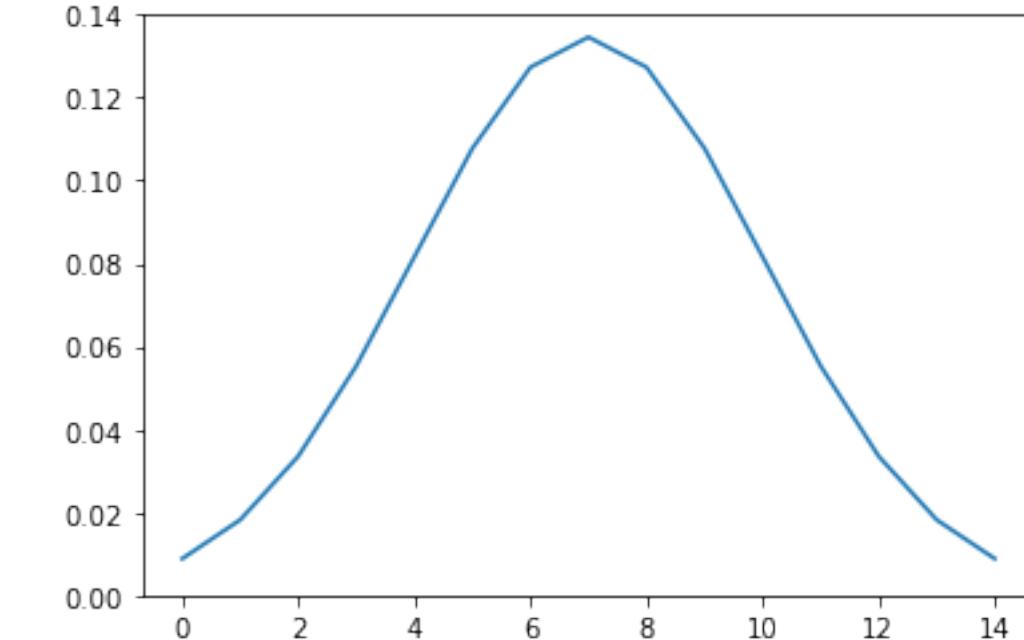
$$\sigma = 1$$



$$\sigma = 2$$



$$\sigma = 3$$



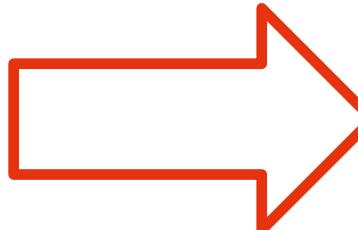
Kernel Types

Gaussian Blur 5x5

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

Takes an **average** of all neighbors with **approximated gauss bell weight**

Normalizing factor: $\frac{1}{256}$



Kernel
10x10
 $\sigma = 10$

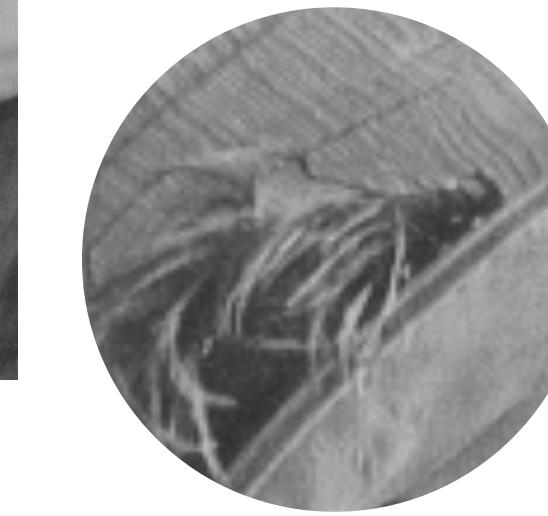
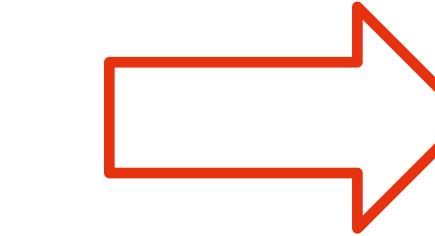
Kernel Types

Sharpen

0	-1	0
-1	5	-1
0	-1	0

Takes the difference to the neighboring pixels which leads to a sharpening effect

Normalizing factor: $\frac{1}{4 * -1 + 5} = \frac{1}{1} = 1$

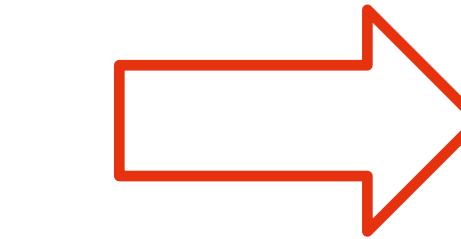


Kernel Types

Sharpen

-1	-1	-1
-1	9	-1
-1	-1	-1

If we enhance the steepness of the filter, we get a stronger effect

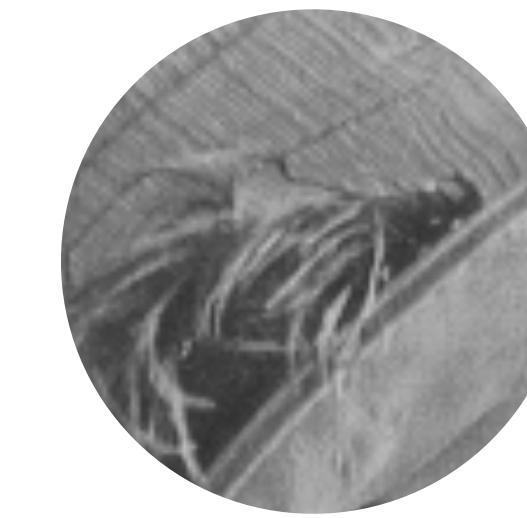
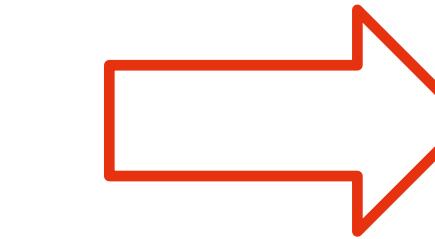


Kernel Types

Ridge detection

-1	-1	-1
-1	8	-1
-1	-1	-1

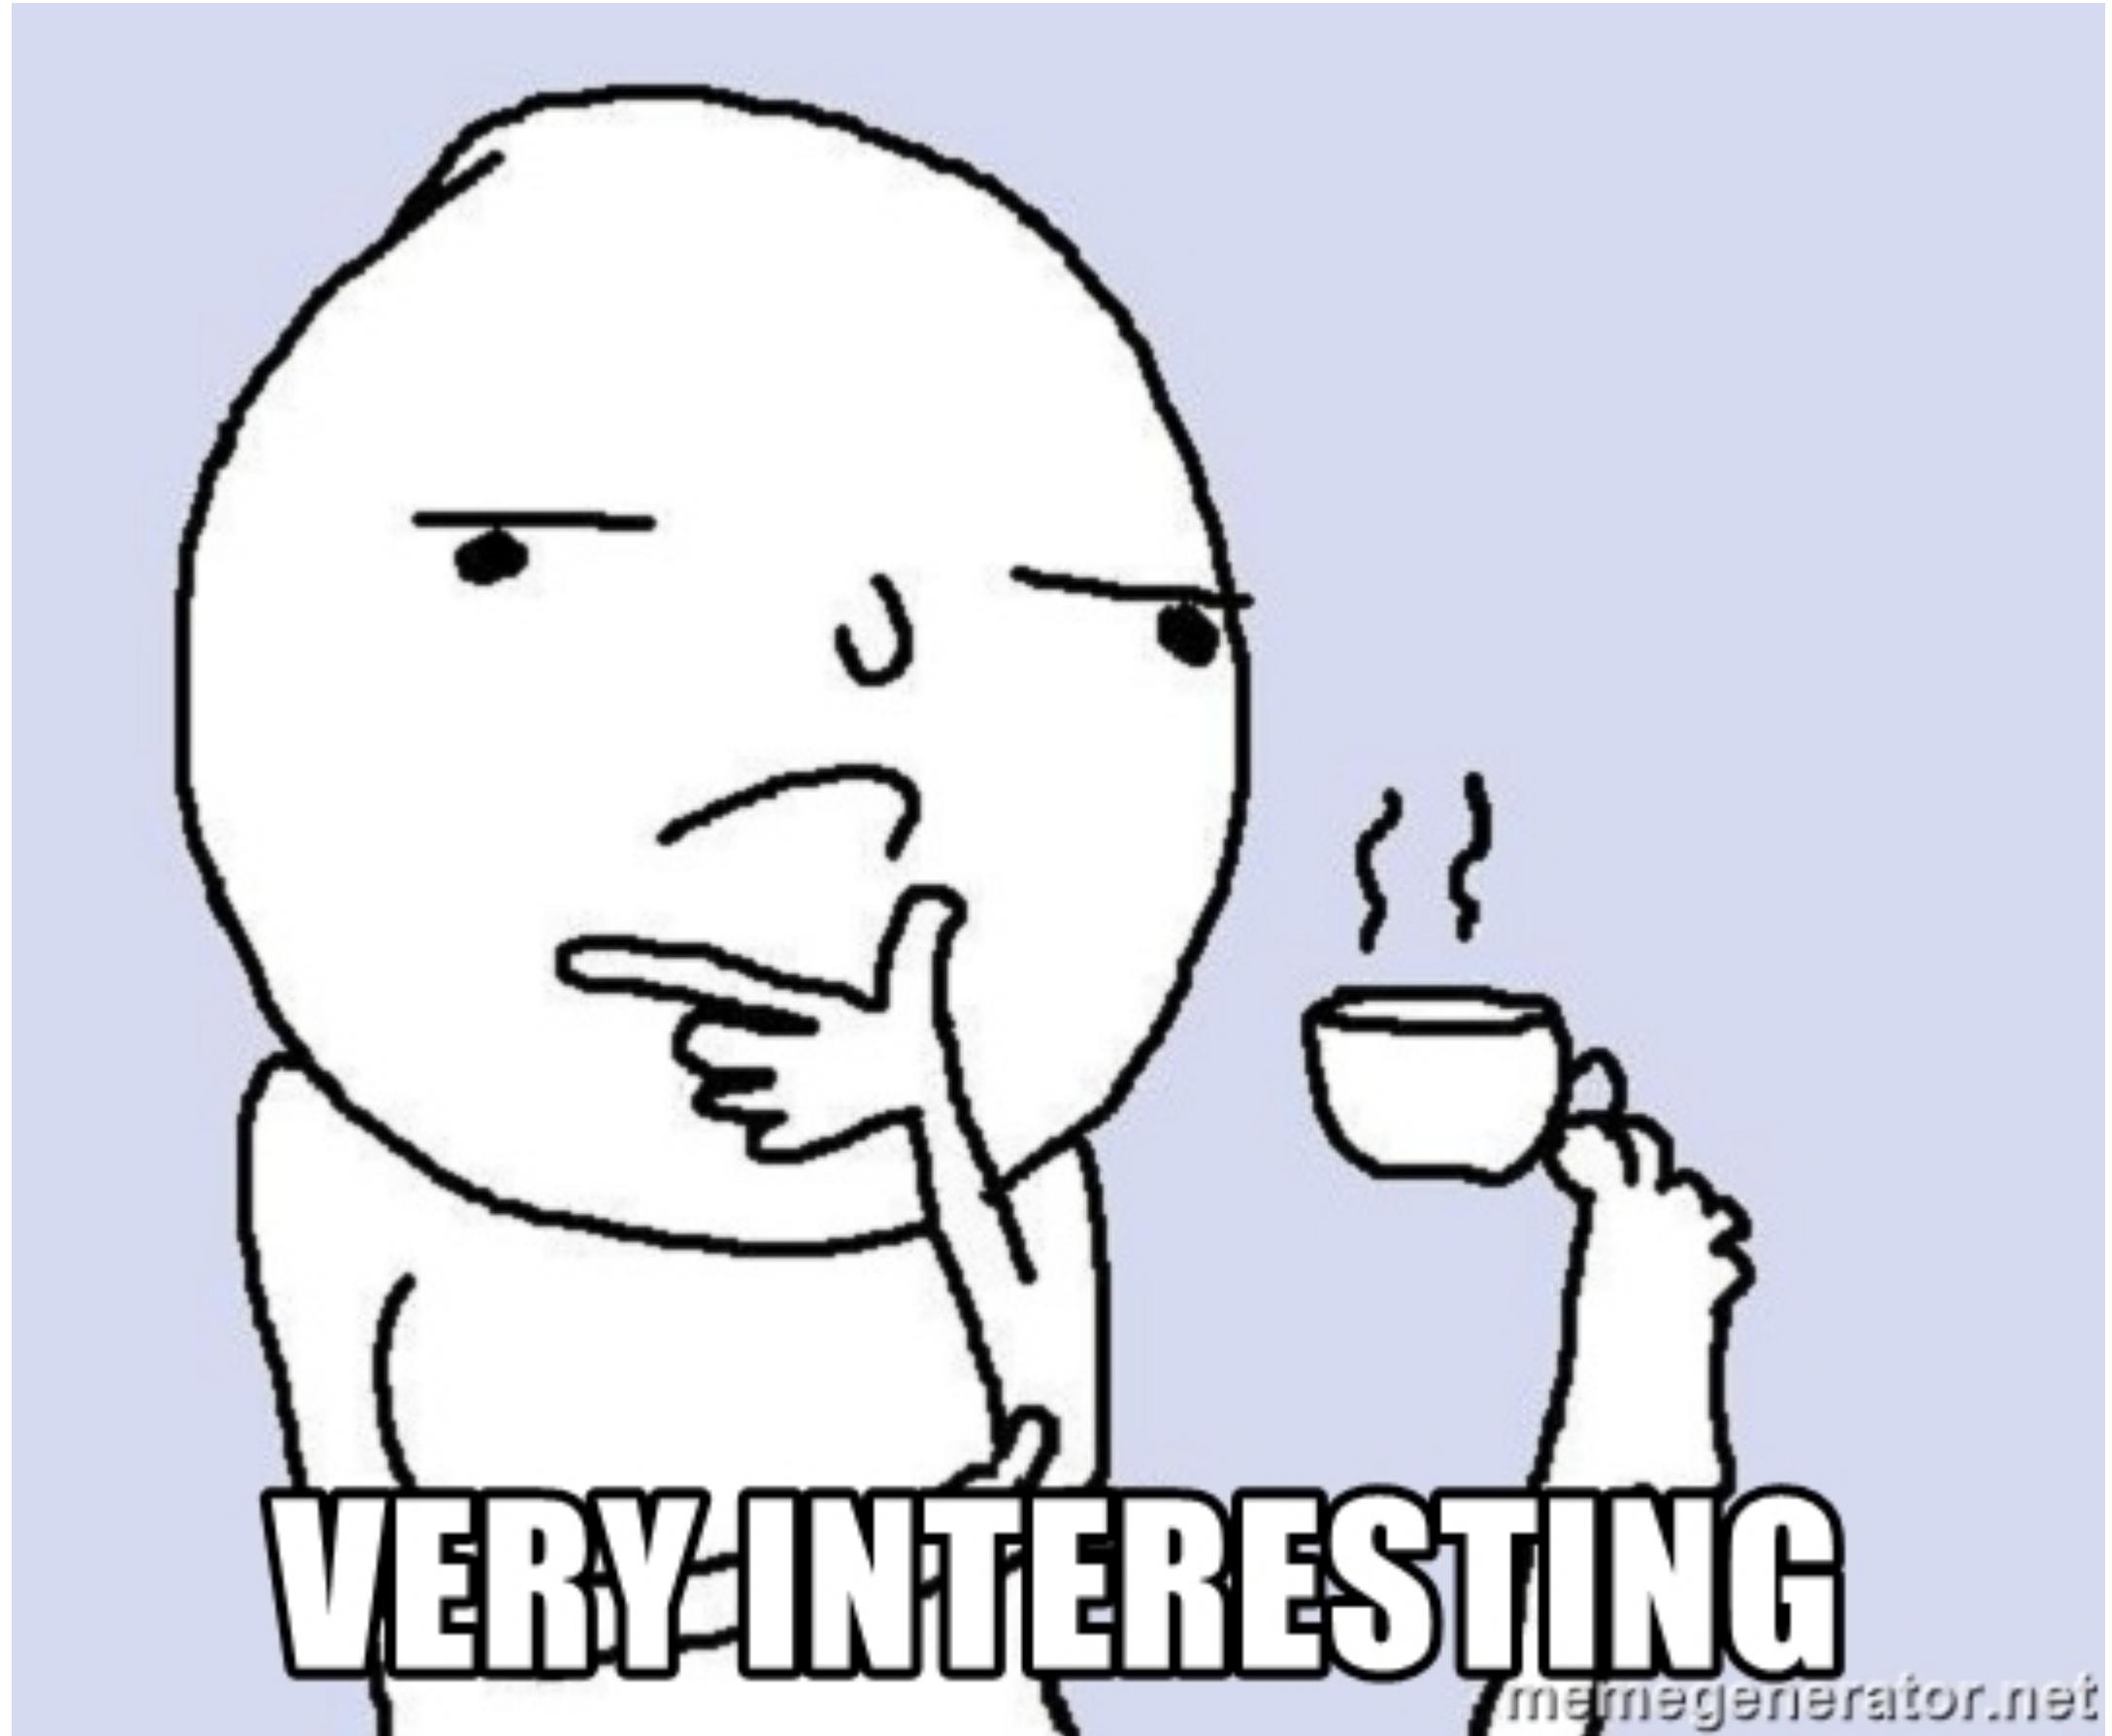
If we vary the sharpening kernel just a tiny bit, we get to a ridge detection filter.



Isn't that a bit unexpected?

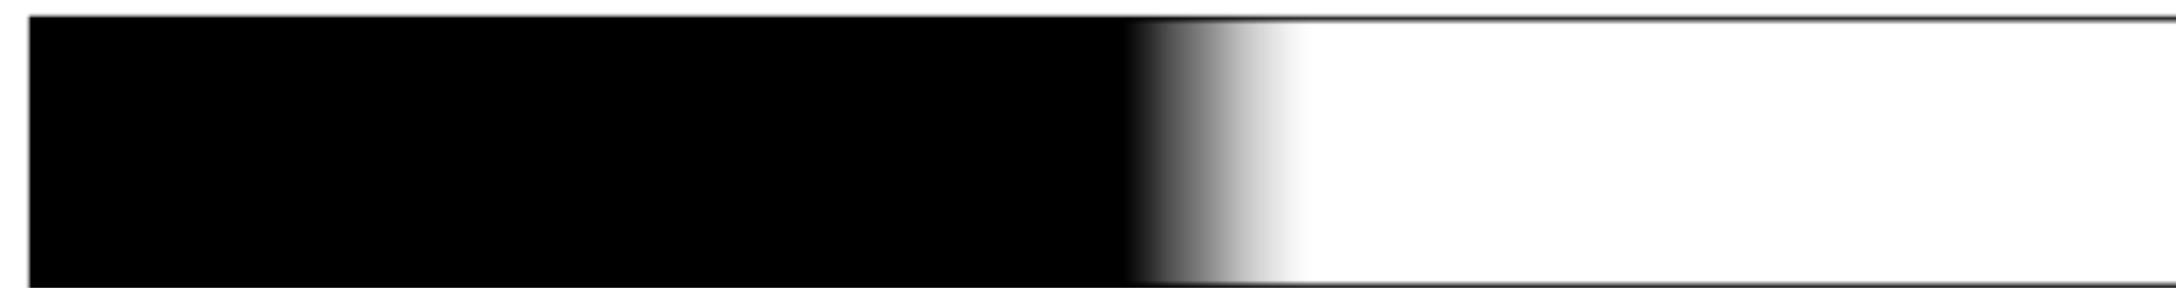
Kernel Types

Sharpen

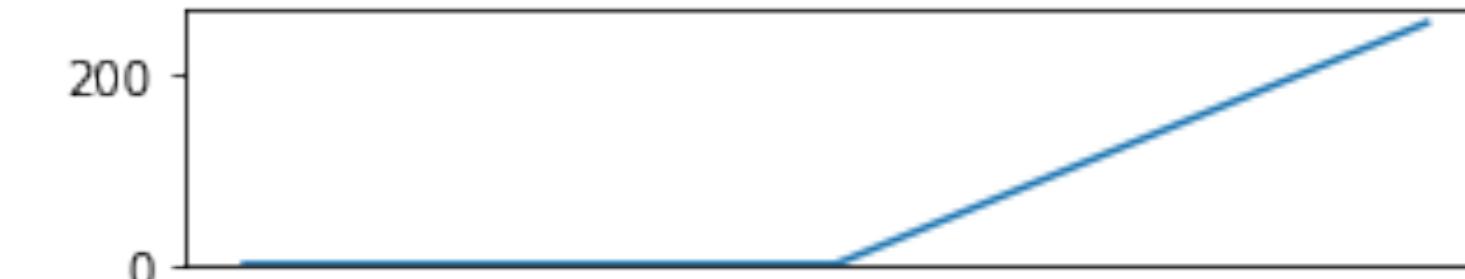
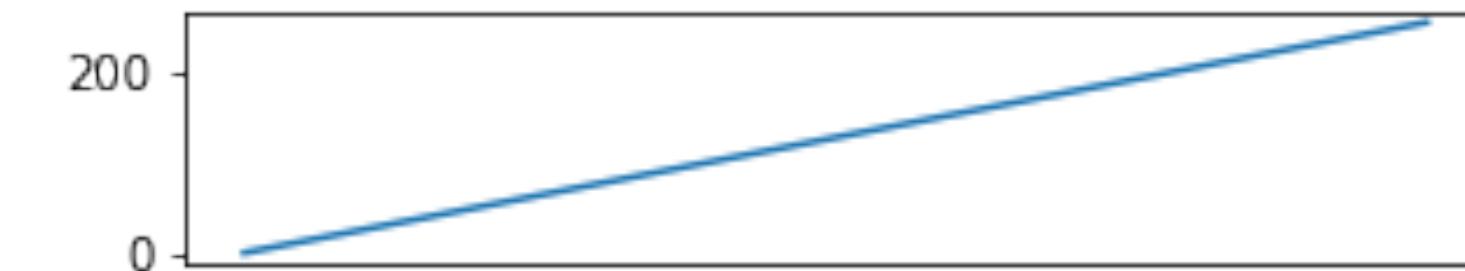


Could this be used for
detecting edges
somehow?

What is an edge?



Distribution





Kernel Types

Edge detection

- The higher the slope at a given pixel, the higher the probability of an edge
- We can get the slope by calculating the first derivative

$$\frac{ds}{dx} = s'(x) = \lim_{\Delta x \rightarrow 0} \frac{s(x + \Delta x) - s(x)}{\Delta x}$$

- As images have discrete integer values, we can simplify $\Delta x = 1$:

$$s(x + 1) - s(x)$$

As an image is defined in two dimensions $S = s(x, y)$, the gradient is defined by

$$\nabla s(x, y) = \begin{pmatrix} s(x+1, y) - s(x, y) \\ s(x, y+1) - s(x, y) \end{pmatrix}$$

→ This means, we can simply use difference operations for edge detections like

0	0	0
0	-1	1
0	0	0

H_x

0	0	0
0	-1	0
0	1	0

H_y

Kernel Types

Edge detection

0	0	0
0	-1	1
0	0	0

0	0	0
0	-1	0
0	1	0

→ These simple versions are very sensitive to noise, thus we often use a more robust one who takes the neighbors into account as well

0	-1	1
0	-1	1
0	-1	1

0	0	0
-1	-1	-1
1	1	1

Kernel Types

Edge detection → Sobel filter

→ A filter delivering good results, is the **Sobel Filter**

1	0	-1
2	0	-2
1	0	-1

H_x

1	2	1
0	0	0
-1	2	-1

H_y

Kernel Types

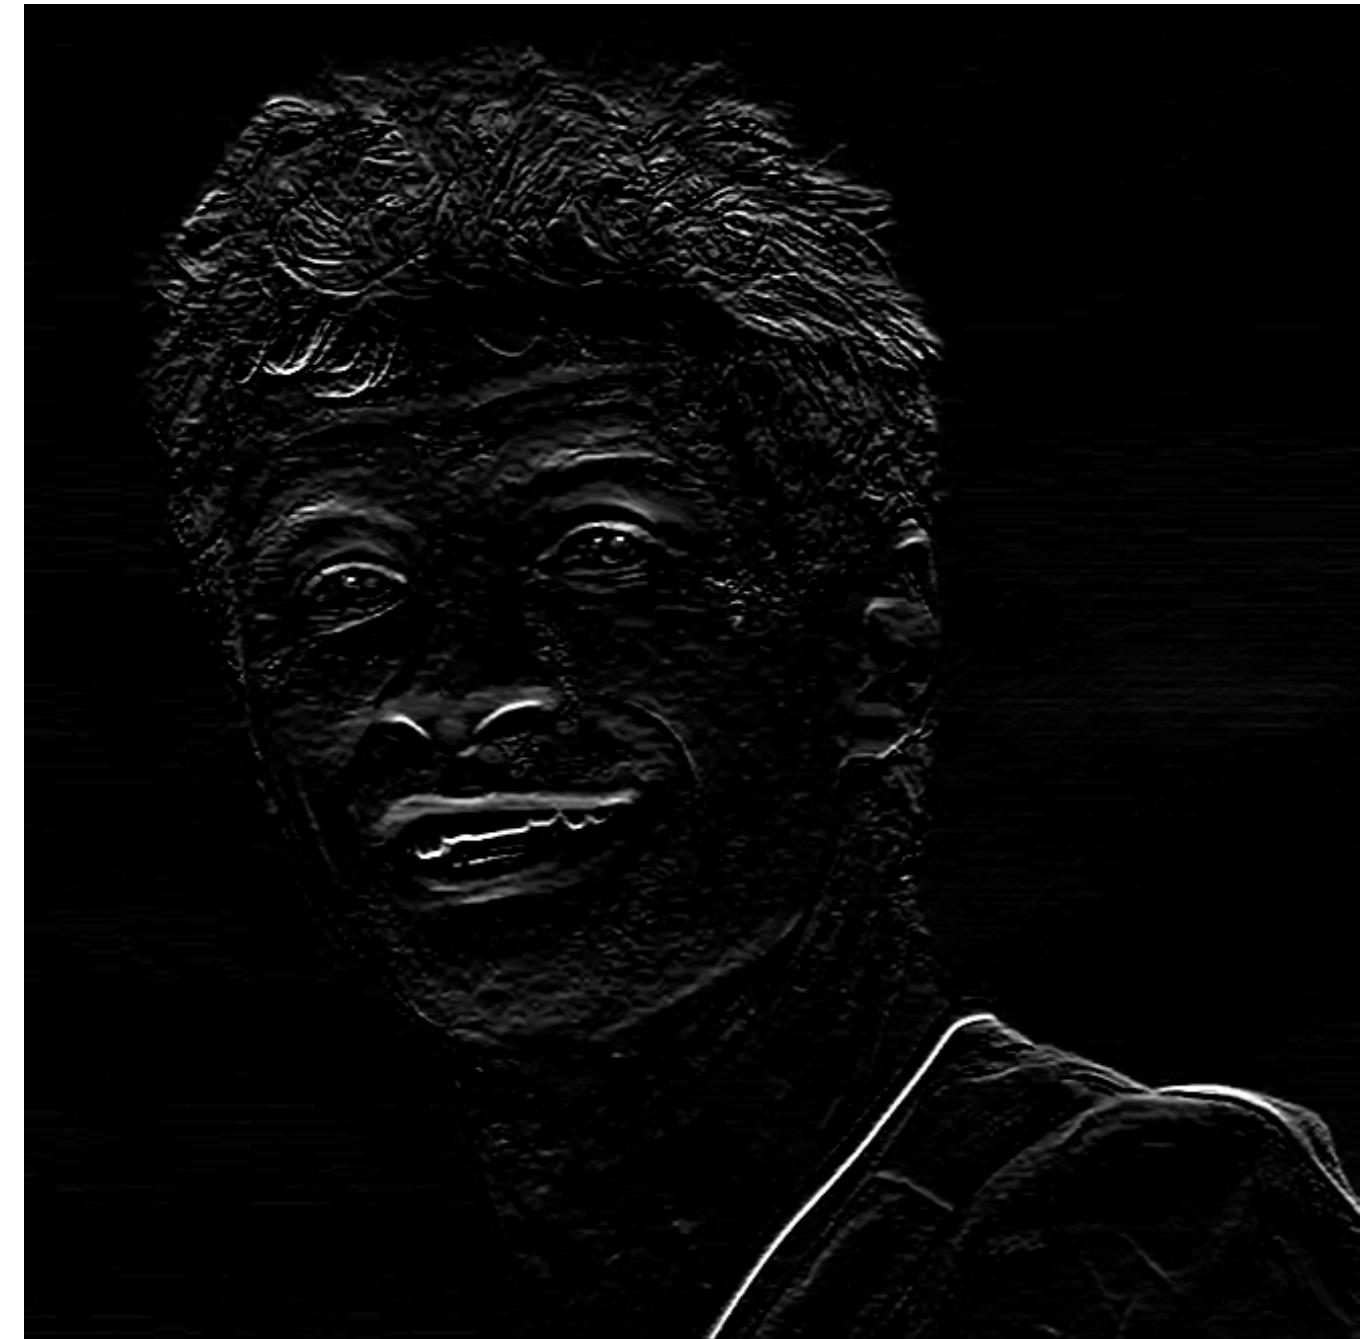
Edge detection → Sobel filter



H_x

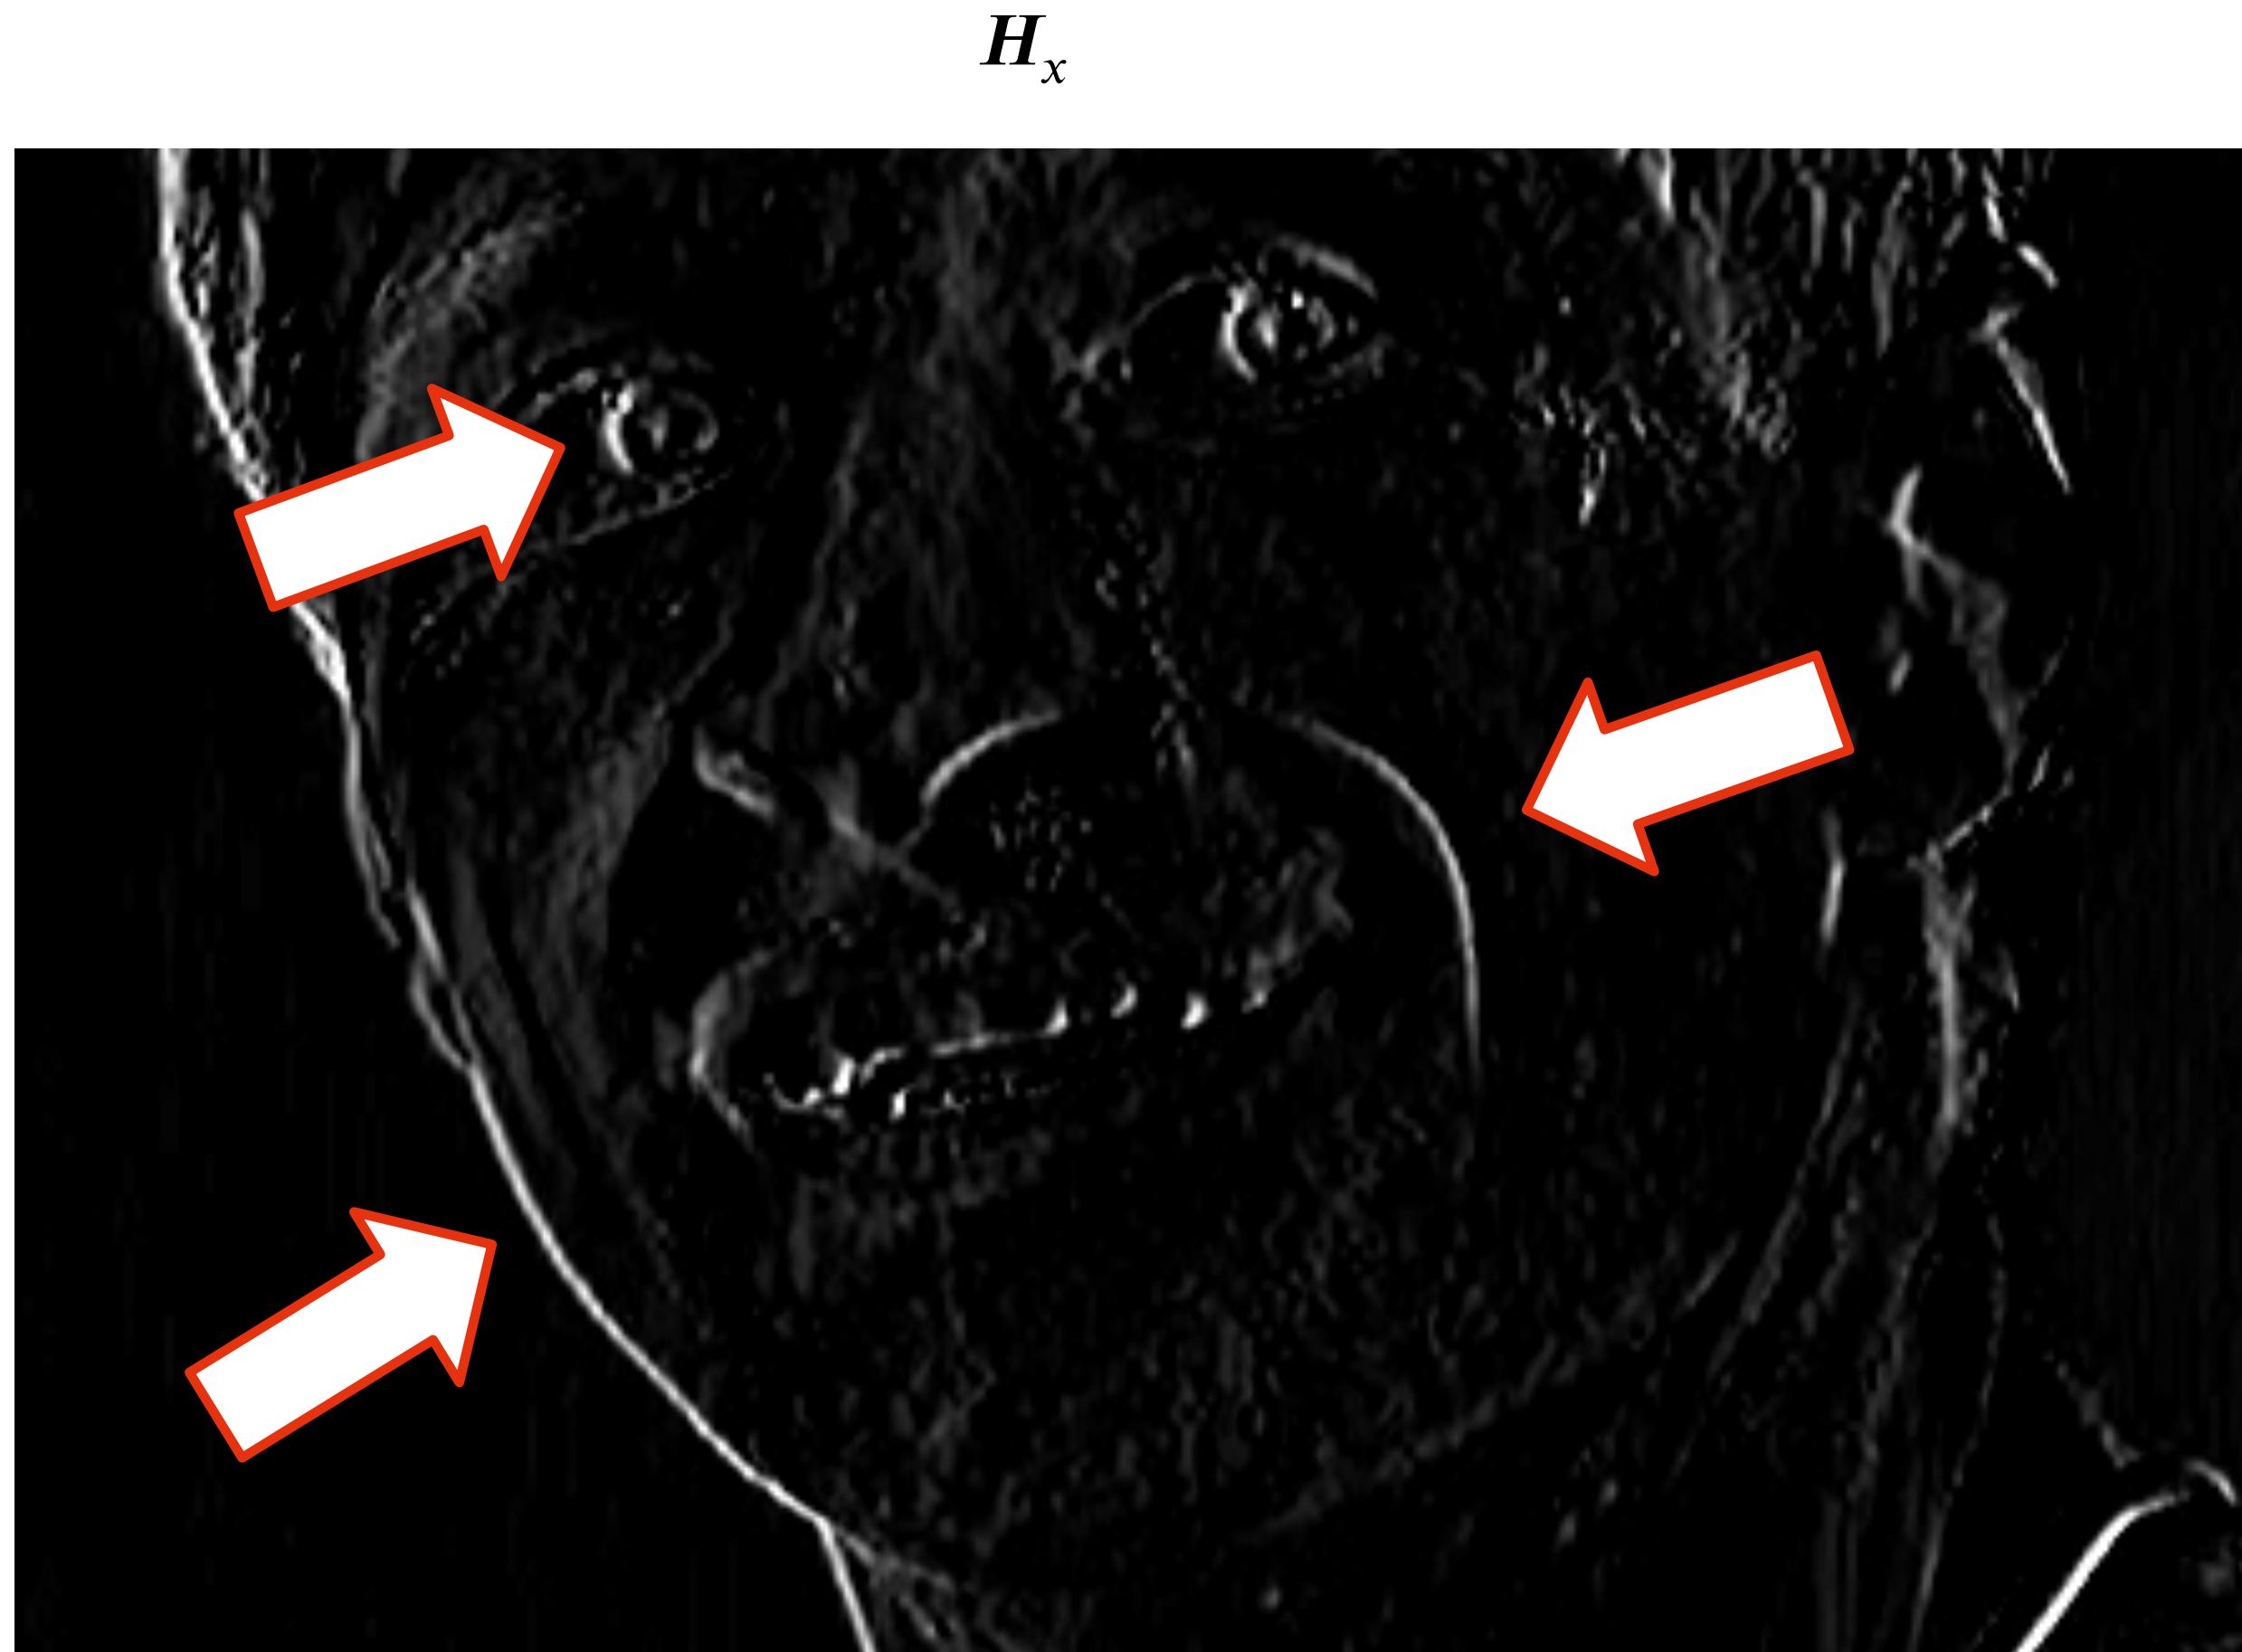


H_y



Kernel Types

Edge detection → Sobel filter





Kernel Types

Edge detection → Sobel filter

→ We can combine the filters in both directions to get a single edge detection result by

$$H = \sqrt{H_x^2 + H_y^2}$$

→ And the orientation of the edge by

$$\theta = \arctan(H_x, H_y)$$

Kernel Types

Edge detection → Sobel filter

$$H = \sqrt{H_x^2 + H_y^2}$$

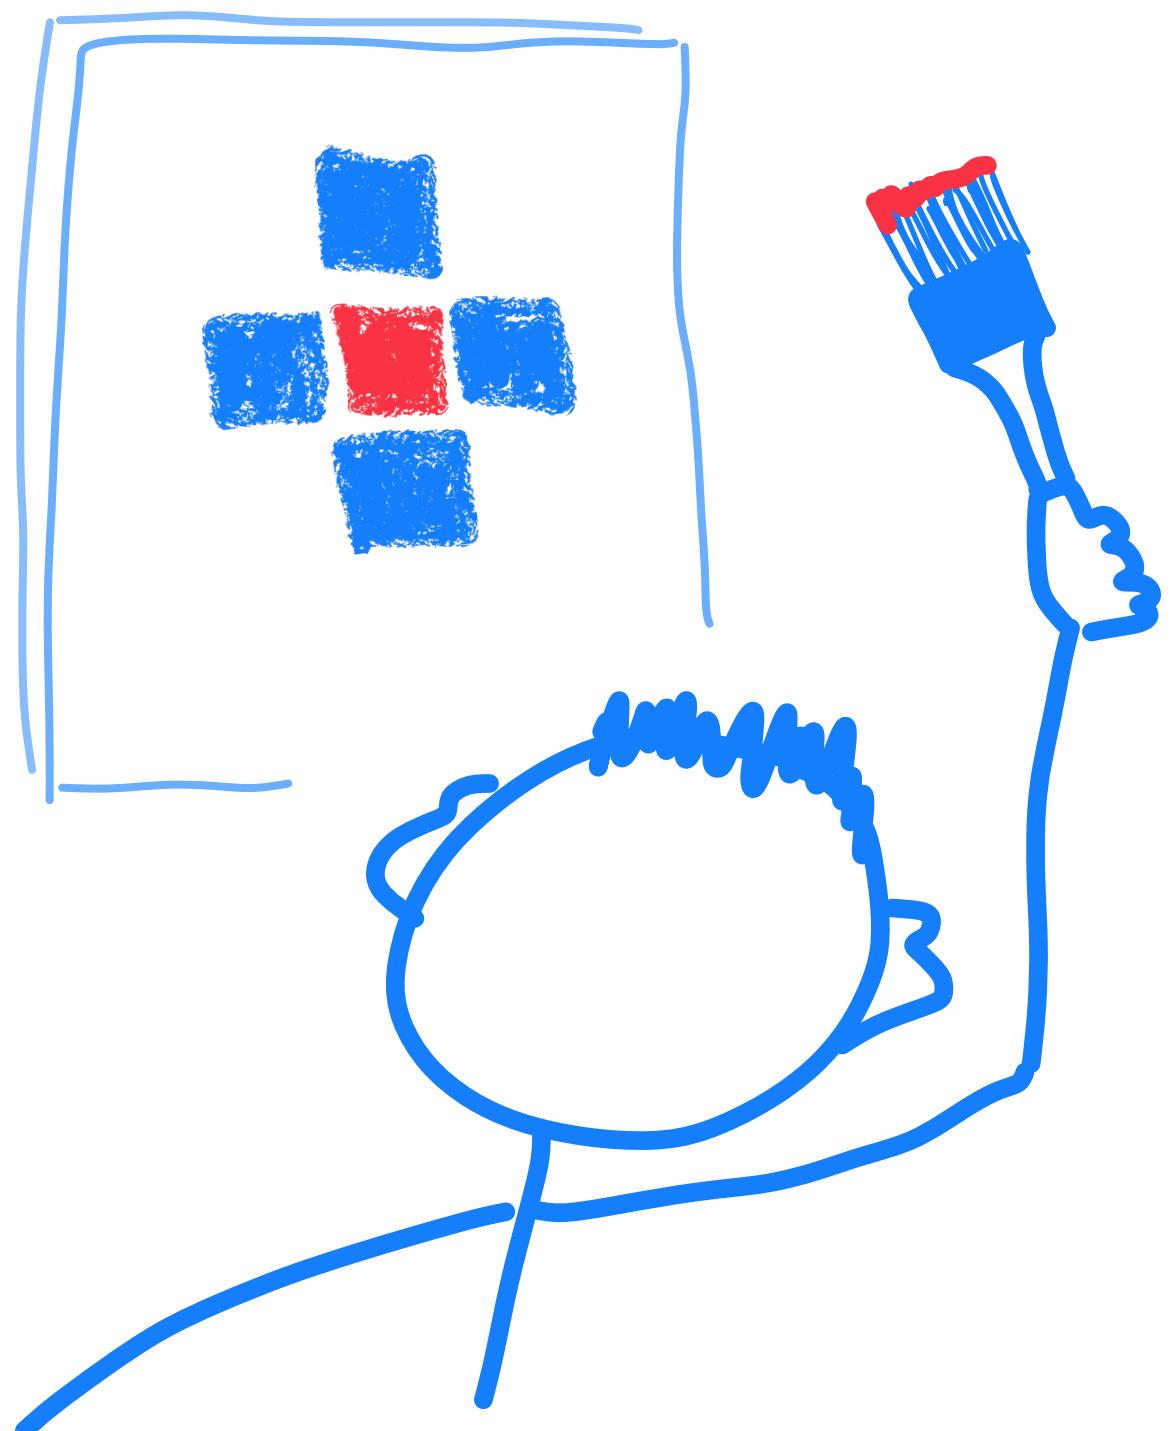
$$\theta = \arctan(H_x, H_y)$$



Multiplied by 5 for better visibility

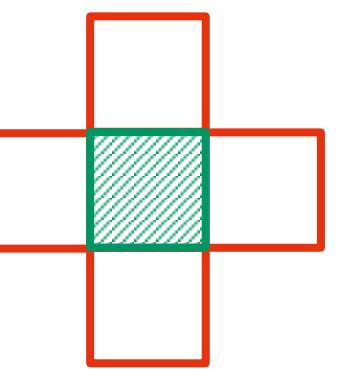


Color coded

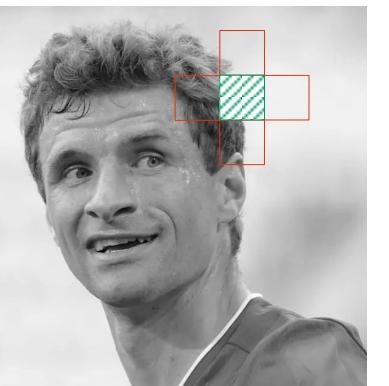


Morphology
kernels

1. Define a structuring element (SE)

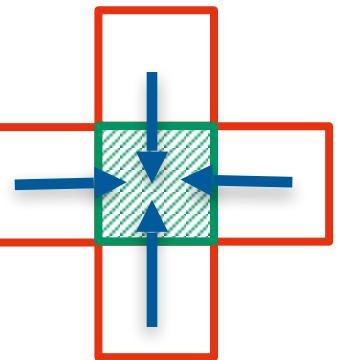


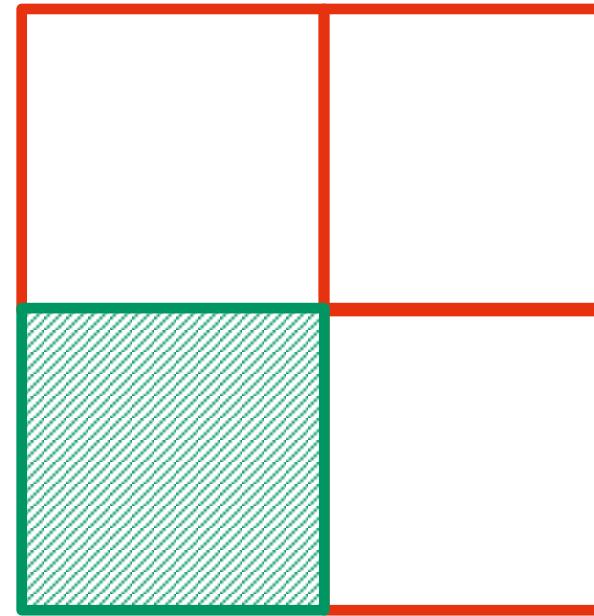
2. Define a morphological operation



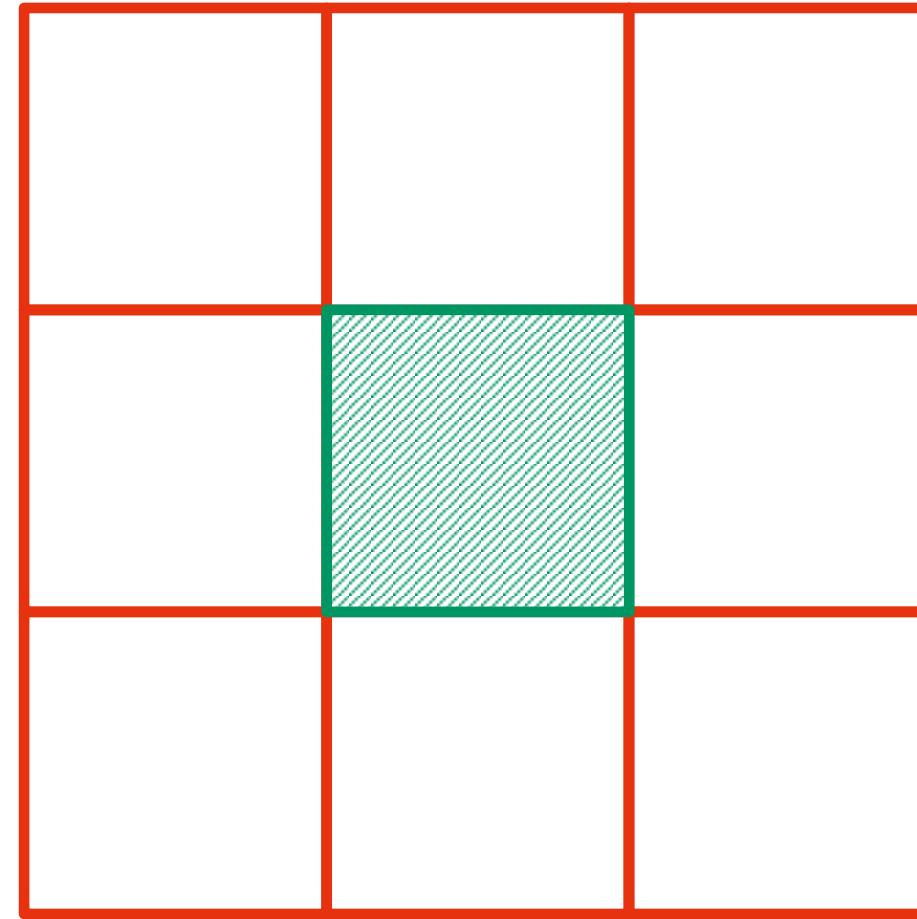
3. Apply it to the image as we did with filters

4. Set the pixel of interest to a new value based
on the other pixel values within the structuring
element and the morphological operation

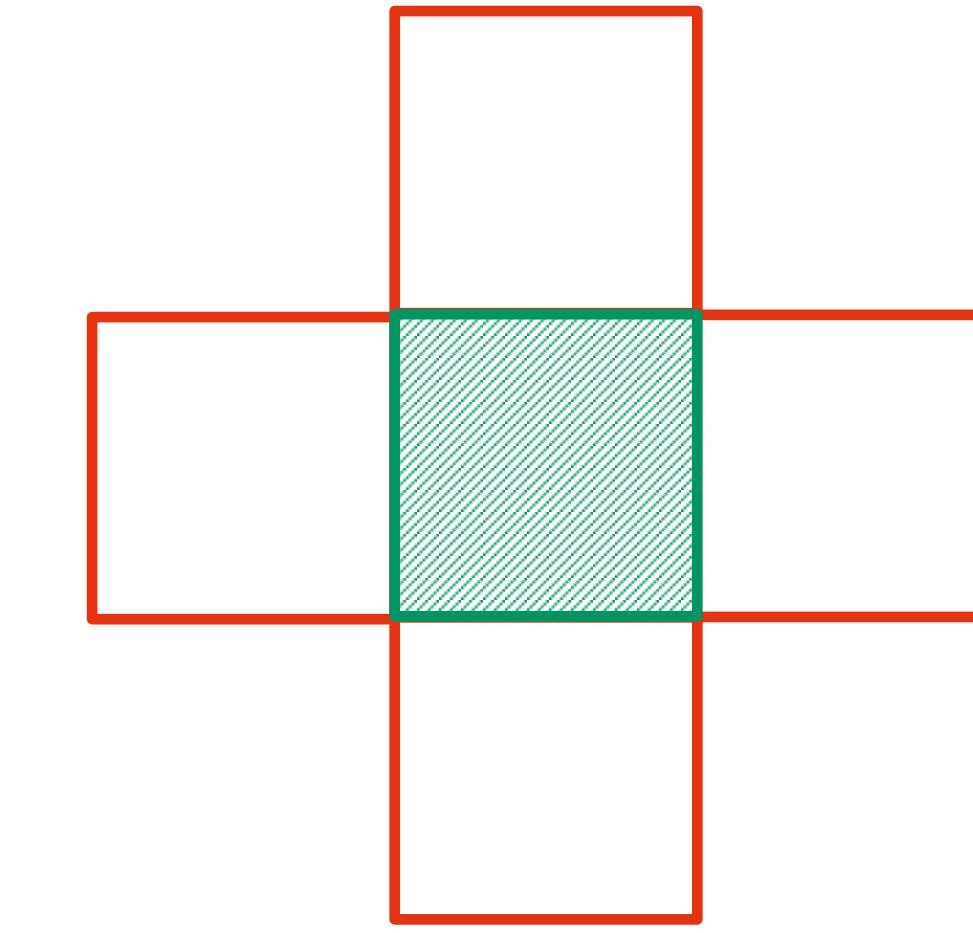




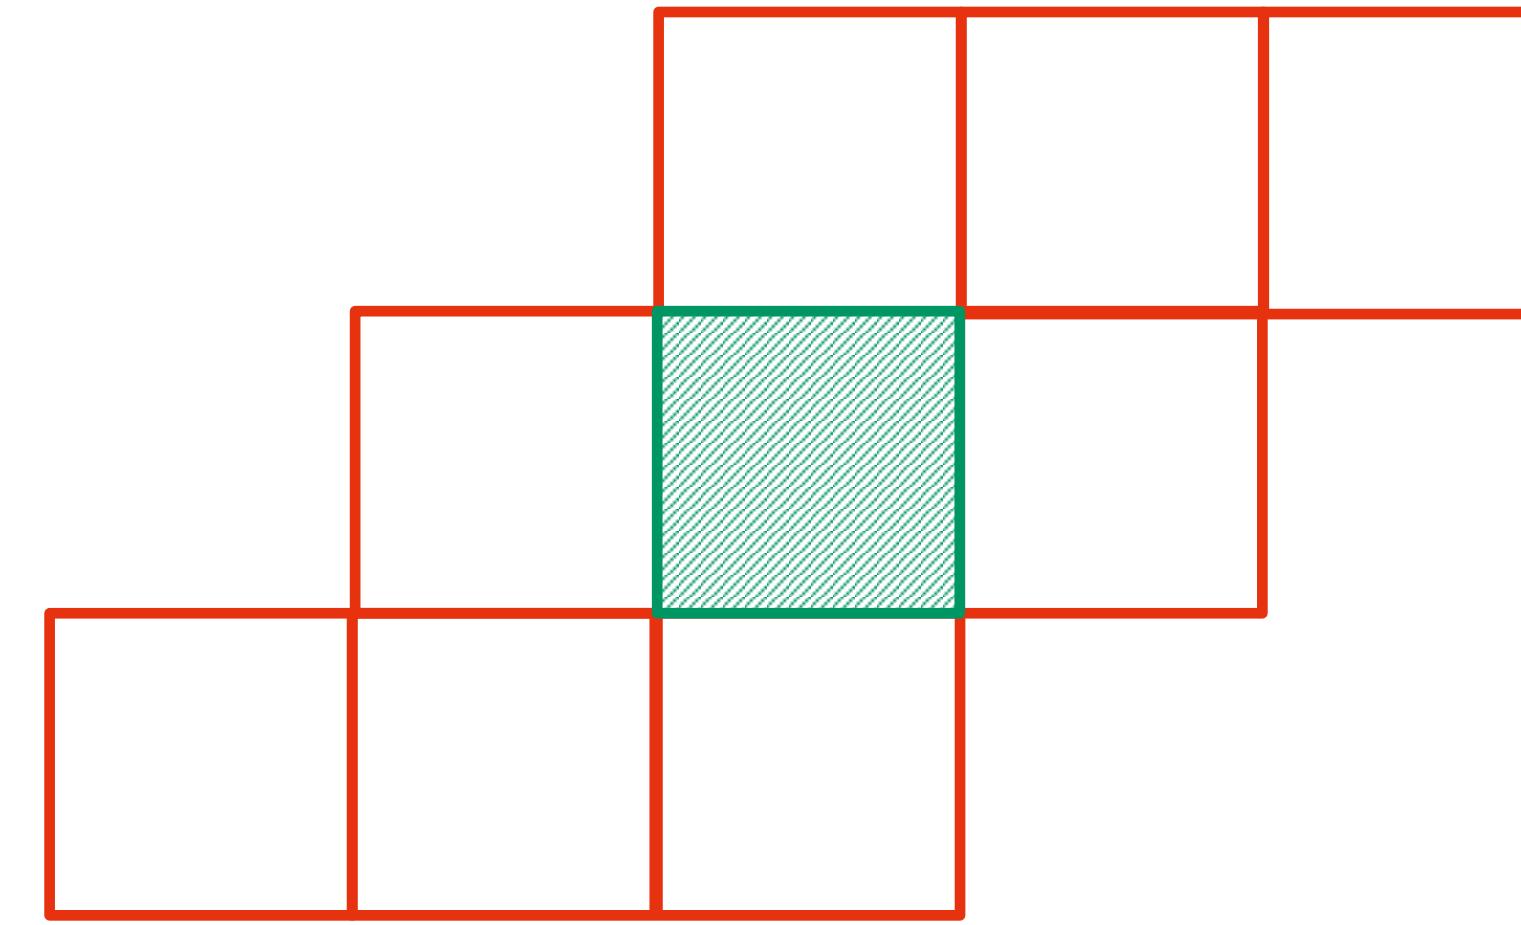
Square



8 Neighbors



Cross



diagonal



Morphology Kernels

Algorithm

- Let $S_{in} = (s_{in}(x, y))$ be a grayscale image
- Define structuring element, we use a Kernel K , which defines which points within the image are considered neighbors of a point (x, y)
- Let's call the point at position (x, y) *reference point*
- The kernel K is applied to the image S_e in the same way as the other filters so far
- The borders of the image are handled as we learned before
- For each position, all intensity values of all neighbors defined by the structuring element are ordered by their value and build a rank r
- The reference point is next replaced with a value $v \in r$
- Depending on how to chose v from r , we can define different morphological operations

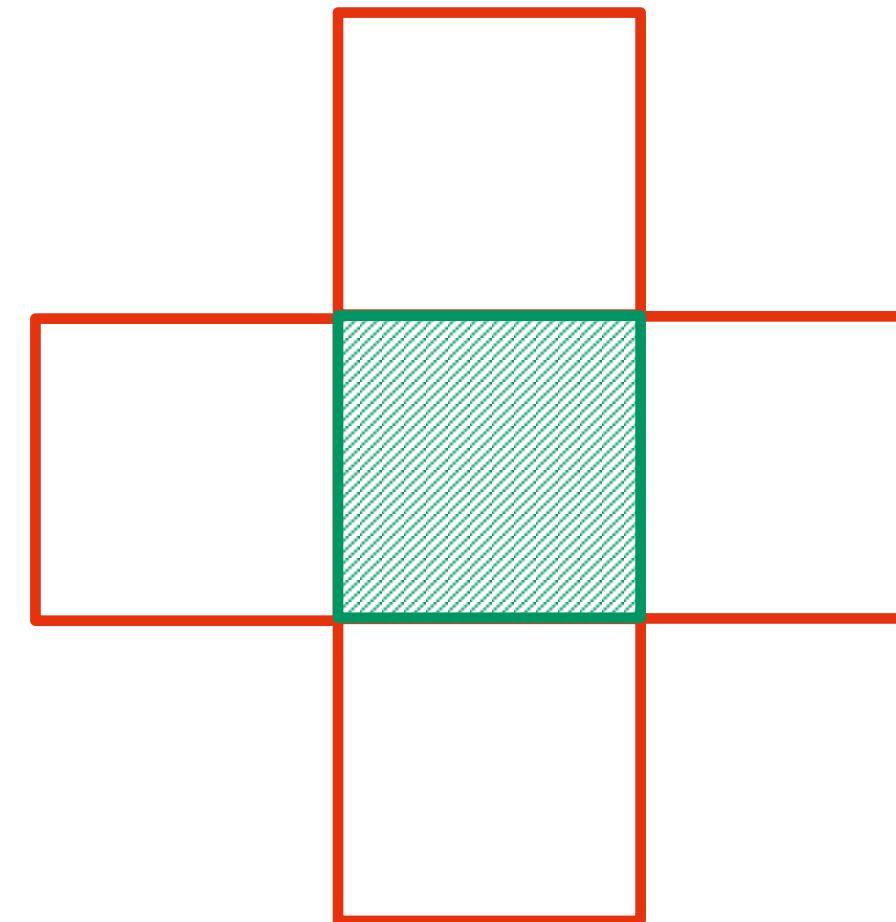
Morphology Kernels

Algorithm, Example:

Example Image

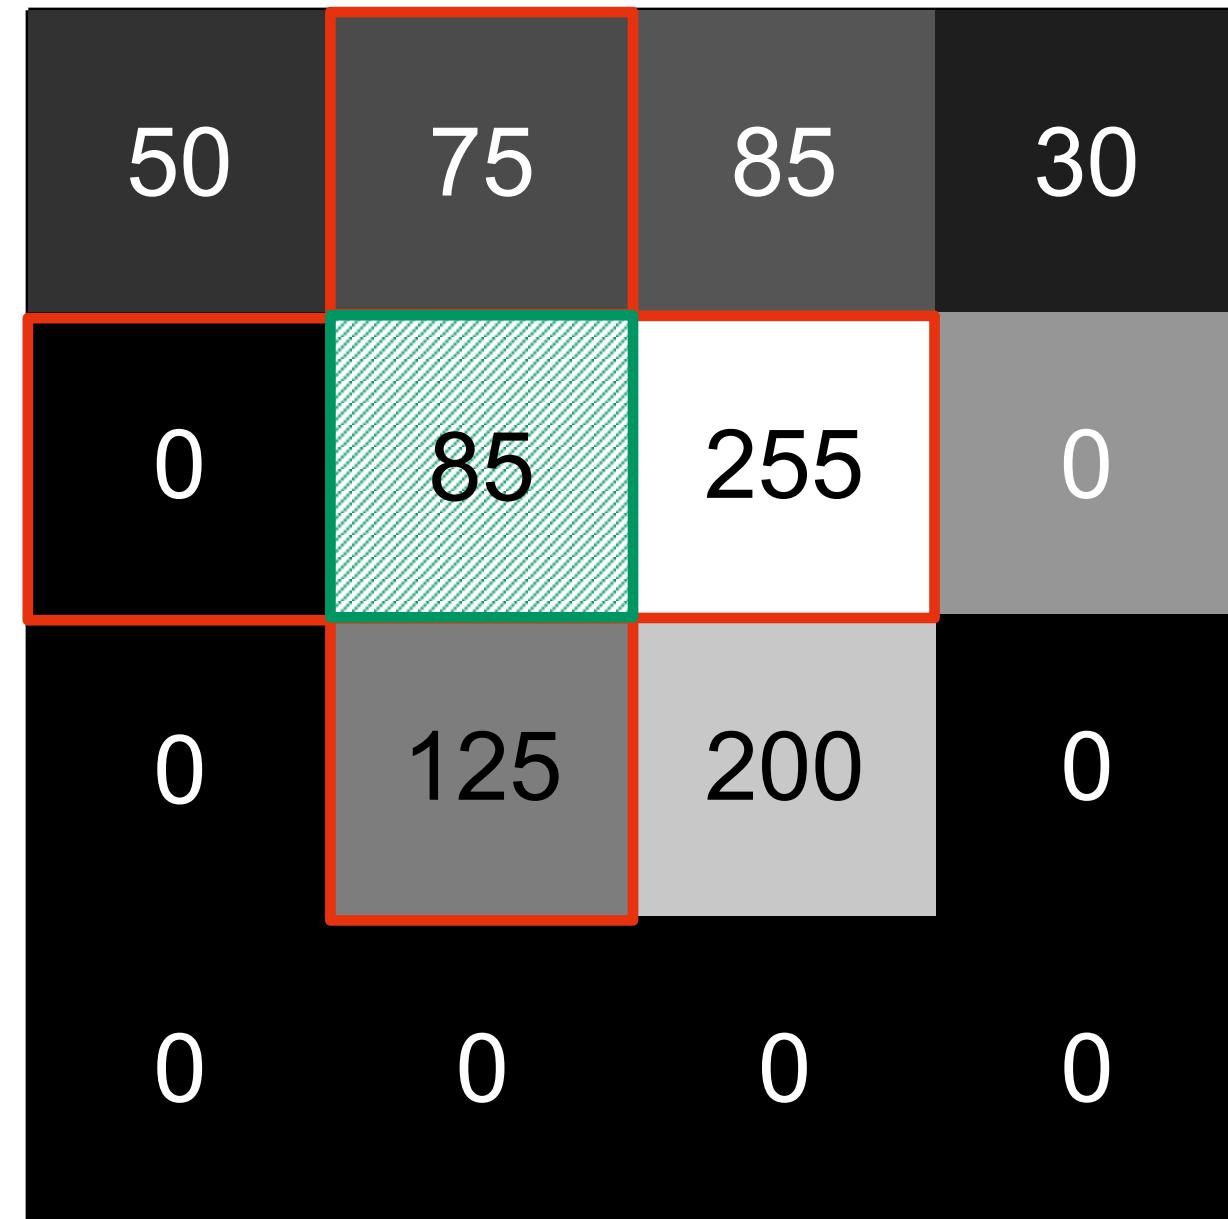
50	75	85	30
0	85	255	0
0	125	200	0
0	0	0	0

Structuring element



Morphology Kernels

Algorithm, Example:

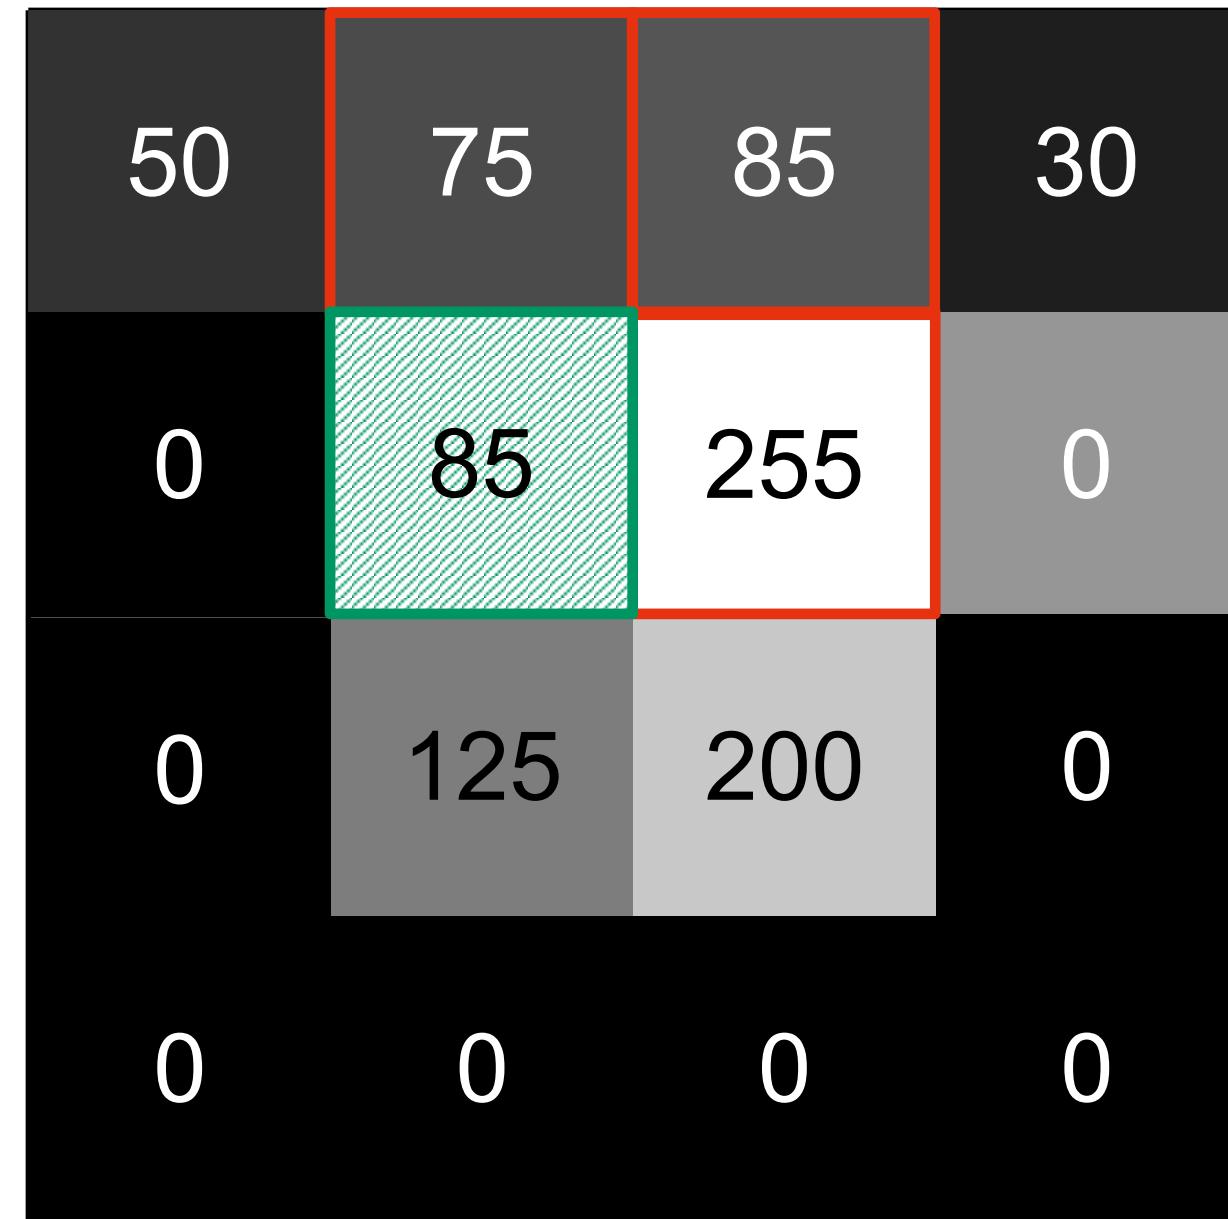


Example at position (1,1):

$$r = \{0, 75, 85, 125, 255\}$$

Morphology Kernels

Algorithm, Example:

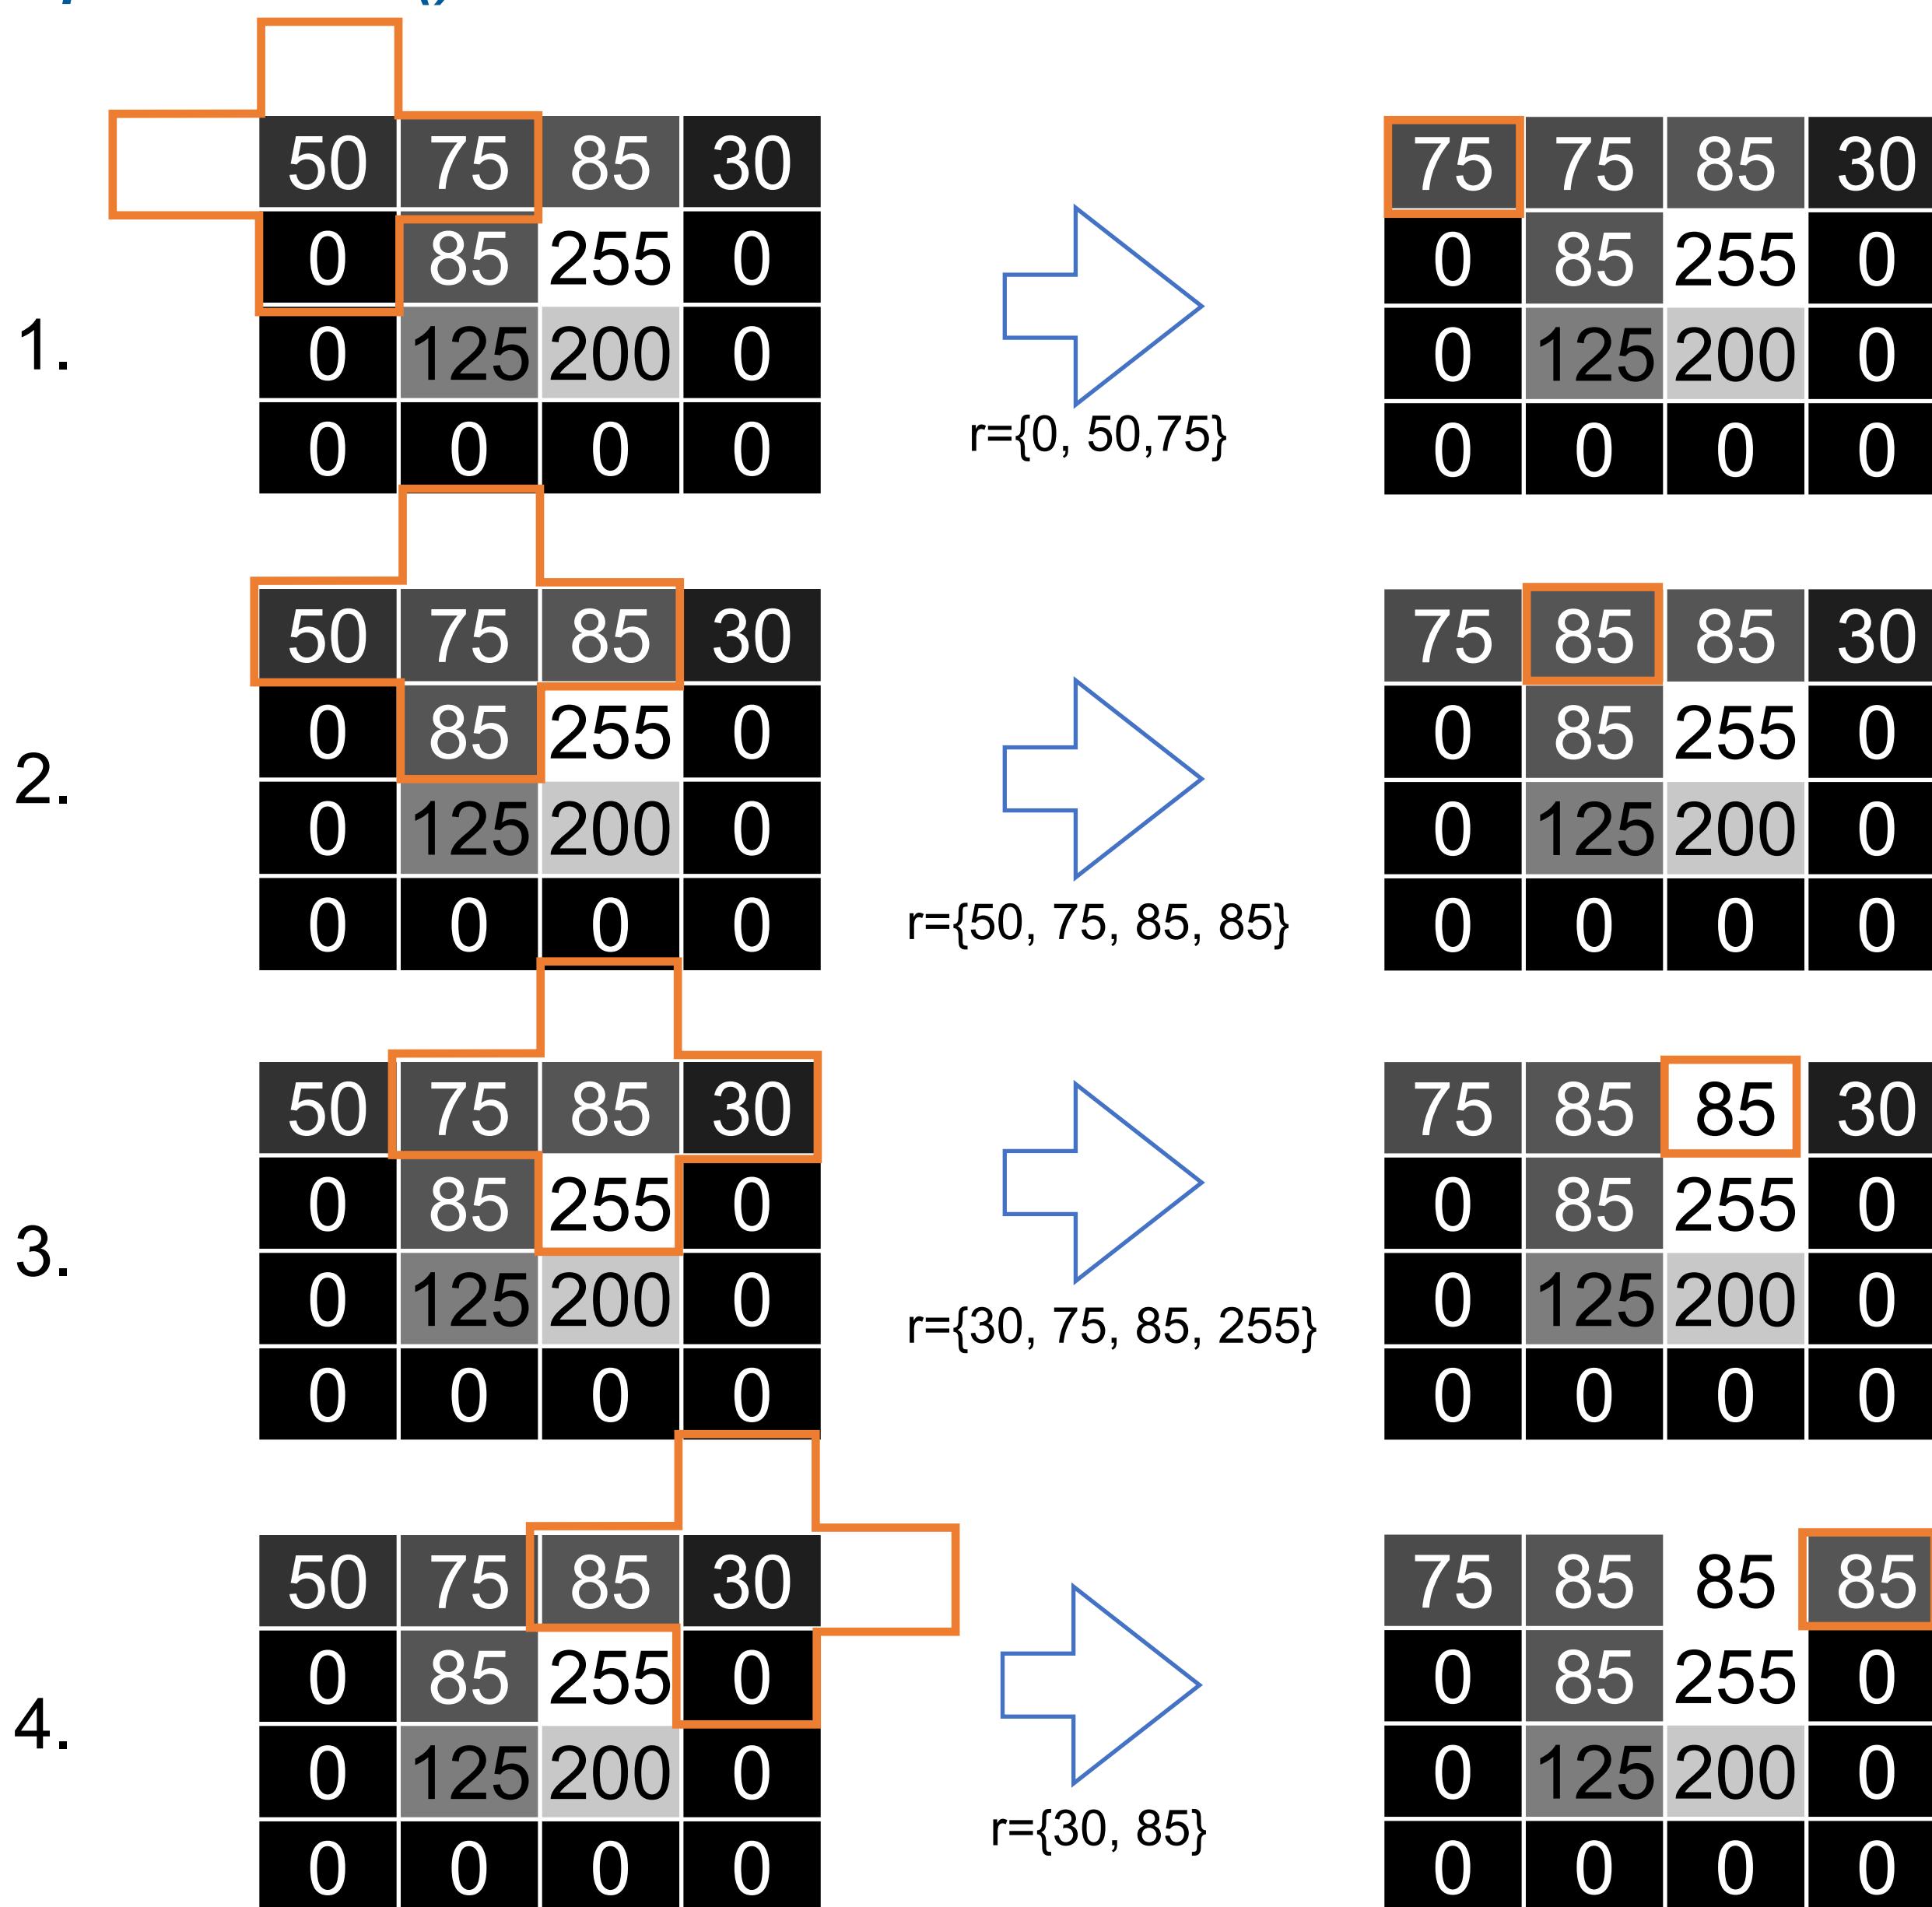


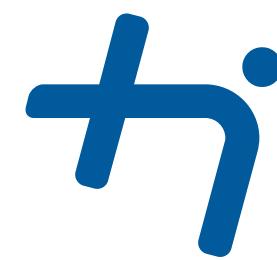
Example at position (1,1):

$$r = \{75, 85, 85, 255\}$$



Example: max()



Example: *max()*

5.		
6.		
7.		

Result after 16 steps:

75	85	85	85
85	255	255	255
125	200	255	200
0	125	200	0

Input image:

50	75	85	30
0	85	255	0
0	125	200	0
0	0	0	0

Comparing various morphological operations

Input image:

50	75	85	30
0	85	255	0
0	125	200	0
0	0	0	0

max() operation

75	85	85	85
85	255	255	255
125	200	255	200
0	125	200	0

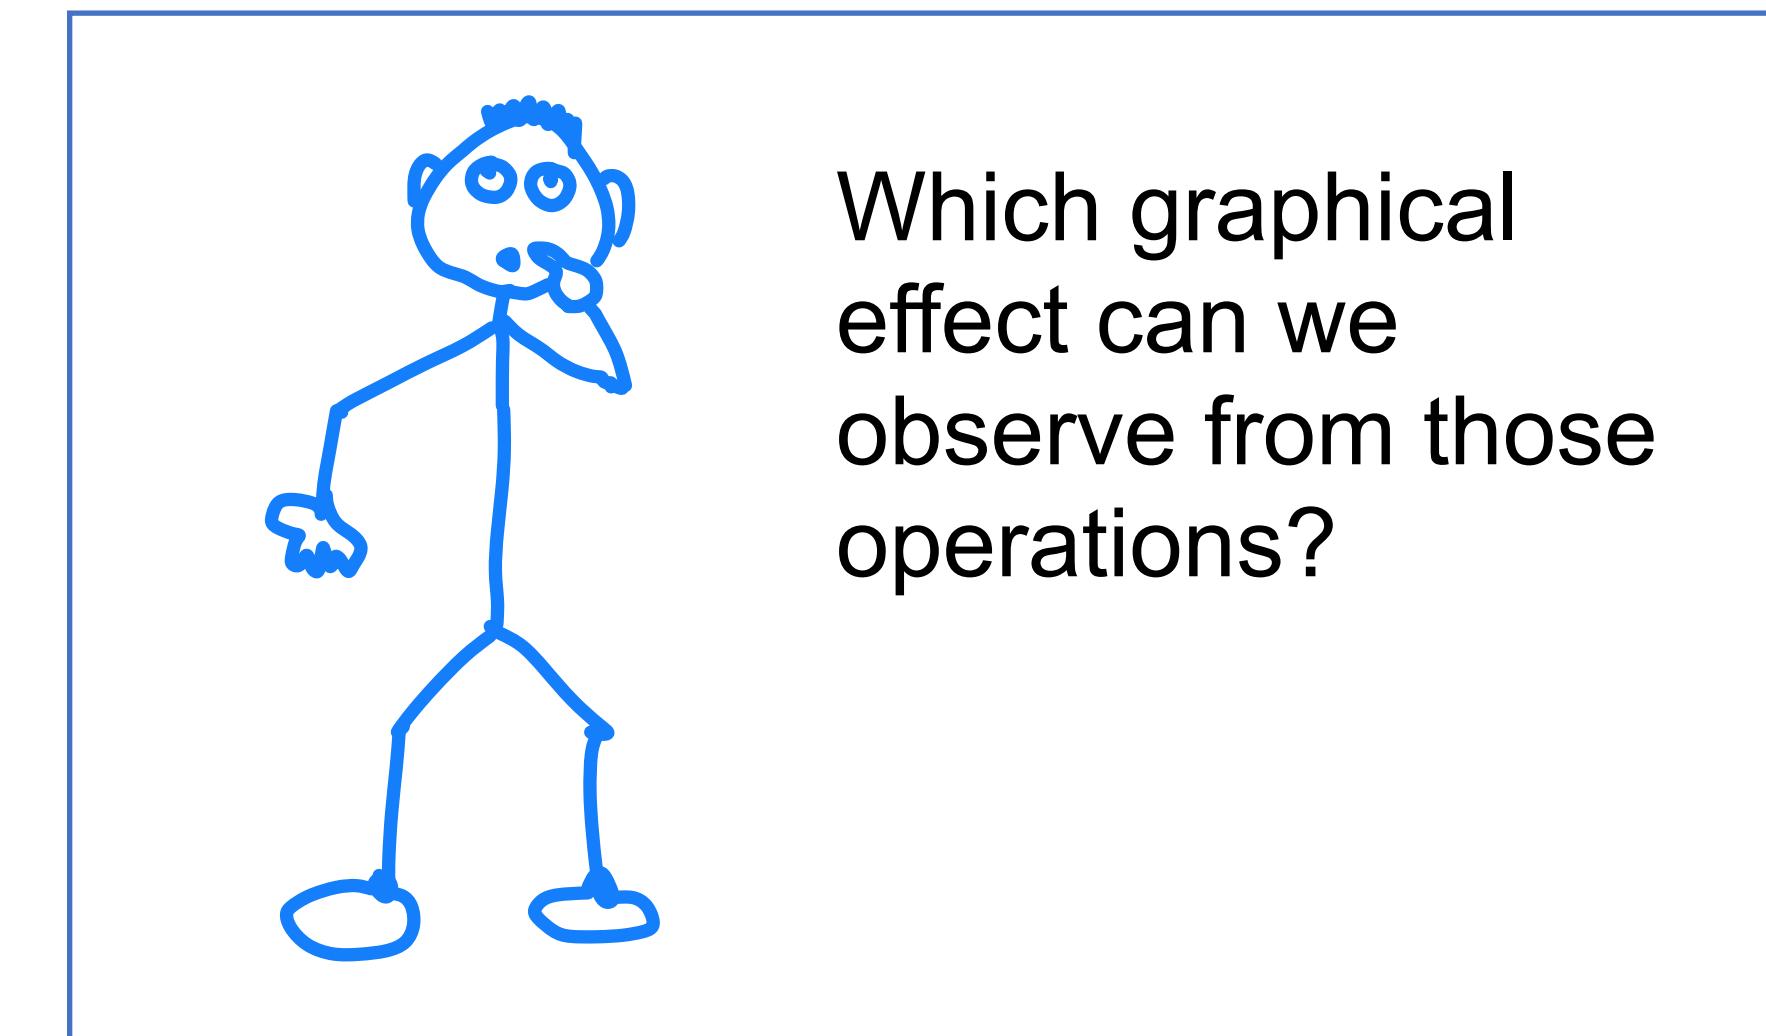
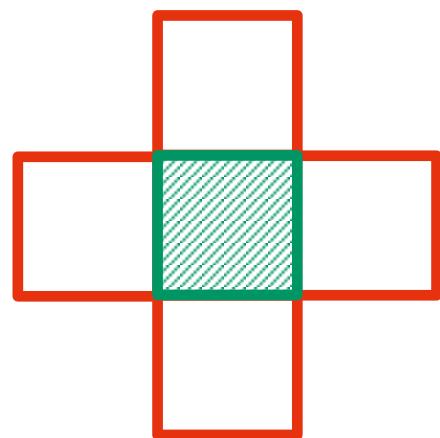
min() operation

0	50	75	30
0	0	0	0
0	0	0	0
0	0	0	0

median() operation

50	80	80	85
0	85	85	15
0	85	125	0
0	0	0	0

Structural element:

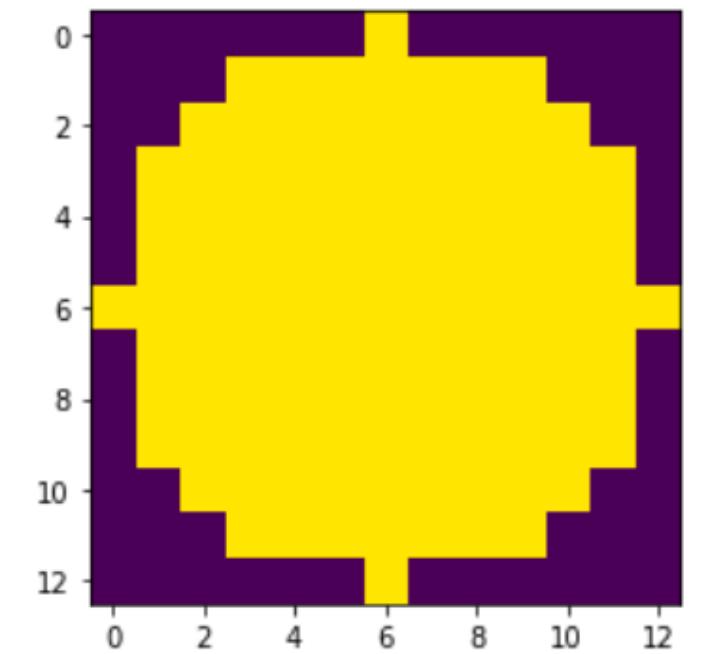


Comparing various morphological operations

Input image:



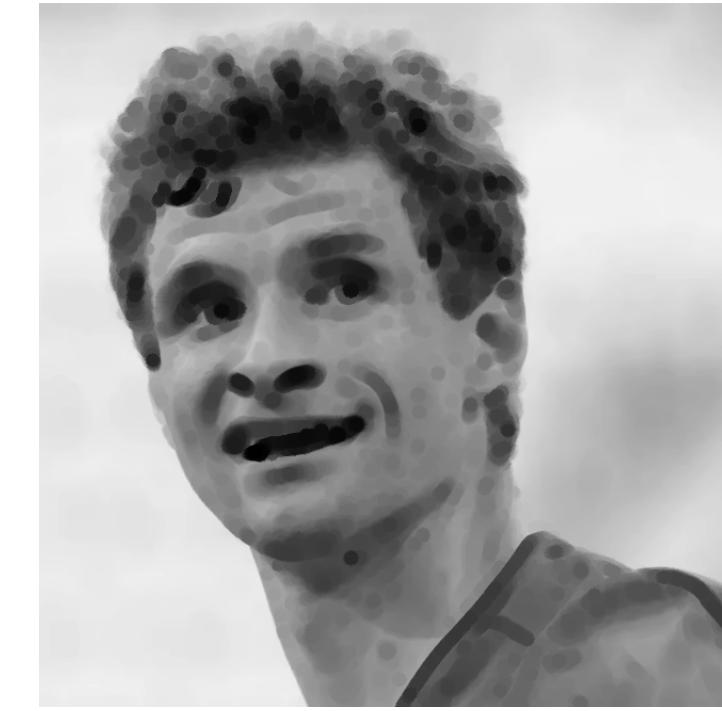
SE:



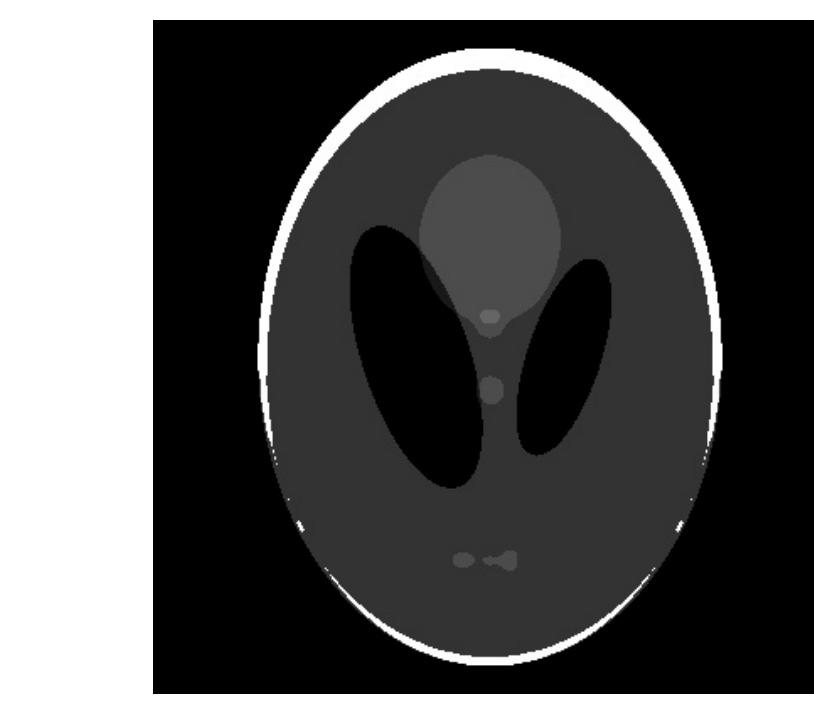
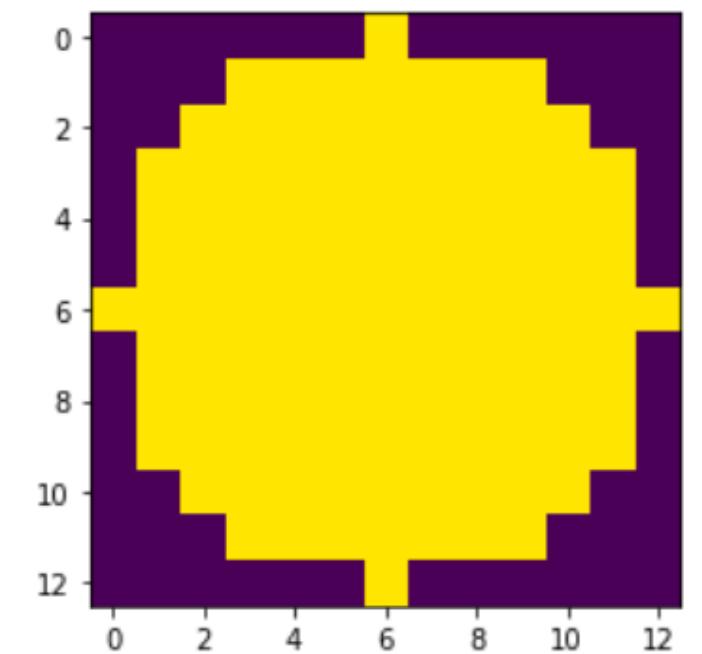
max() operation



min() operation



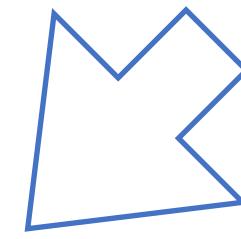
median() operation



Comparing various morphological operations

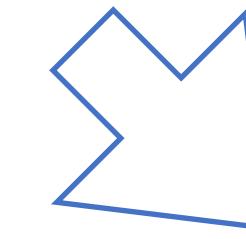
Bright areas are extended
Dark areas are thinned

max() operation



Dark areas are extended
Bright areas are thinned

min() operation



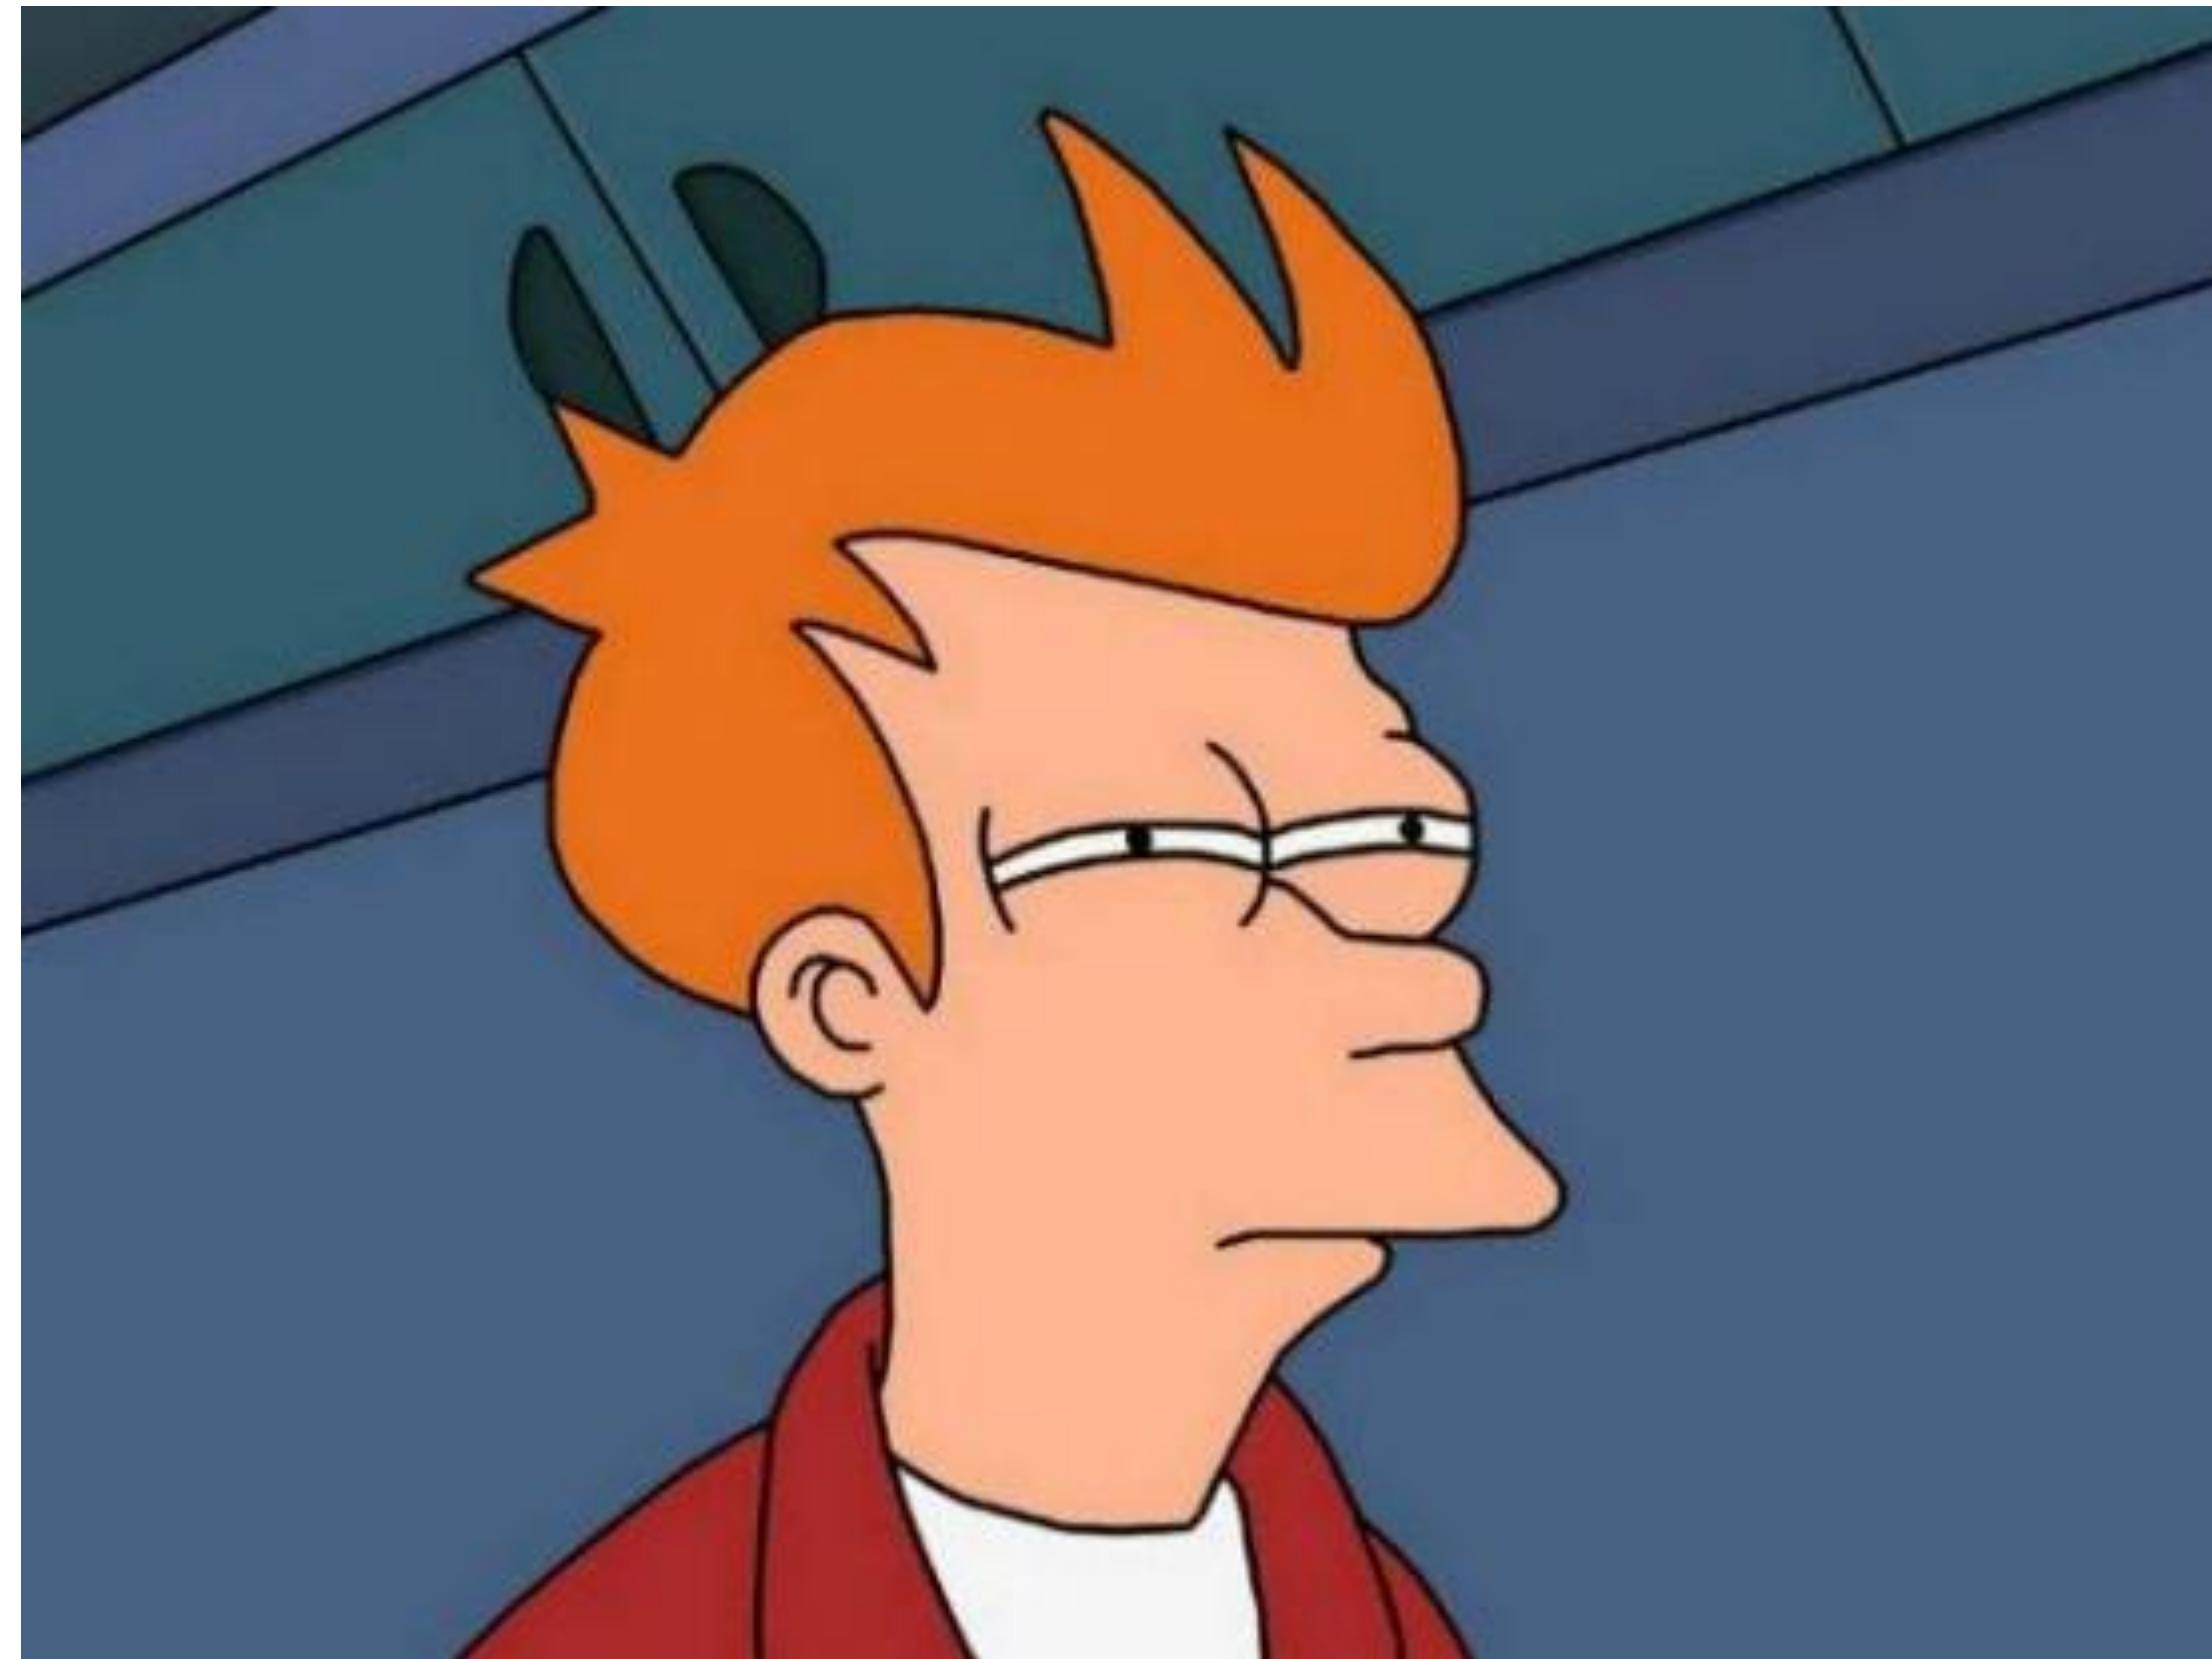
dilation filter



erosion filter

Morphology Kernels

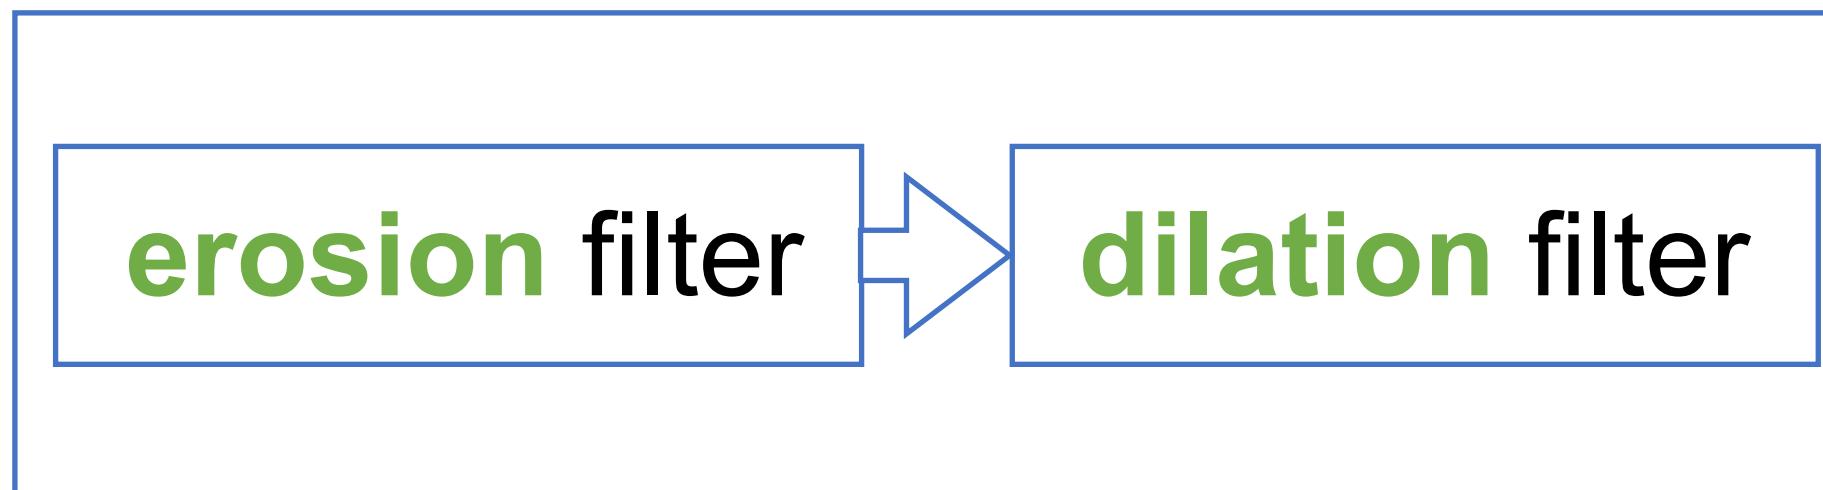
Opening & Closing



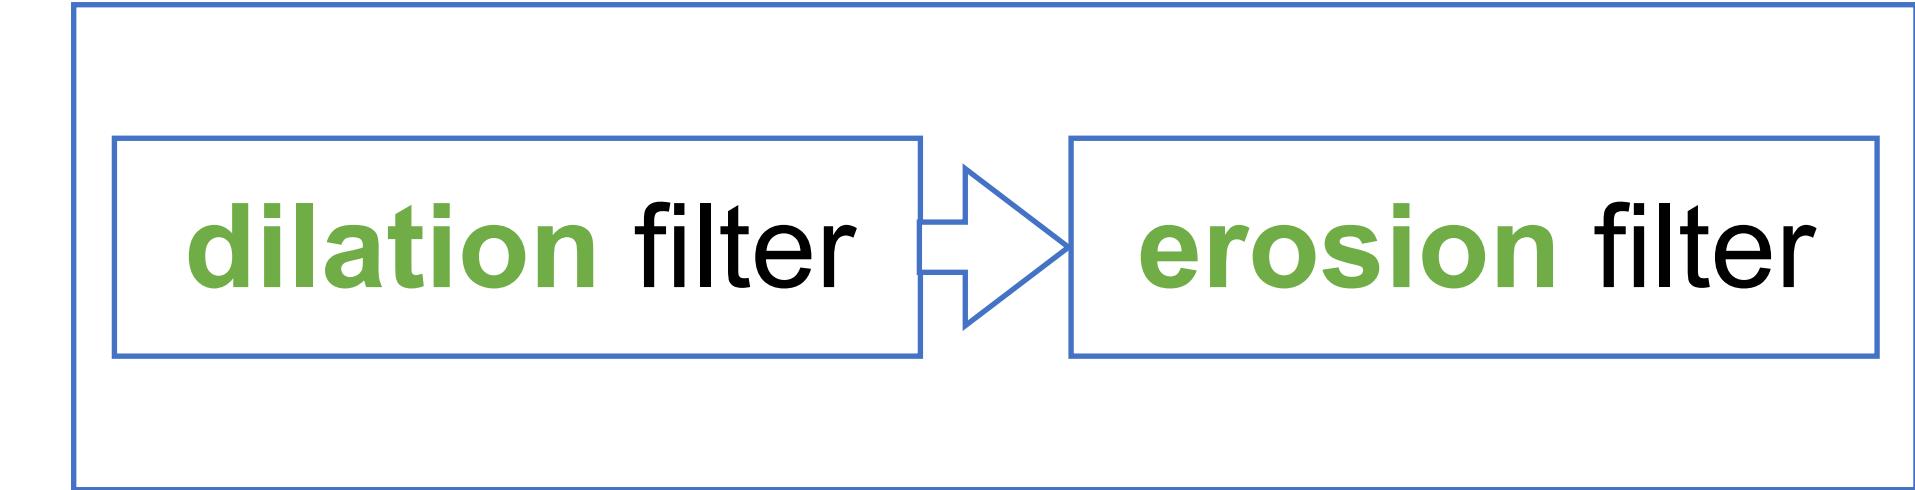
What if we combine
dilation and erosion?

Combining erosion and dilation

- Erosion and dilation filters can be applied after each other.



opening filter

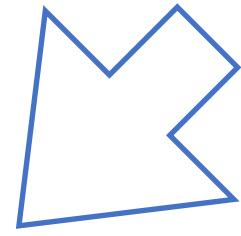


closing filter

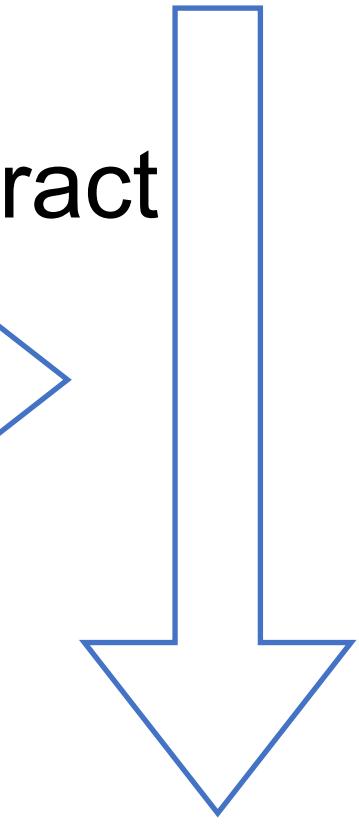
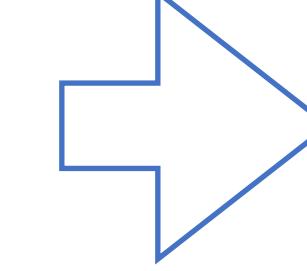


Let's have a bit more fun with opening and closing

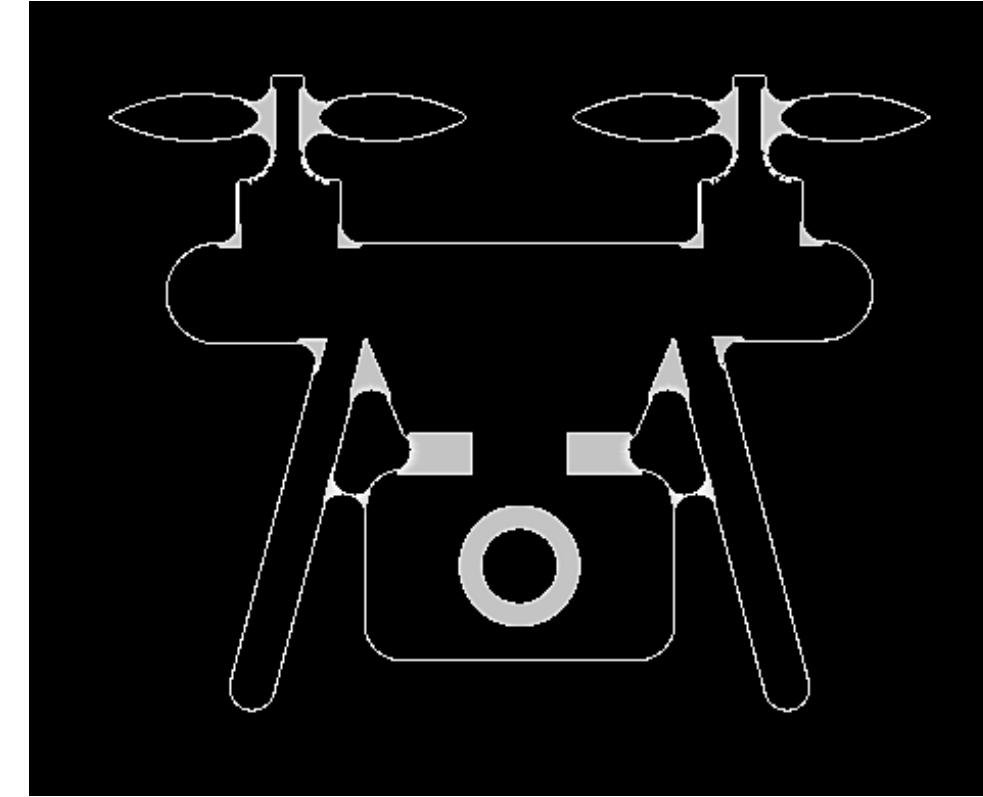
closing



subtract

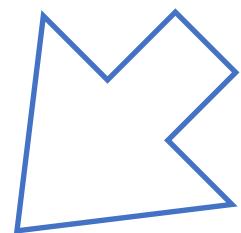


**hello
world**

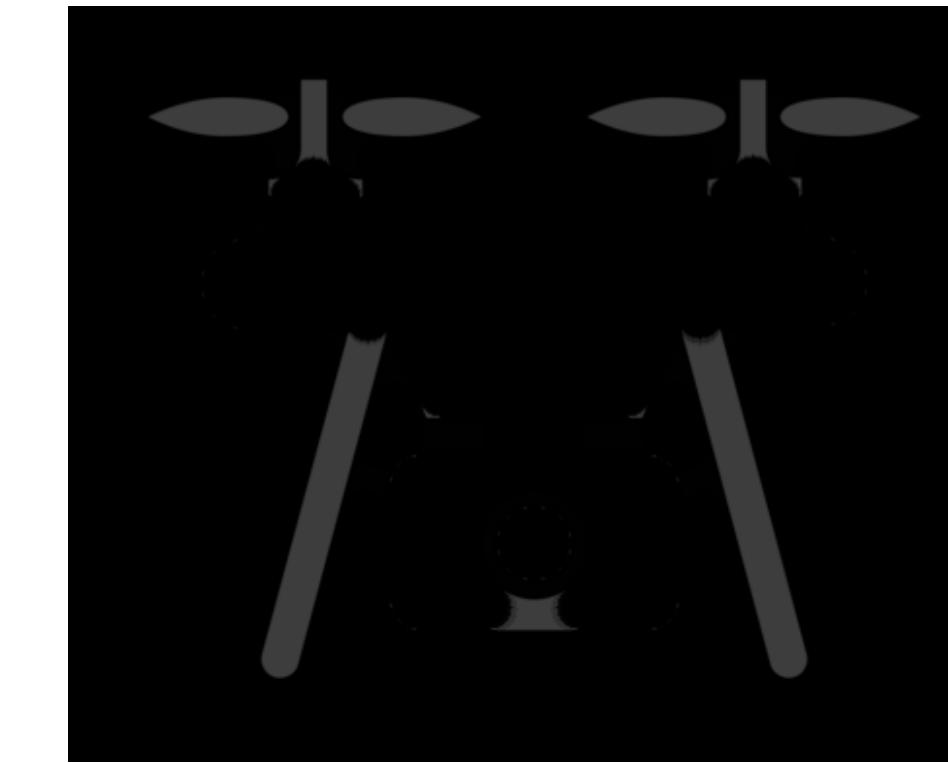
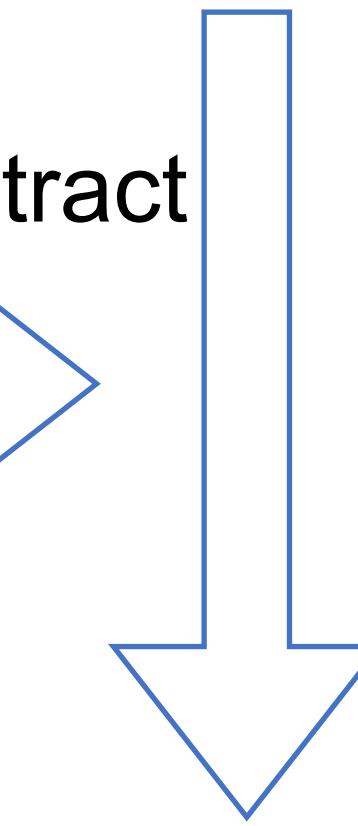
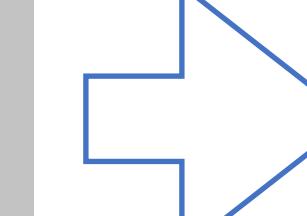


black tophat filter

opening

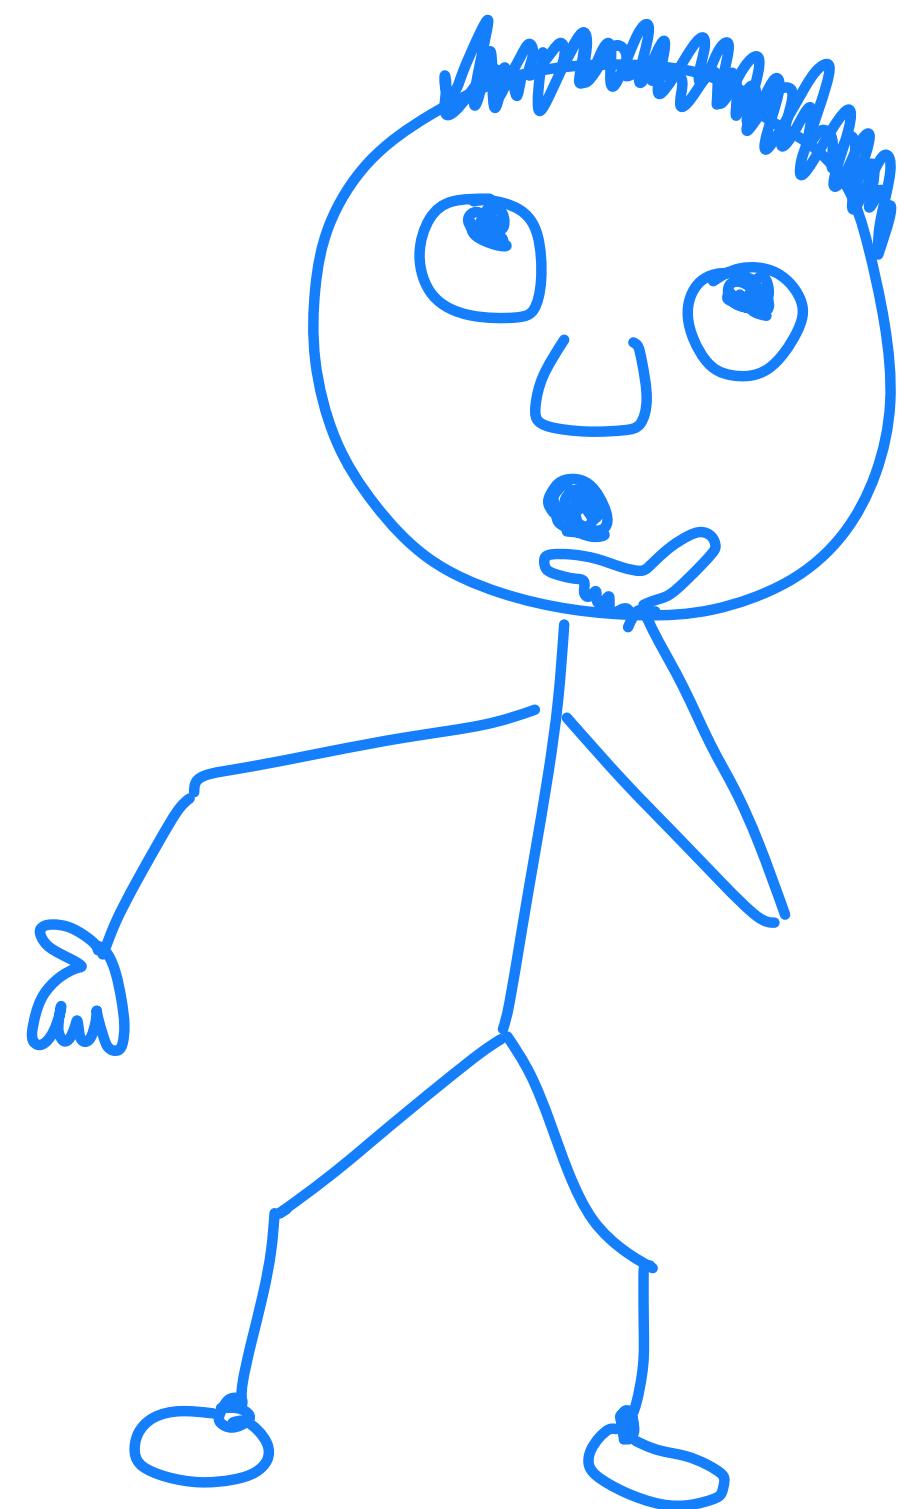


subtract



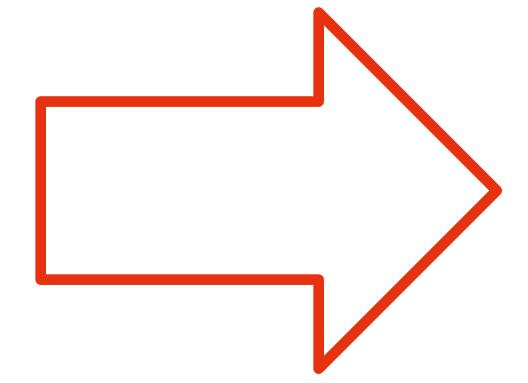
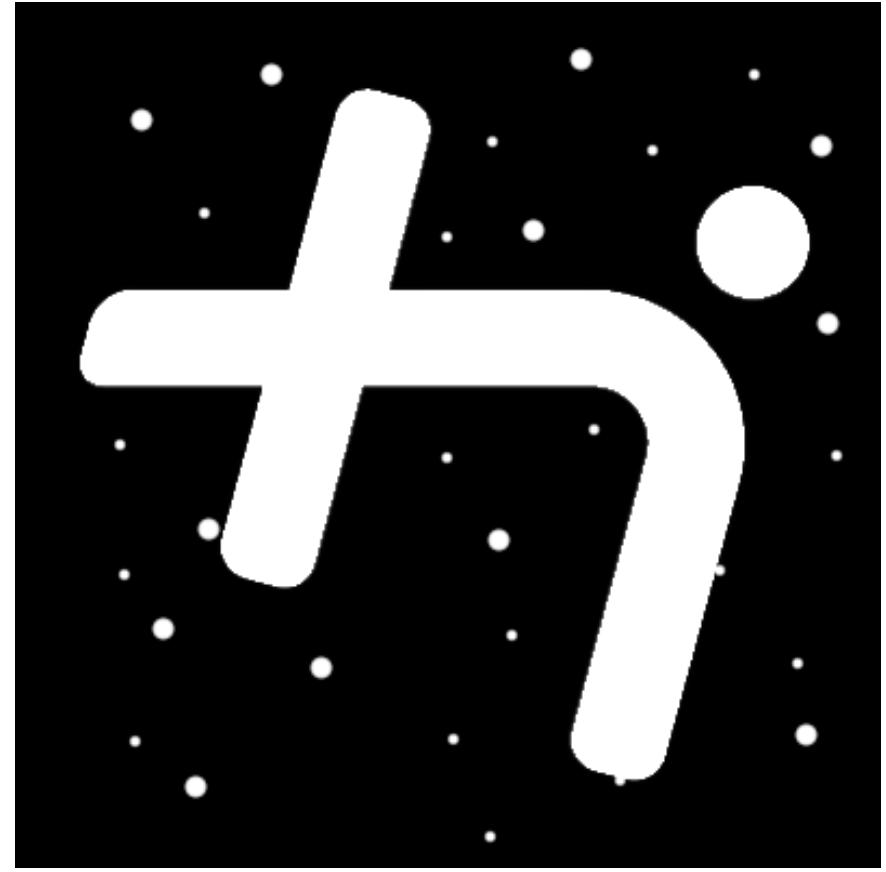
white tophat filter

Yay, so we have a new tool! What could we use this for?

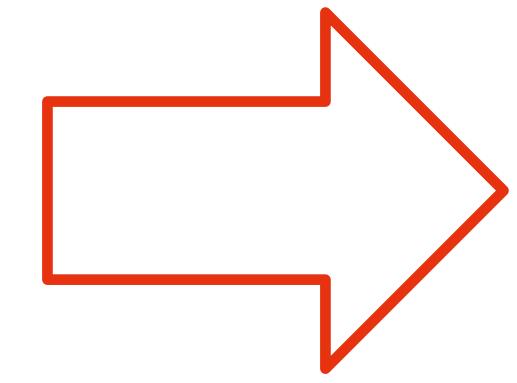
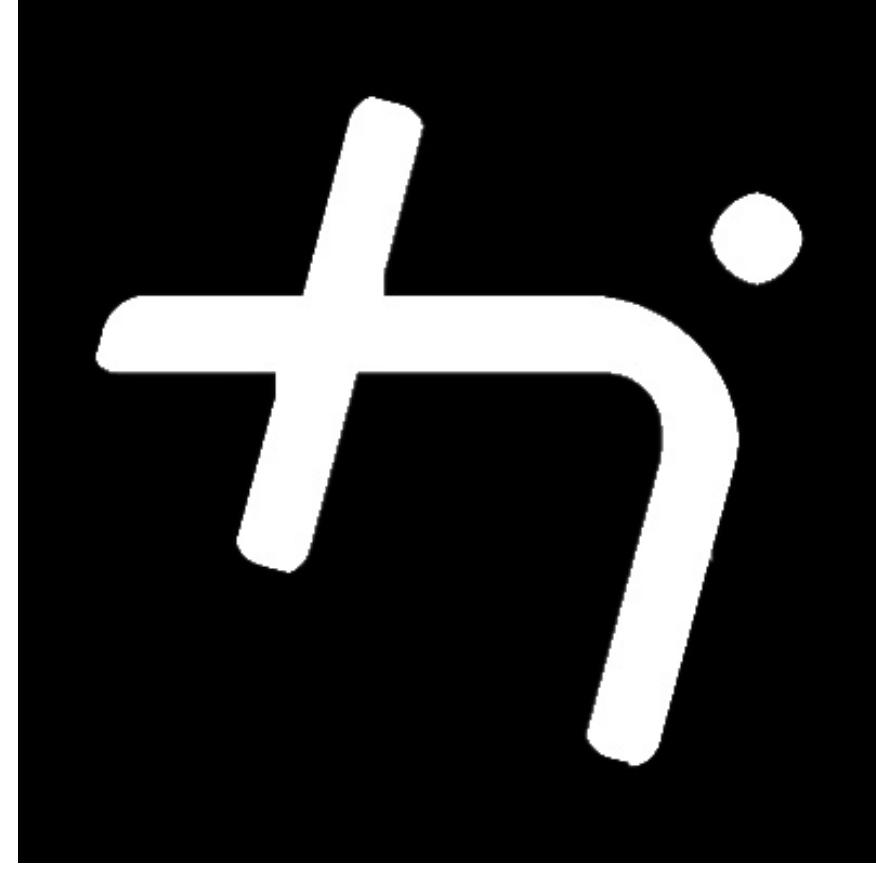


Morphology Kernels

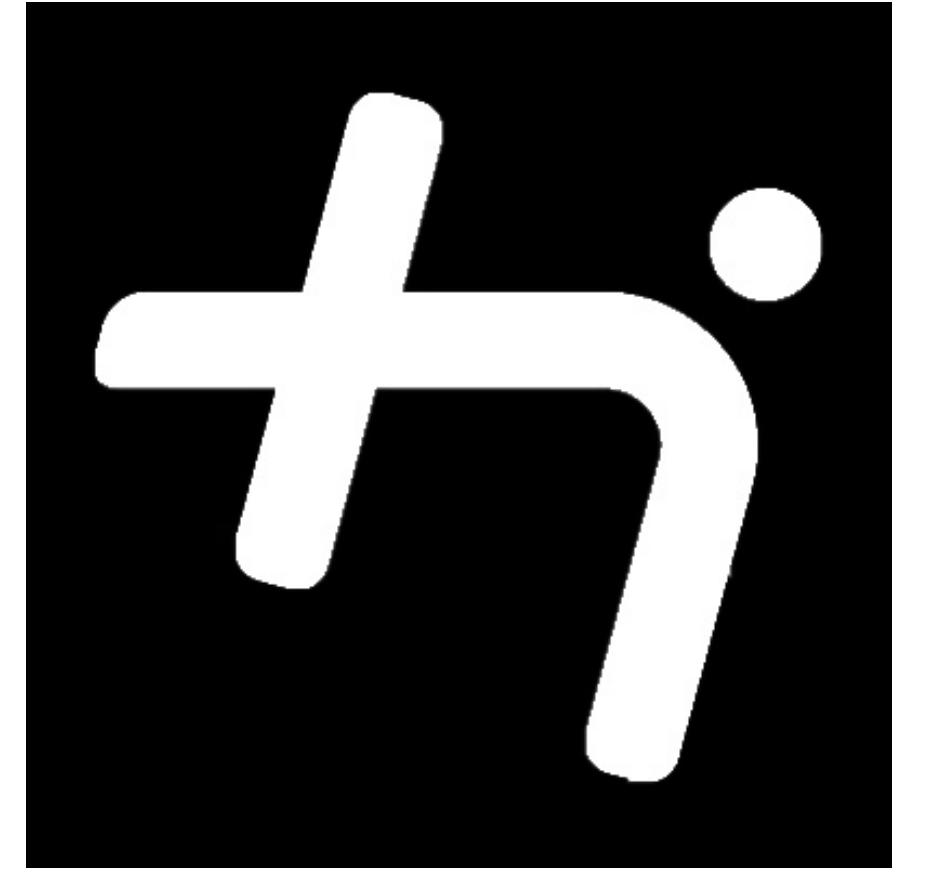
Opening



Erosion



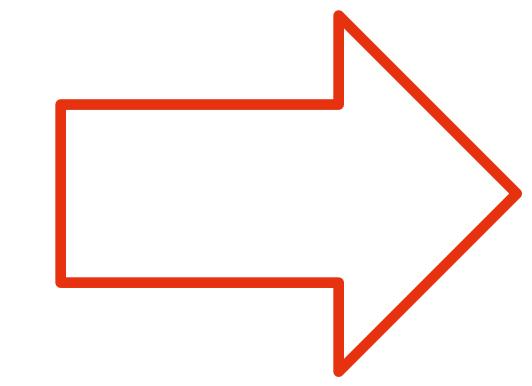
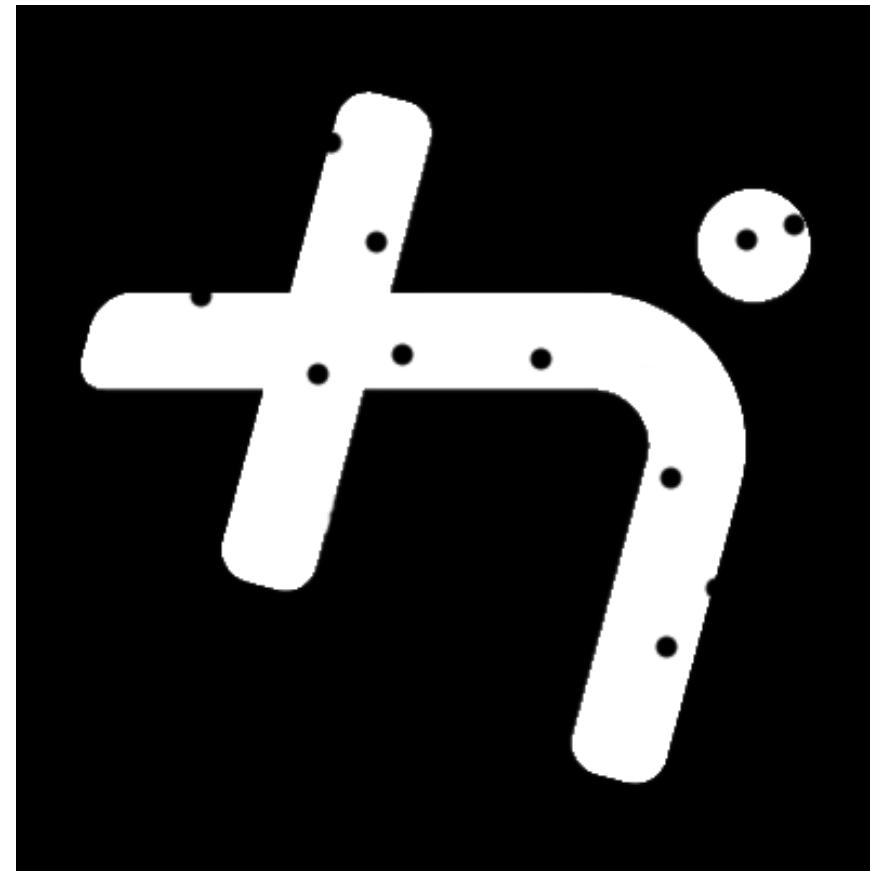
Dilation



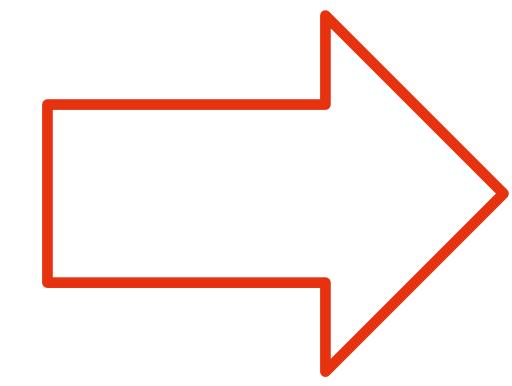
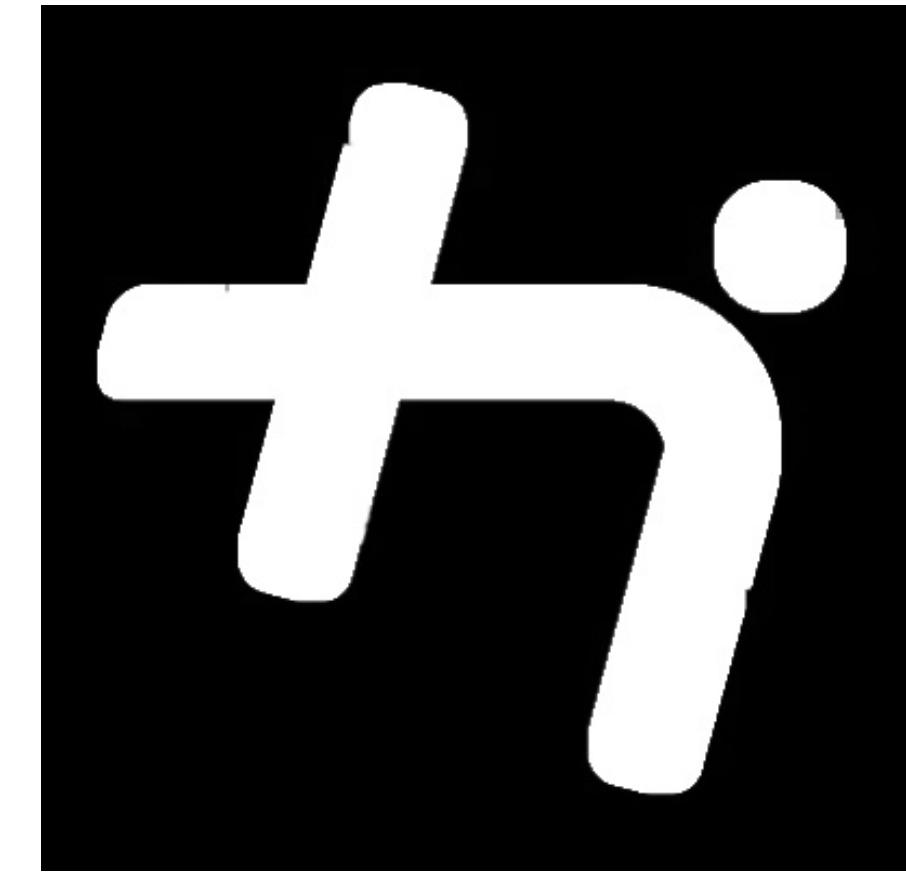
→ Opening: removing noise

Morphology Kernels

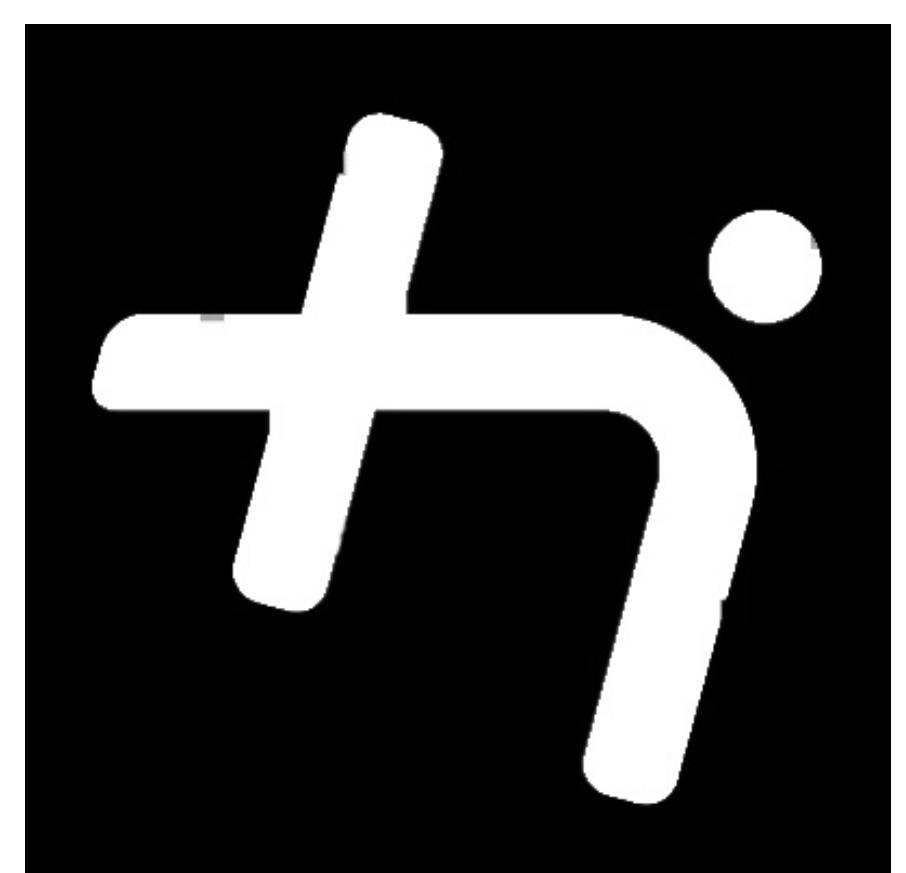
Closing



Dilation



Erosion

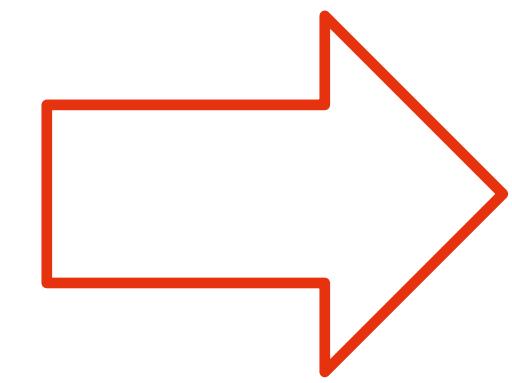
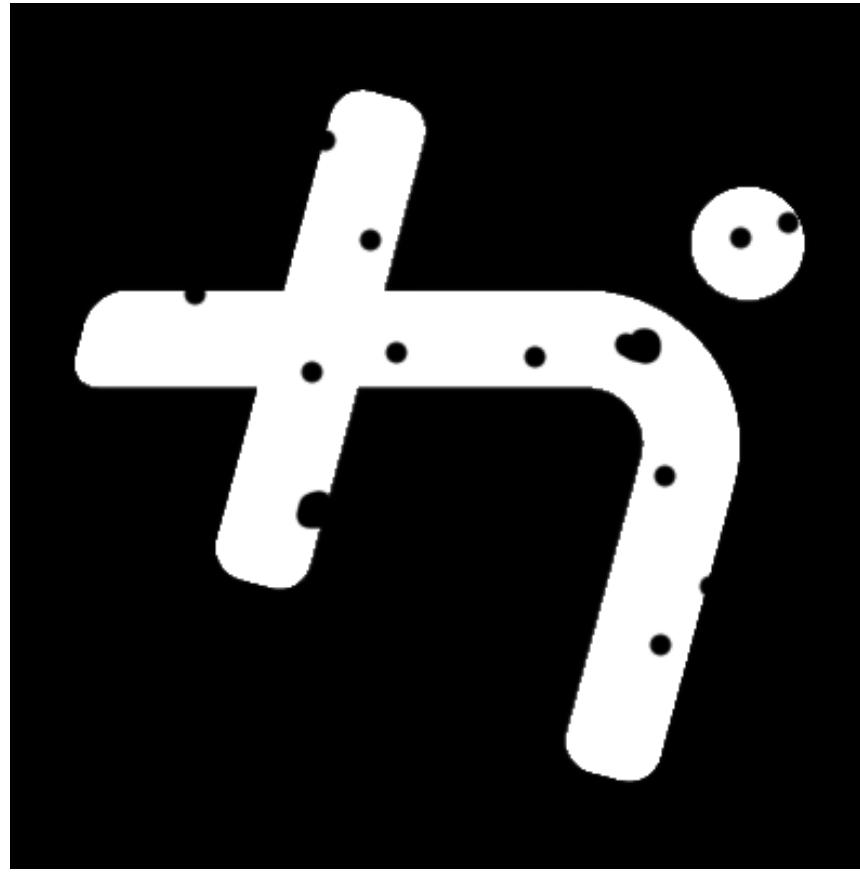


→ Closing: closes small holes in foreground objects

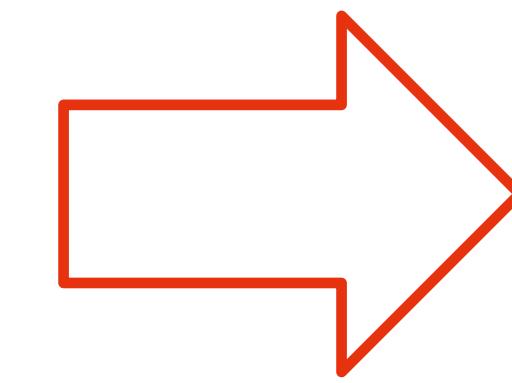
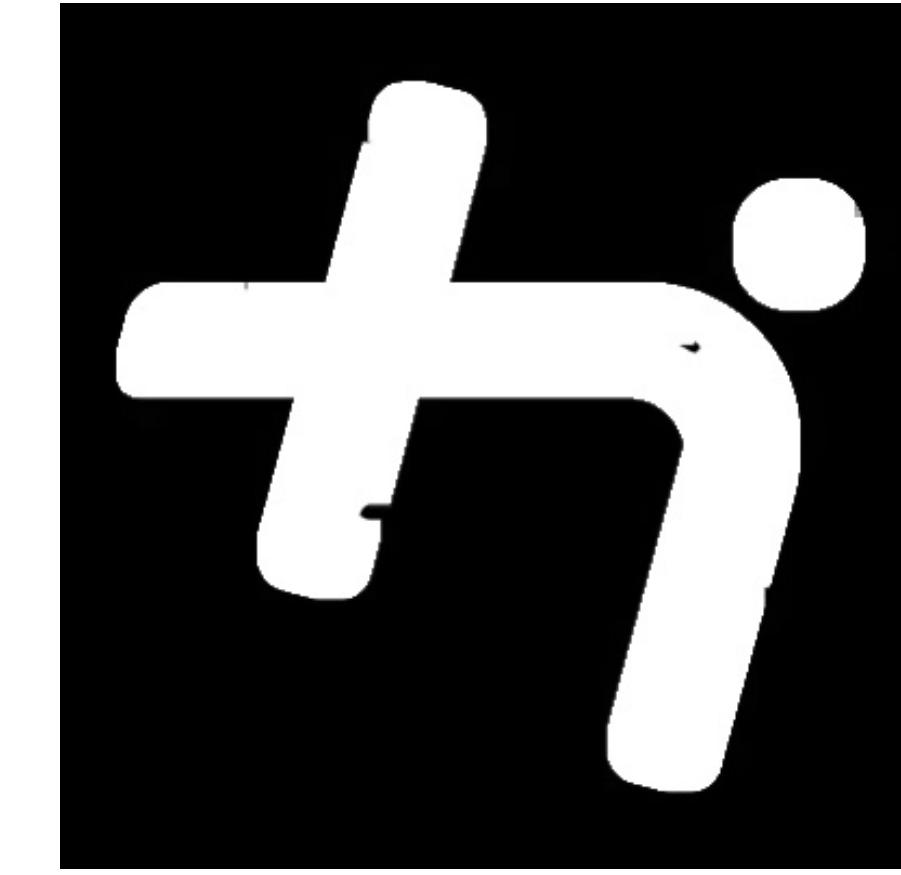
Morphology Kernels

Closing

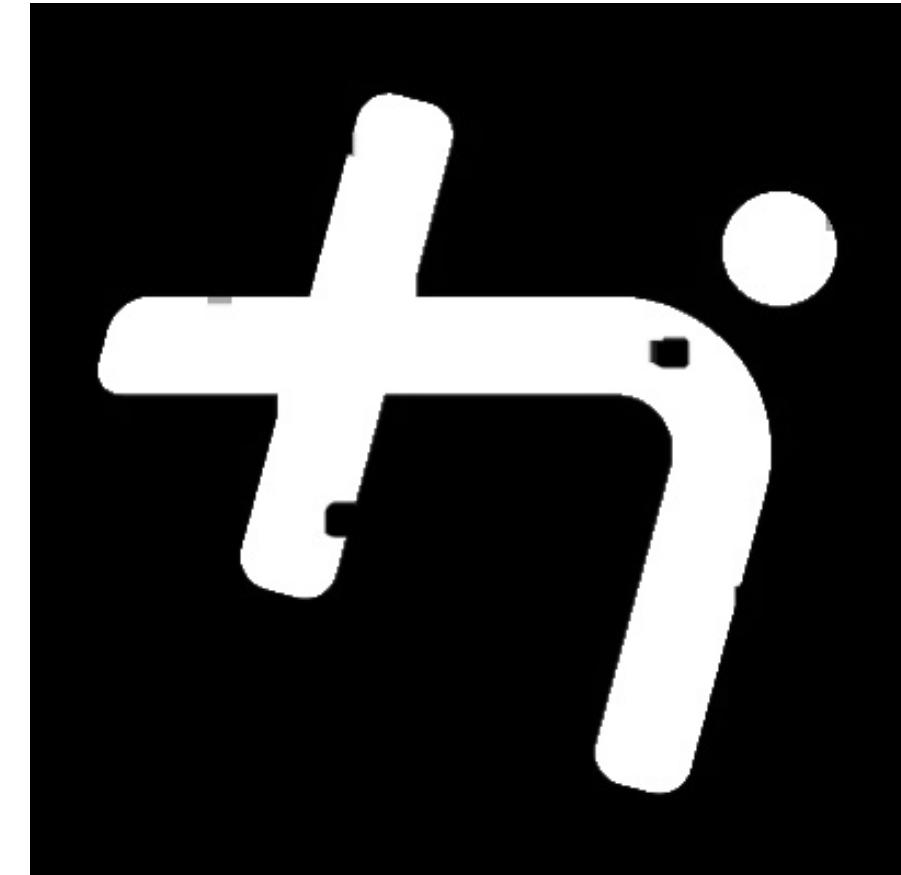
K: 10x10



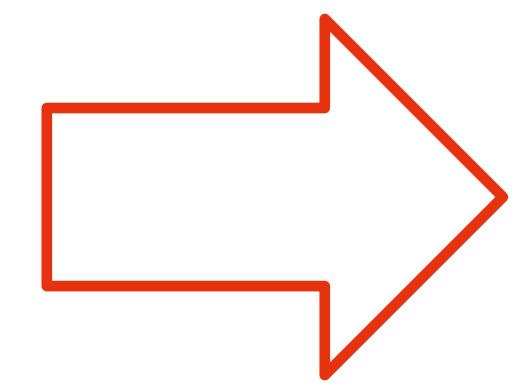
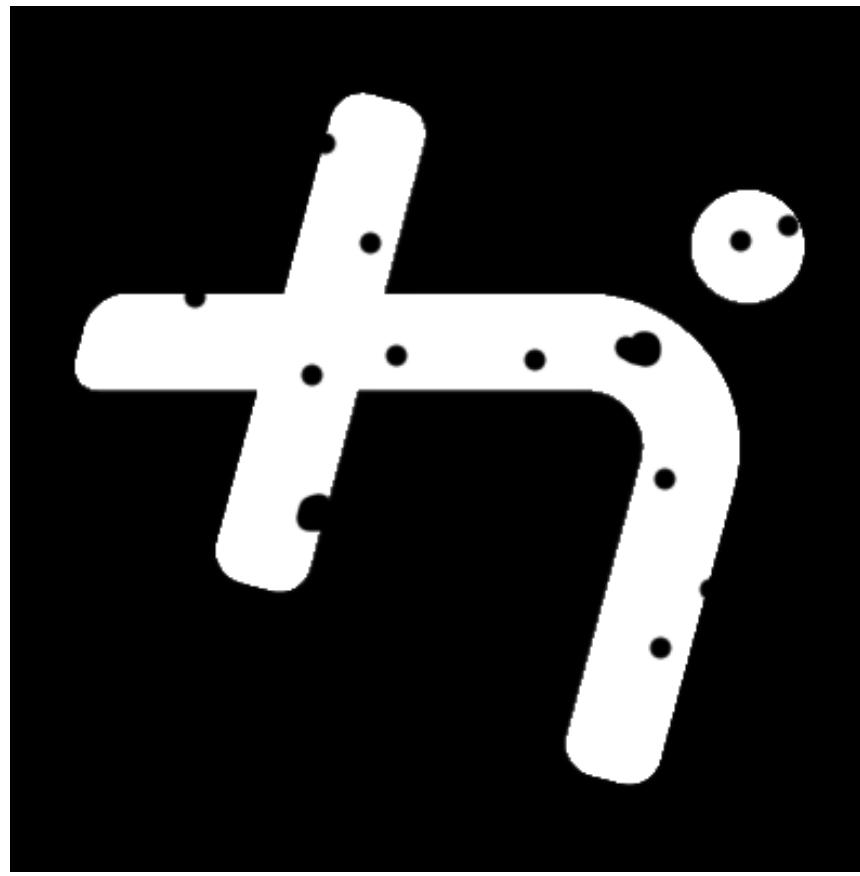
Dilation



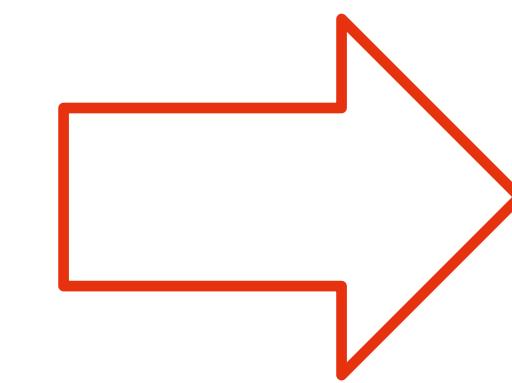
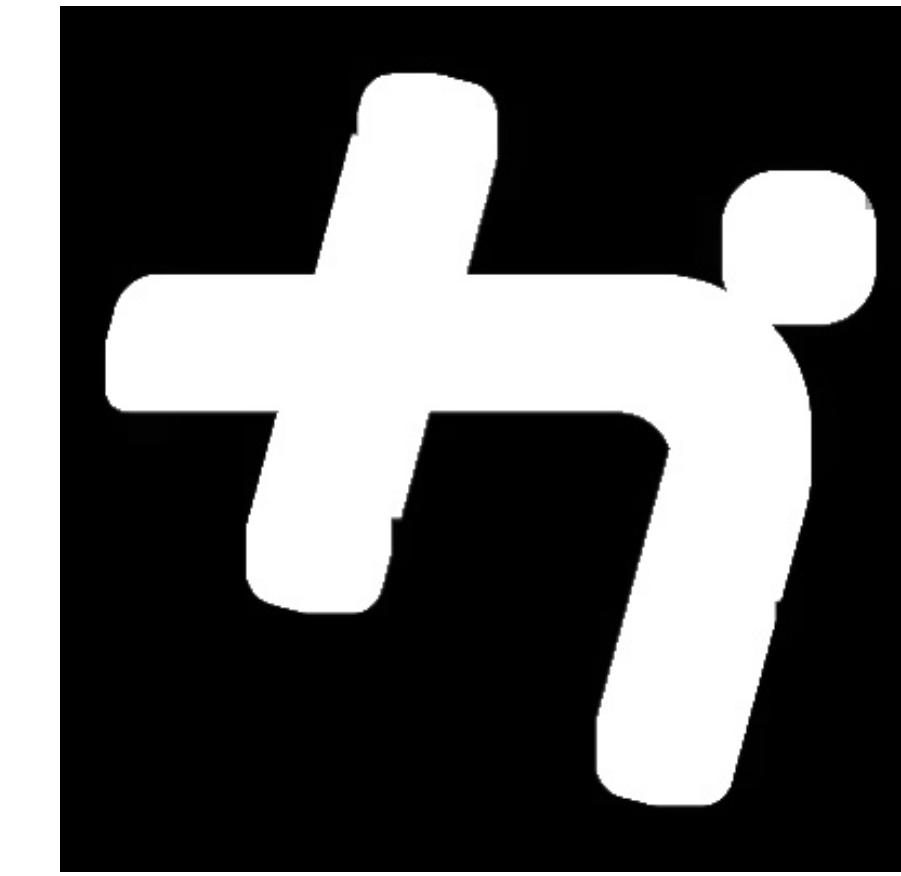
Erosion



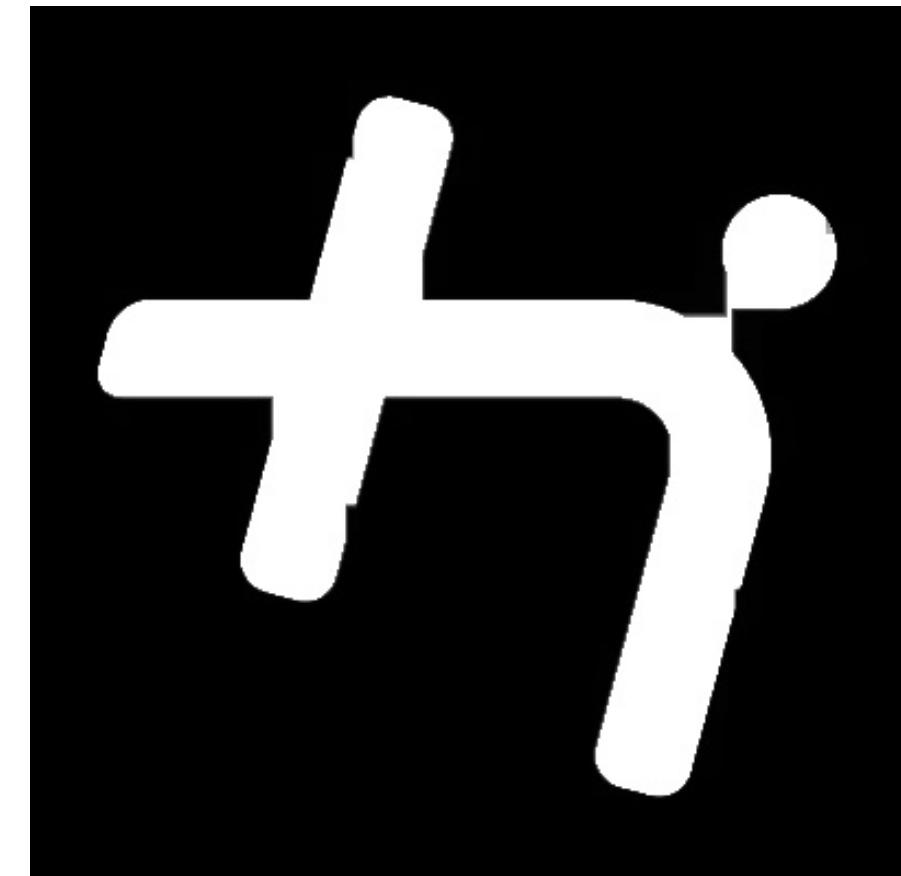
K: 20x20



Dilation



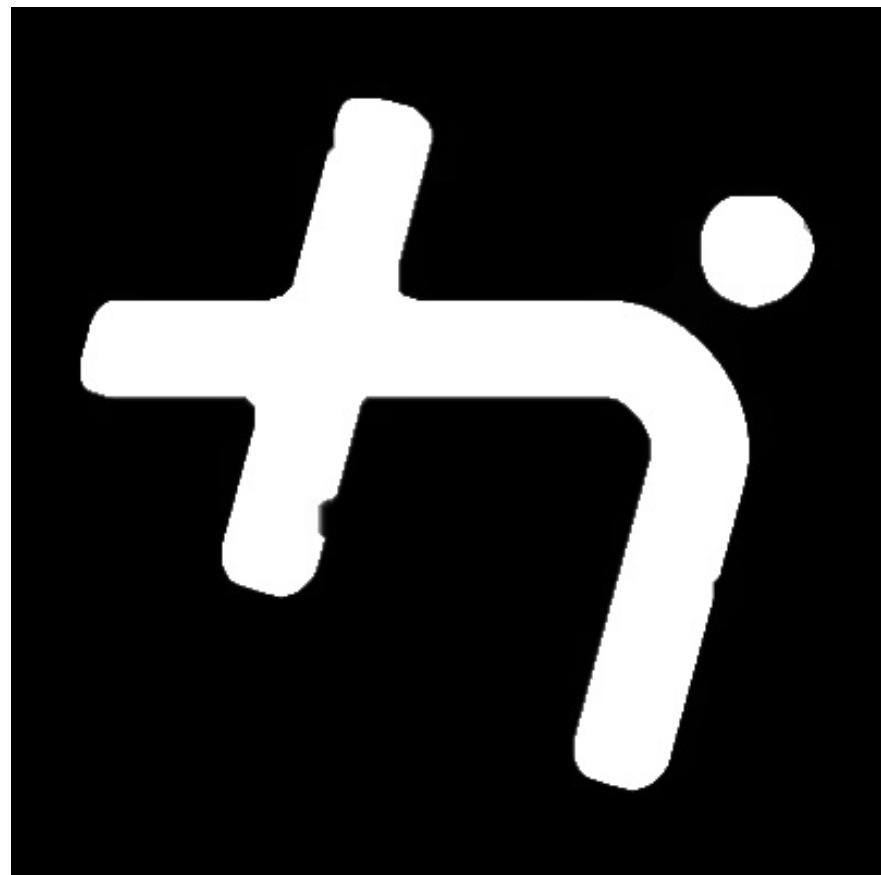
Erosion



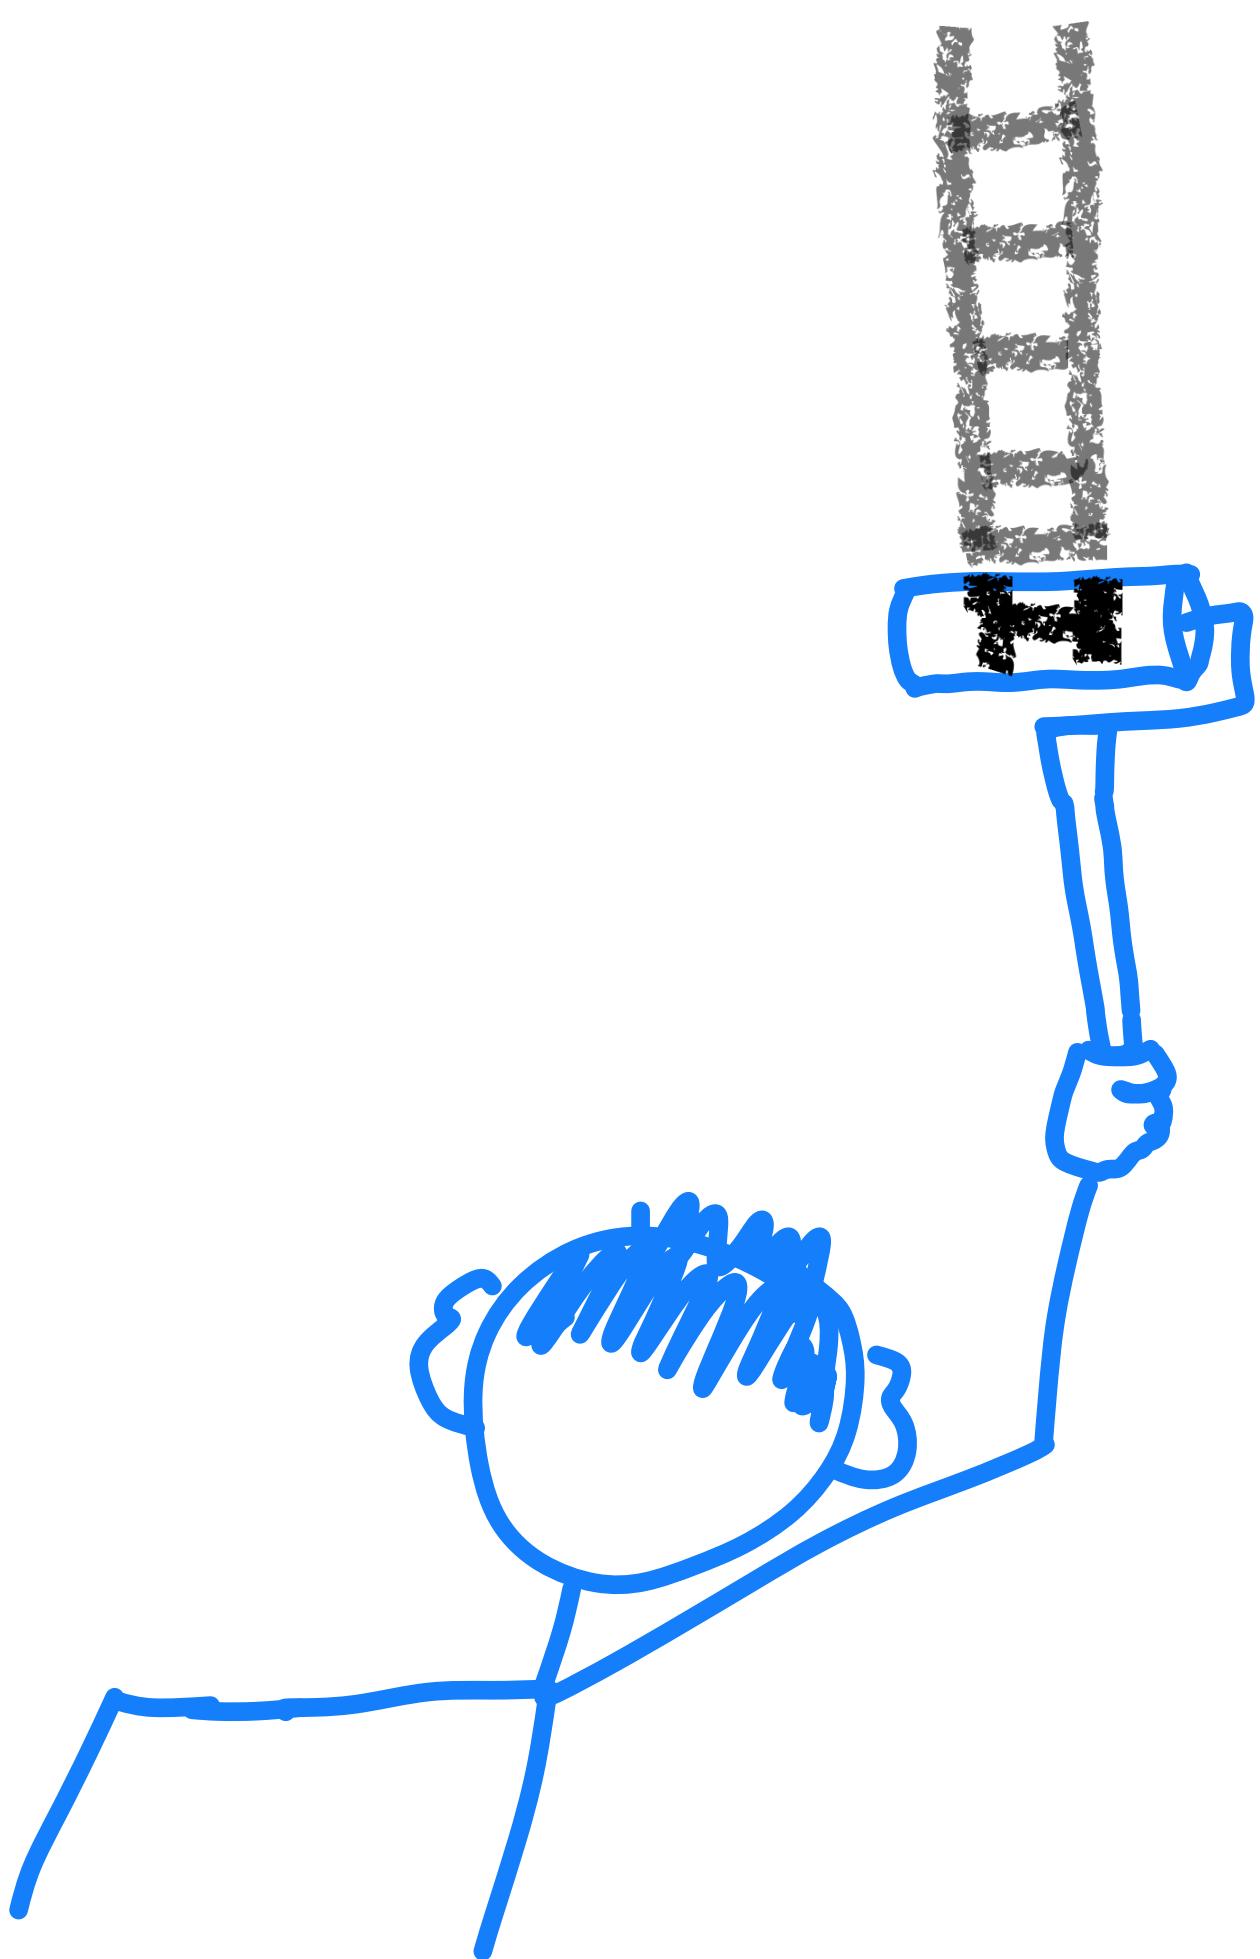
Morphology Kernels

Closing

- Options to get better result:
 - Try different structured elements
 - Try different kernel sizes
 - Try different number of iterations for dilate and erode



Example Solution:
Ellipse kernel 10x10
2 dilate followed by 2 erode



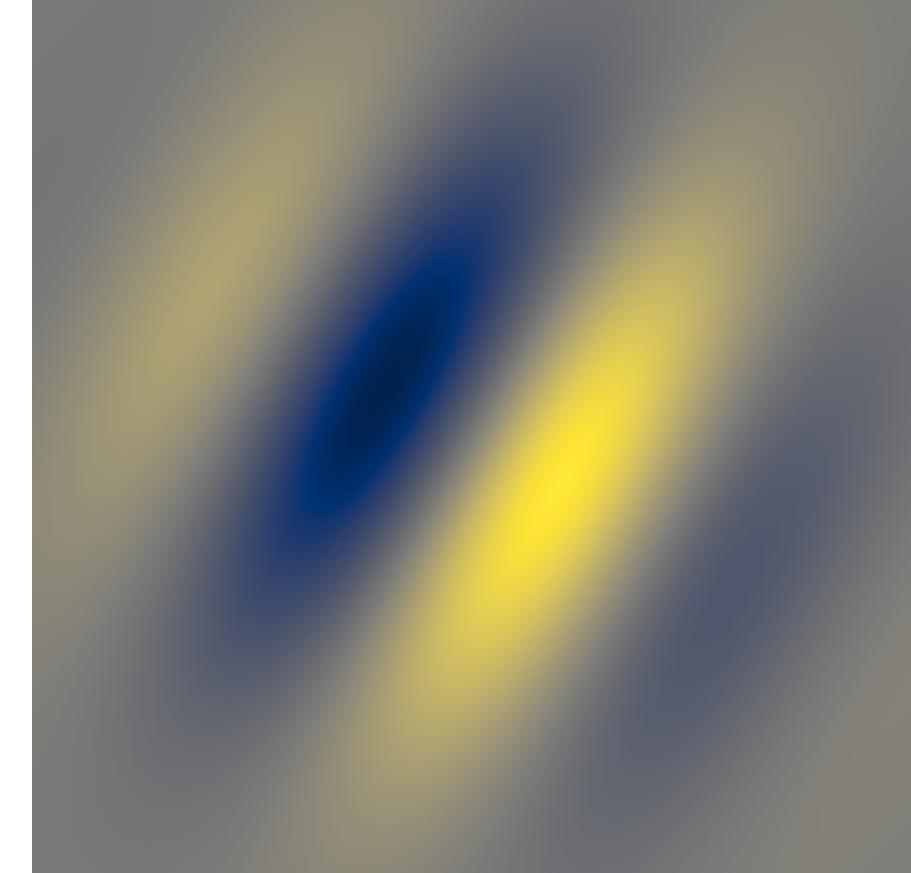
Texture

ernels

Texture Kernels

Gabor Filter

- Named after *Dennis Gabor*
- Used for texture analysis
- Localizes a signal in both, space and frequency domain
- Analyzes for specific frequency content in specific regions around the pixel of interest
- Particularly appropriate for texture representation and distinction

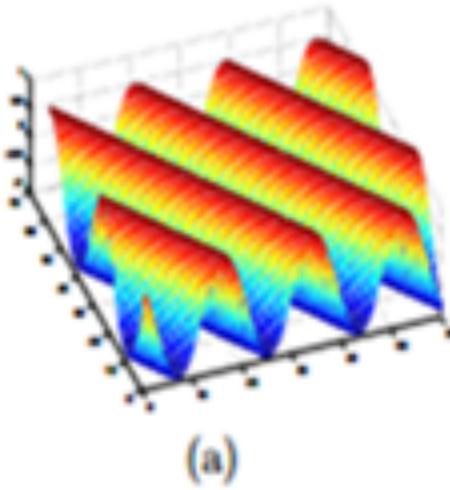


Quelle:
https://en.wikipedia.org/wiki/Gabor_filter

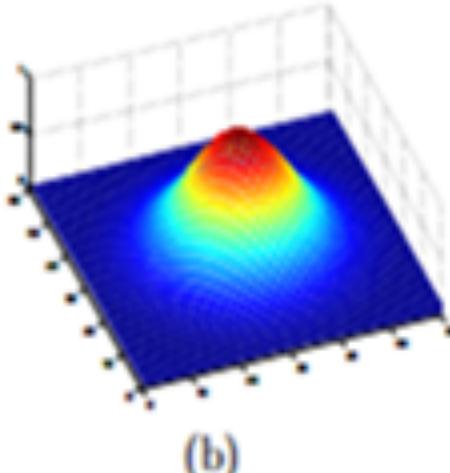
Texture Kernels

Gabor Filter

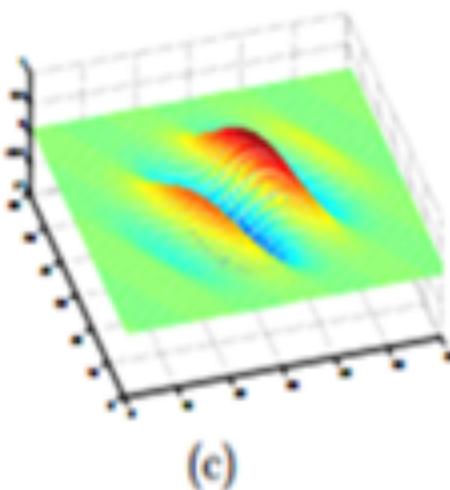
- In the spatial domain, a 2D Gabor filter is a **sinusoidal plane wave** modulated with a **Gaussian kernel function**.
- It is thus some special kind of band pass filter



A Sinusoid oriented 30° with X-axis



A 2-D Gaussian



The corresponding 2-D Gabor filter

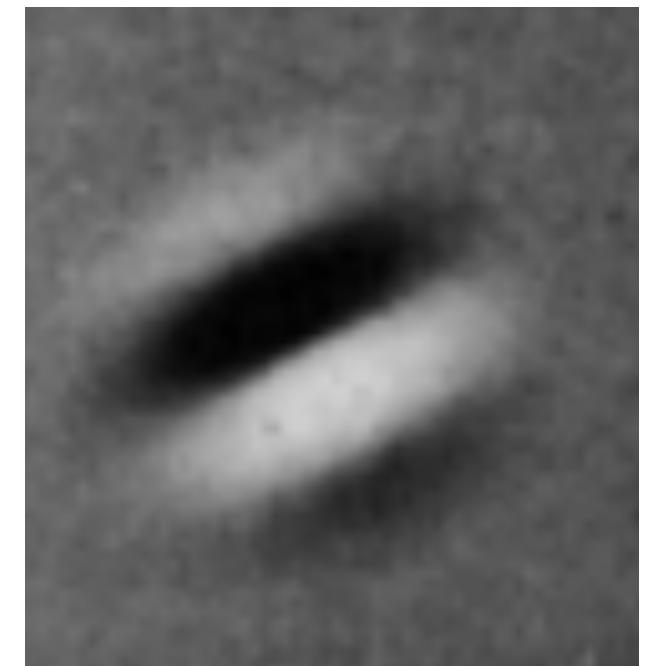
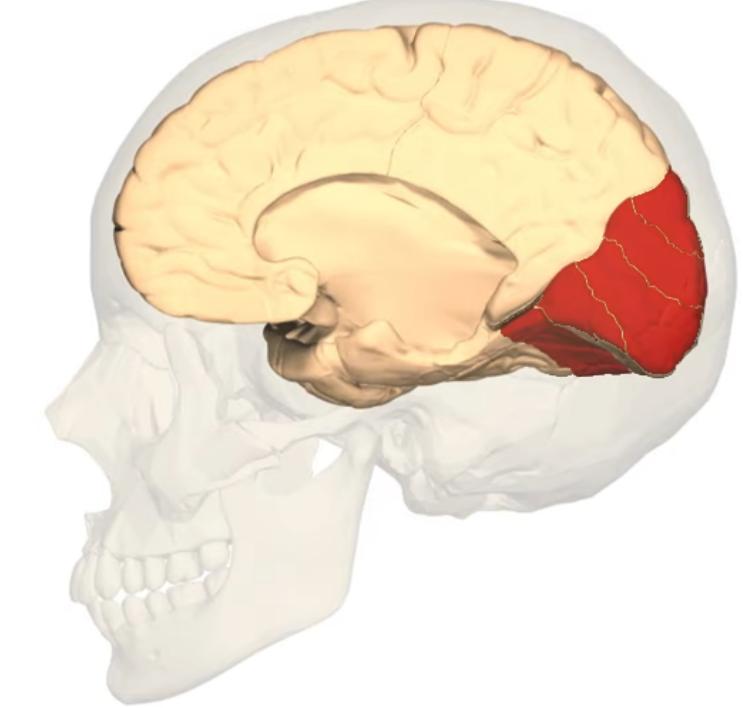
Quelle:
https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97



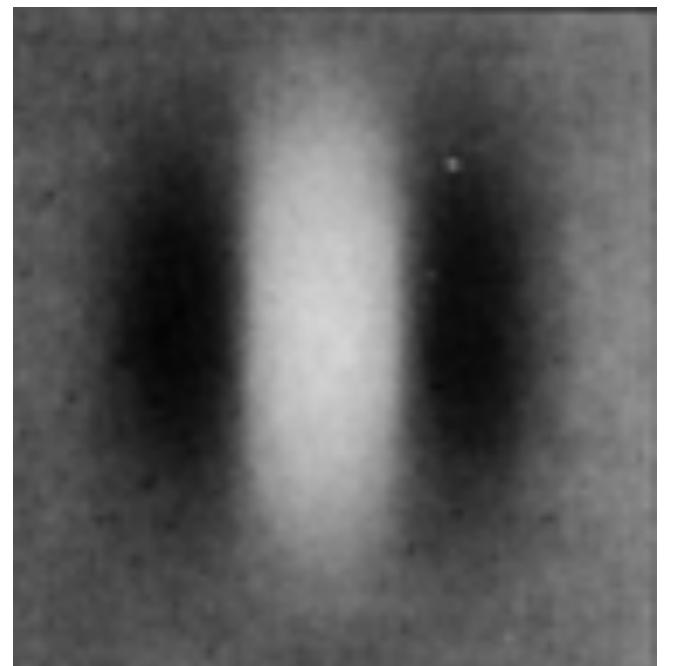
Texture Kernels

Gabor Filter

- Some papers have shown that some simple cells in the visual cortex of mammalian brains can be modelled by Gabor filters [1][2][3]
- Gabor filters are thus often related to human perception
- Gabor filters are very powerful and can detect textures as well as edges in many different orientations and scales



Neuron #1 of visual cortex
(model)



Neuron #2 of visual cortex
(model)

<http://viscomp.csail.mit.edu/resource/slides/lecture5.pptx>

[1] Olshausen, B. A. & Field, D. J. (1996). "Emergence of simple-cell receptive-field properties by learning a sparse code for natural images". *Nature*. **381** (6583): 607–609

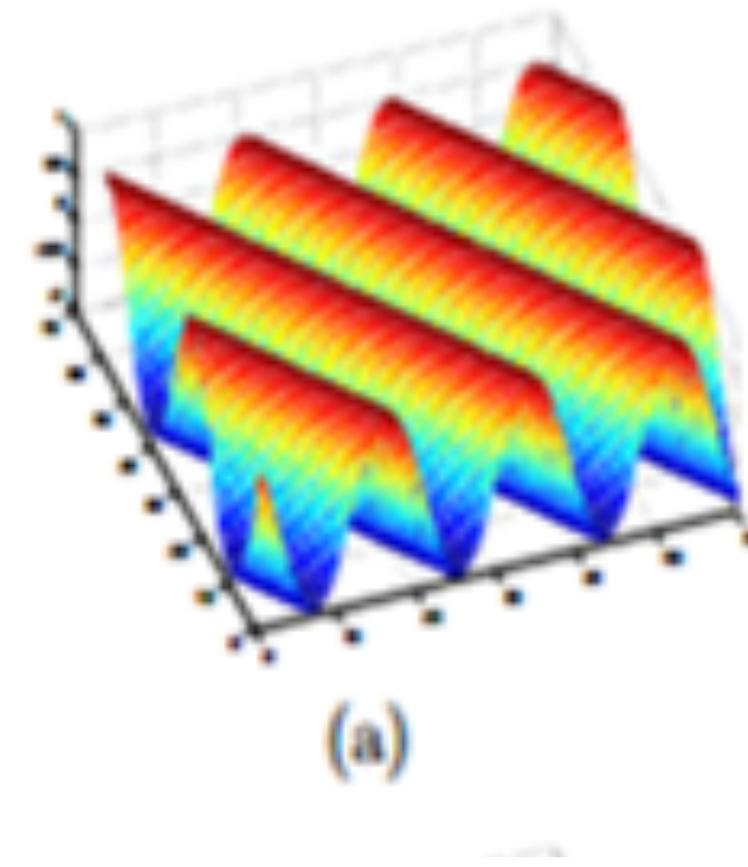
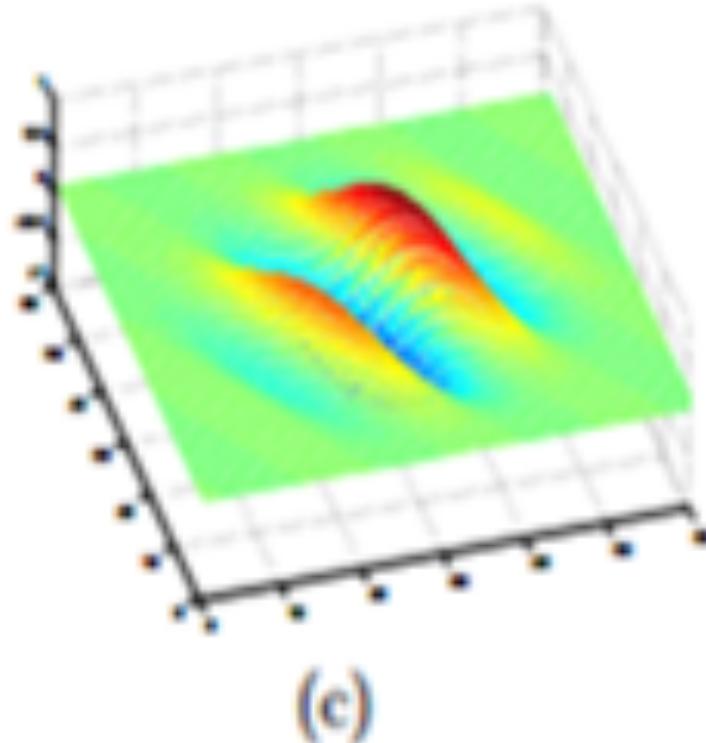
[2] Marčelja, S. (1980). "Mathematical description of the responses of simple cortical cells". *Journal of the Optical Society of America*. **70** (11): 1297–1300

[3] Daugman, John G. (1985-07-01). "Uncertainty relation for resolution in space, spatial frequency, and orientation optimized by two-dimensional visual cortical filters". *Journal of the Optical Society of America A*. **2** (7): 1160–9

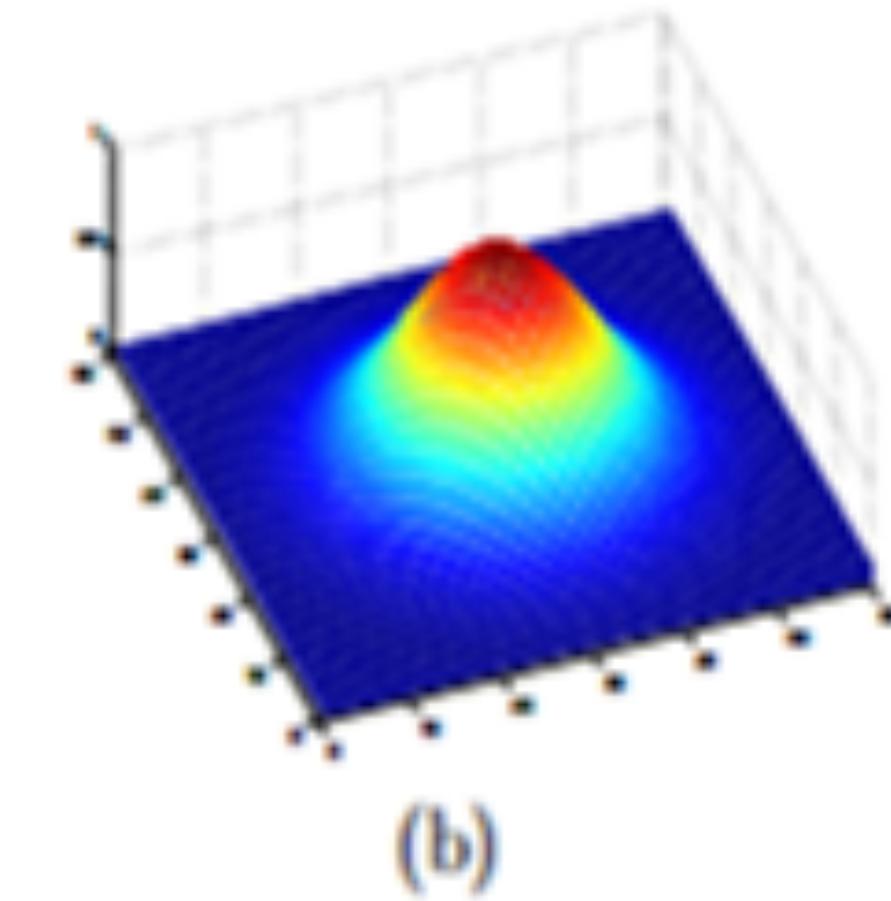
Texture Kernels

Gabor Filter

**Sinusoidal
plane wave**



Gaussian function





Texture Kernels

Gabor Filter

- The Gabor filter consists of a **real** and a **complex** component and is defined by:

$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \exp\left(i\left(2\pi\frac{x'}{\lambda} + \psi\right)\right)$$

where

$$x' = x\cos\theta + y\sin\theta \quad y' = -x\sin\theta + y\cos\theta$$



- The **real** and a **complex** component can also be used individually and are then given by

Real:
$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \cos\left(2\pi \frac{x'}{\lambda} + \psi\right)$$

Imaginary:
$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \sin\left(2\pi \frac{x'}{\lambda} + \psi\right)$$



$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \exp\left(i\left(2\pi\frac{x'}{\lambda} + \psi\right)\right)$$

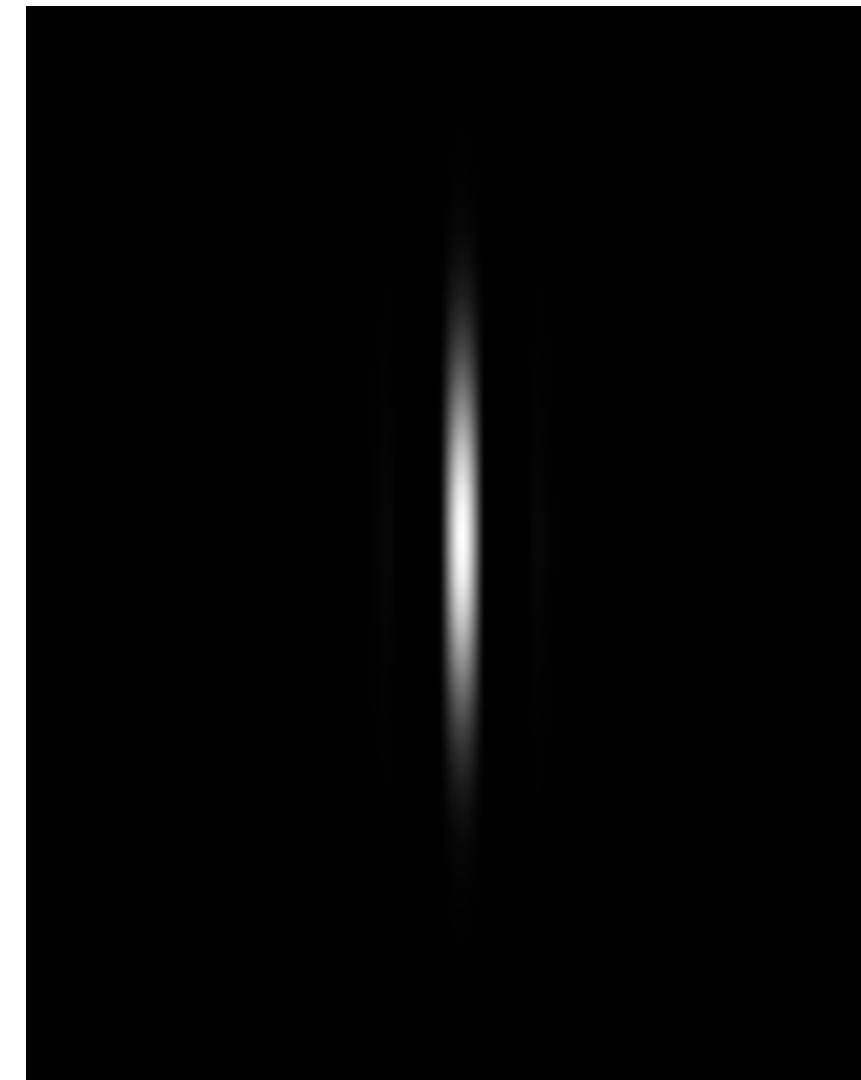
What do the parameters mean?

- λ : Wavelength of the sinusoidal component
- θ : Orientation of the normal to the parallel stripes
- ψ : Phase offset of the sinusoidal function
- σ : Standard deviation of the Gaussian
- γ : Spatial aspect ratio → ellipticity

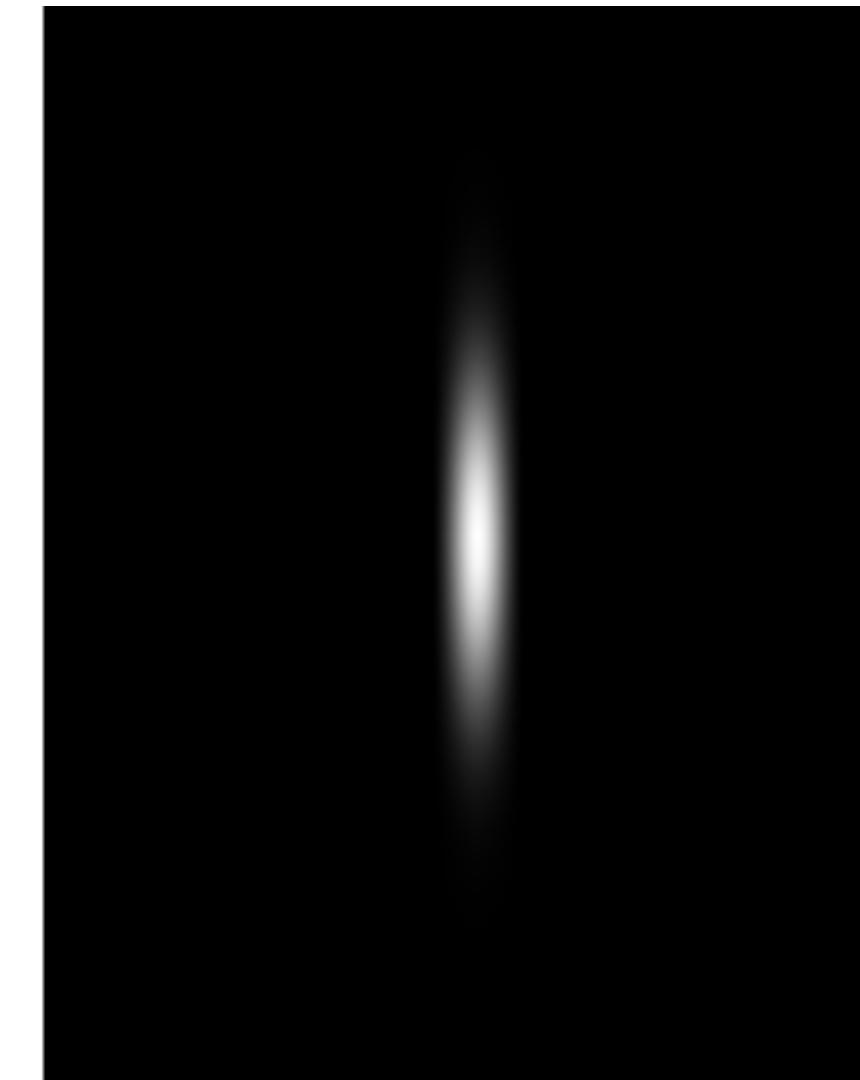
Let's dig into those parameters now ...

λ : Wavelength of the sinusoidal component

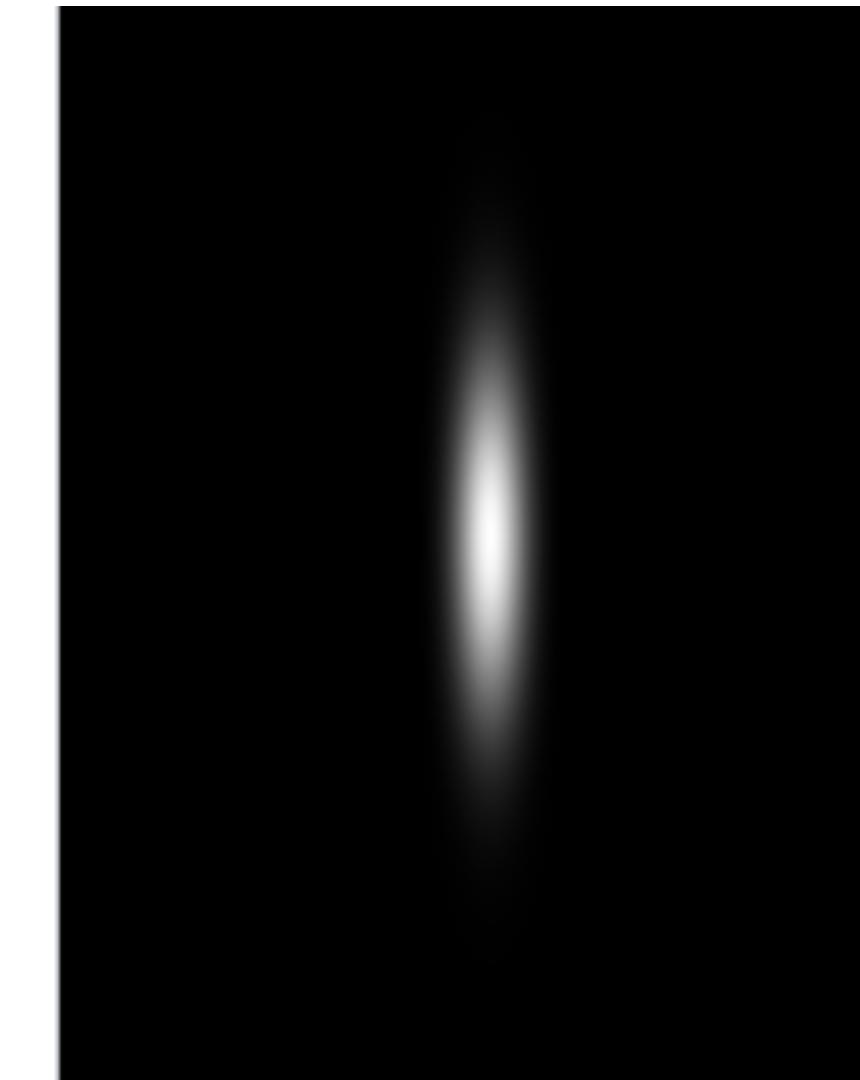
The wavelength influences the width of the peaks in the sinus signal, thus increasing λ lead to thicker stripes, decreasing λ leads to thinner stripes



Lambda (λ) = 30



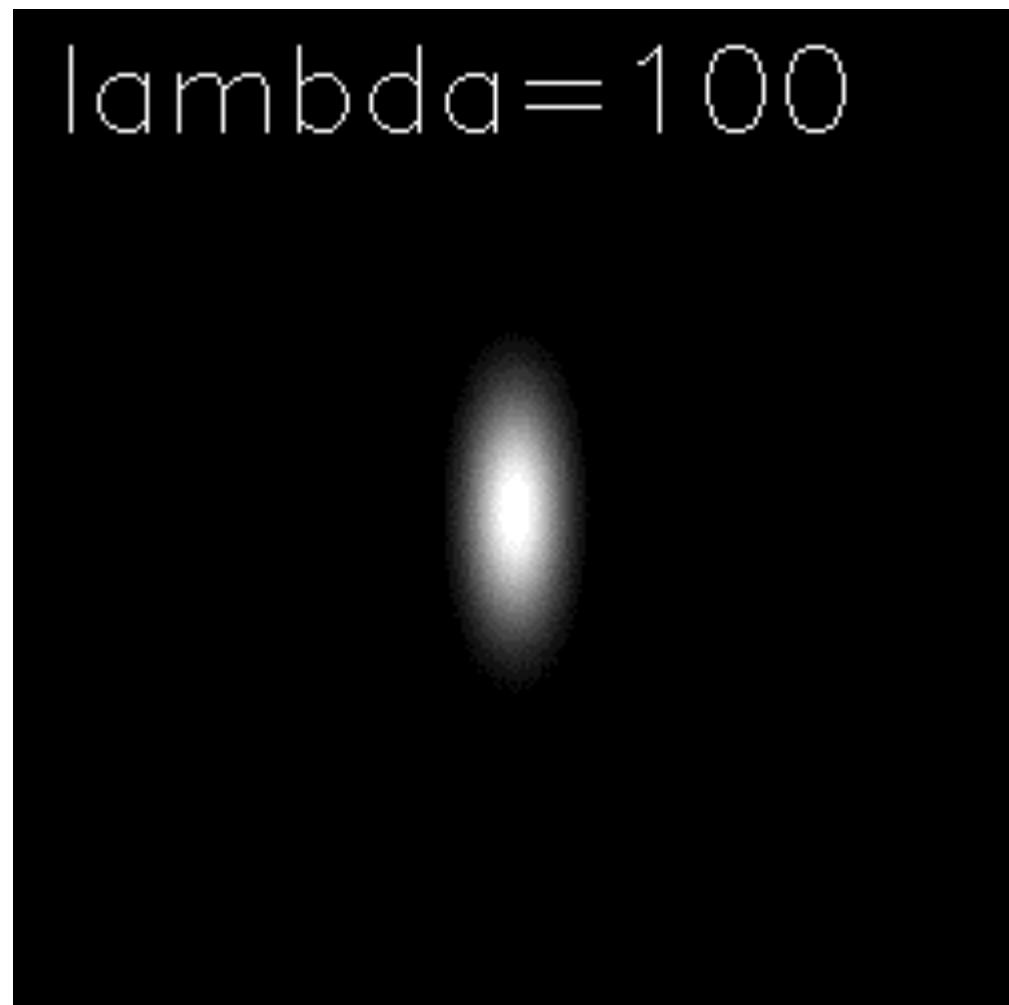
Lambda (λ) = 60



Lambda (λ) = 100

Quelle: https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97

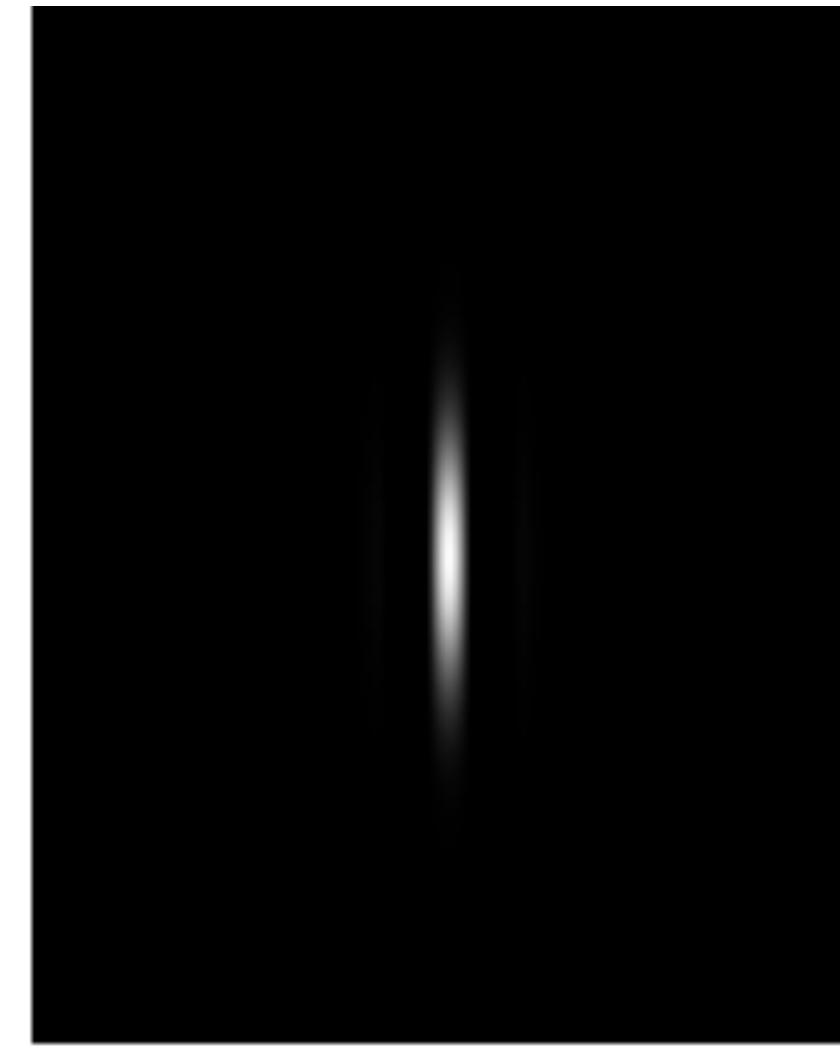
λ : Wavelength of the sinusoidal component



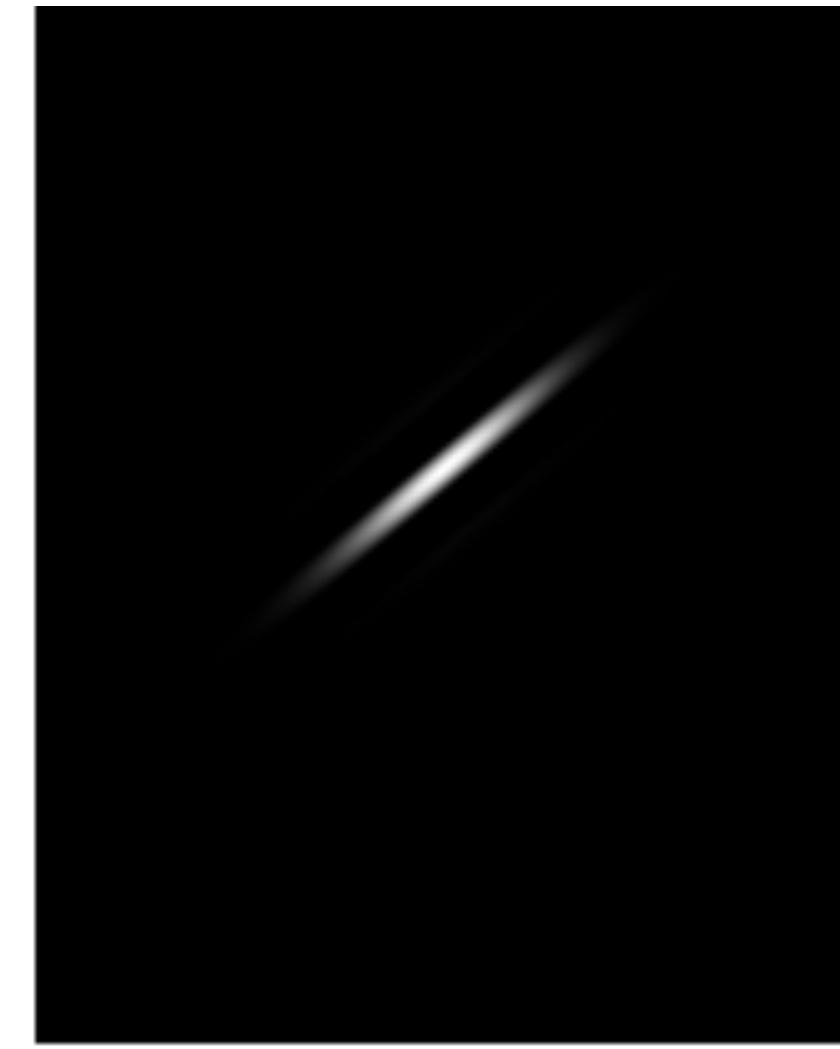
θ Theta = 0
 ψ Psi = 0
 σ Sigma = 10
 γ Gamma = 0.5

θ : Orientation of the normal to the parallel stripes

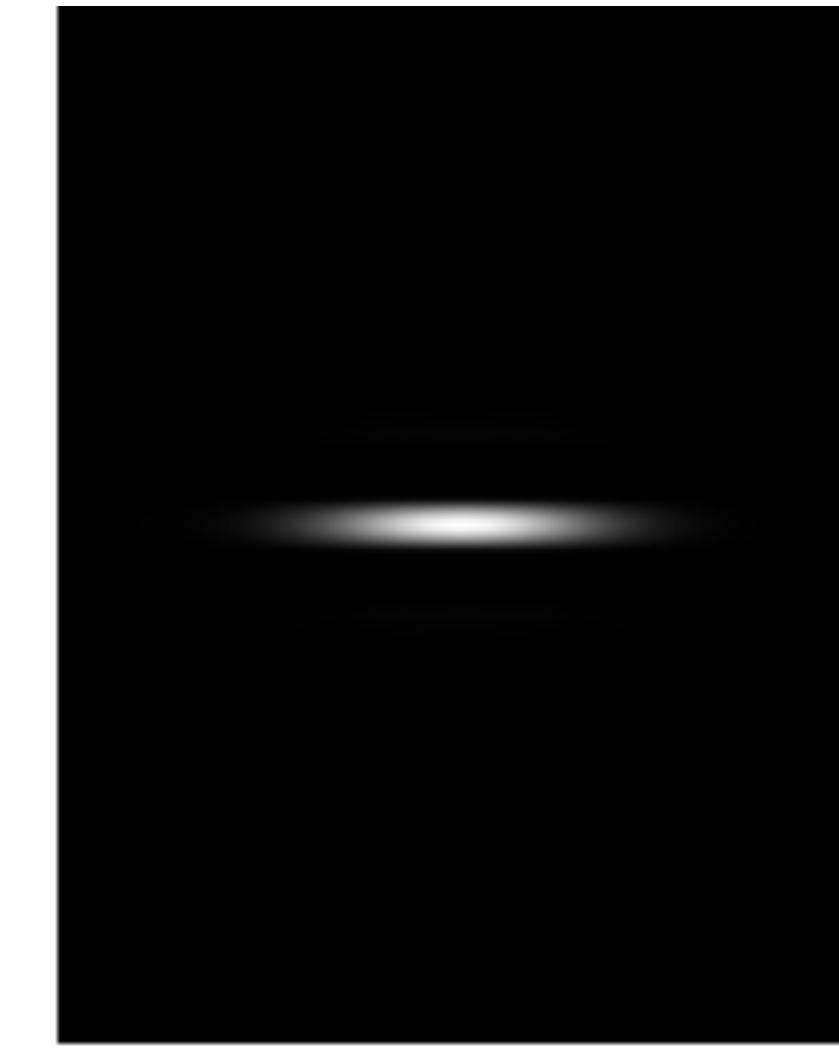
Controls the orientation of the stripes, zero means vertical position



Theta (Θ) = 0



Theta (Θ) = 45

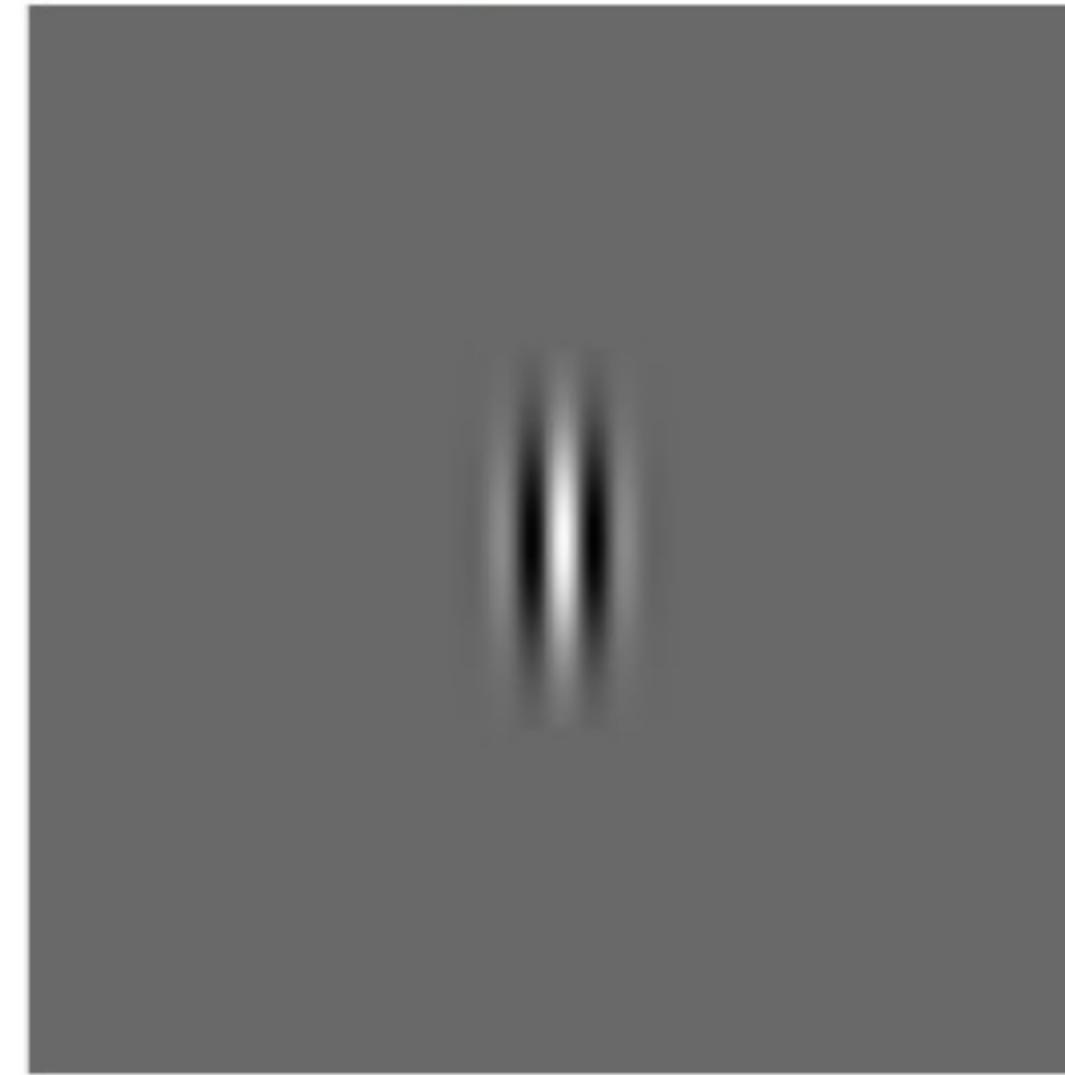


Theta (Θ) = 90

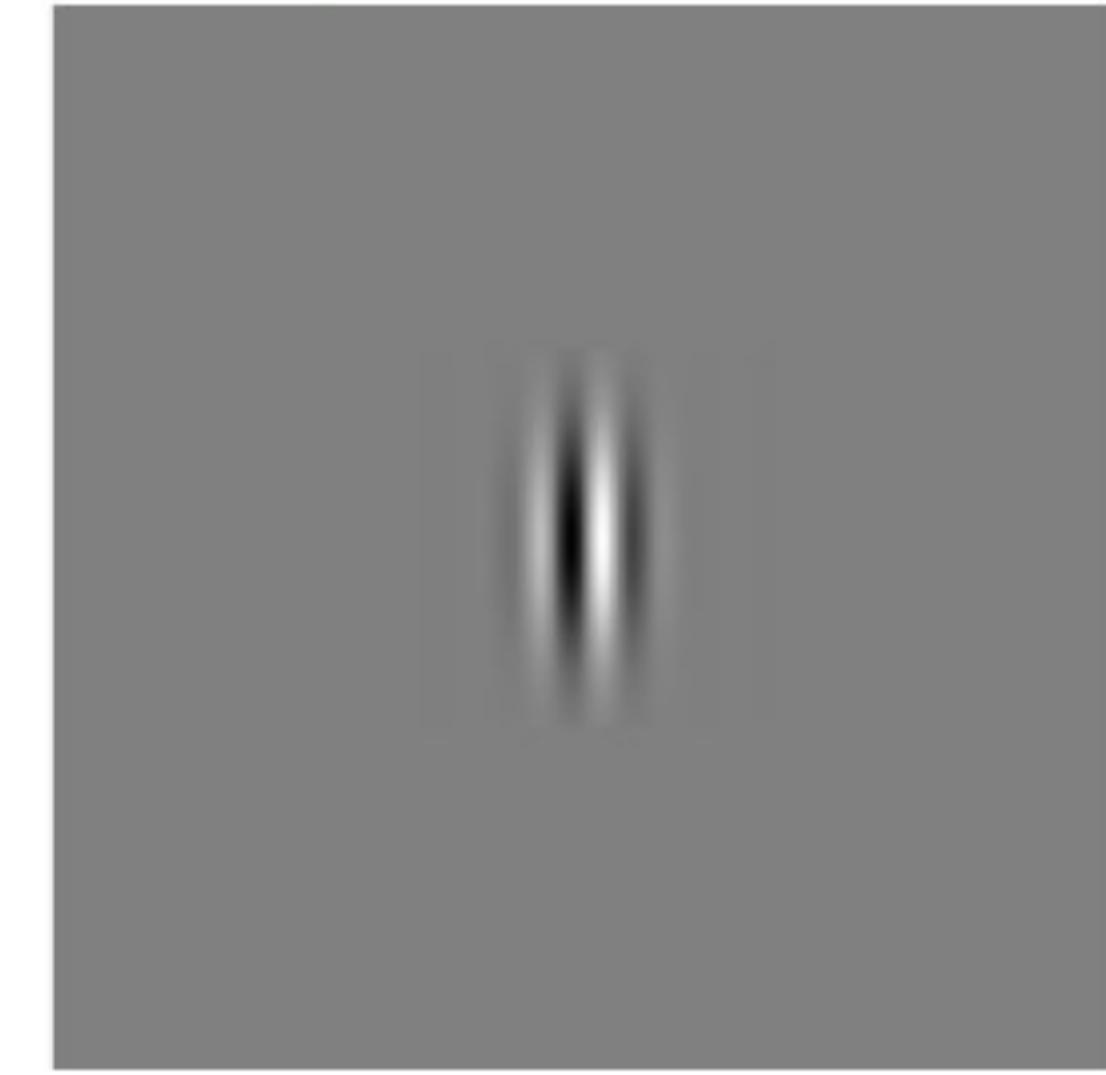
Quelle: https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97

ψ : Phase offset of the sinusoidal function

Controls offset of the phase and thus the symmetry of the filter



Phase offset = 0
(symmetric function)

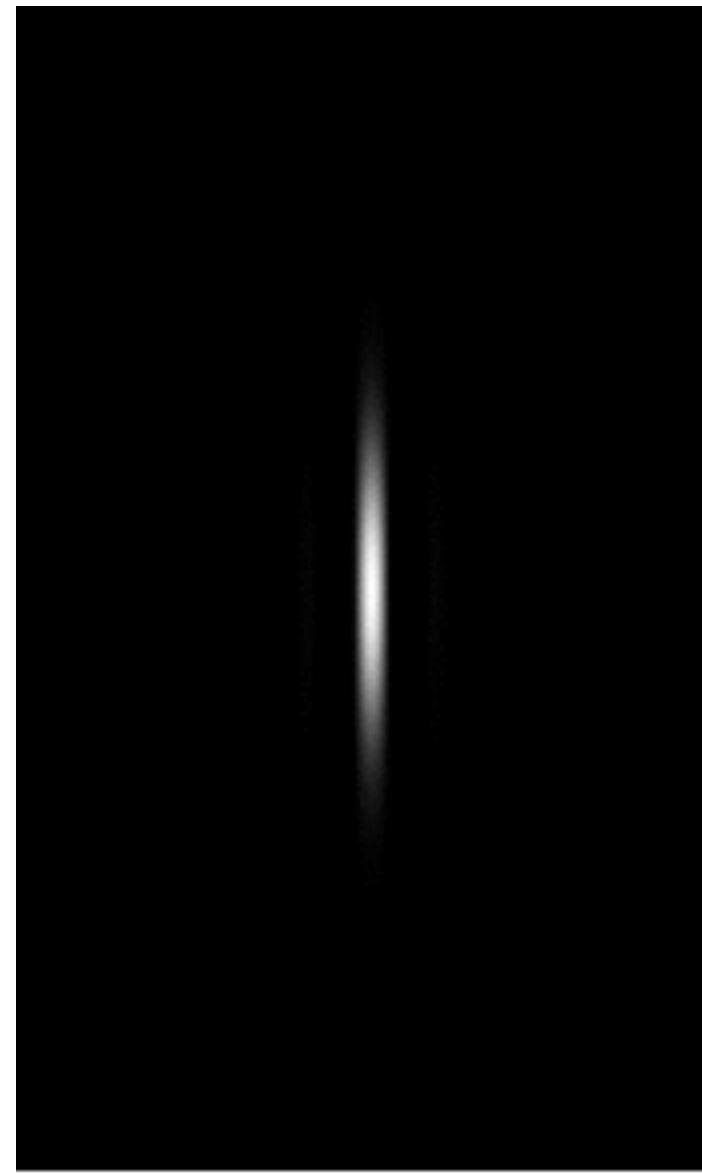


Phase offset = -90
(anti-symmetric function)

Quelle: <https://slidetodoc.com/2-d-gabor-functions-for-image-processing-and/>

σ : Standard deviation of the Gaussian

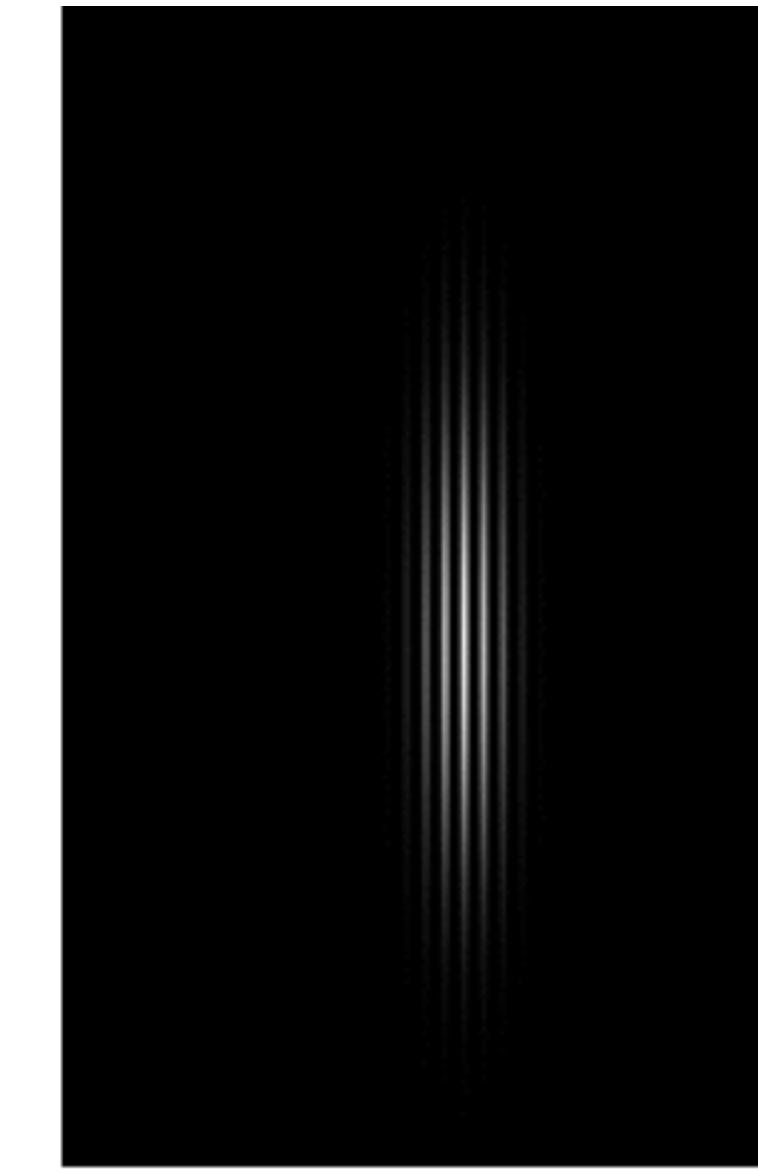
As the standard deviation increases, it gets more broad and thus keeps more of the sinusoidal peaks activated, which leads to more stripes



Sigma (σ) = 10



Sigma (σ) = 30

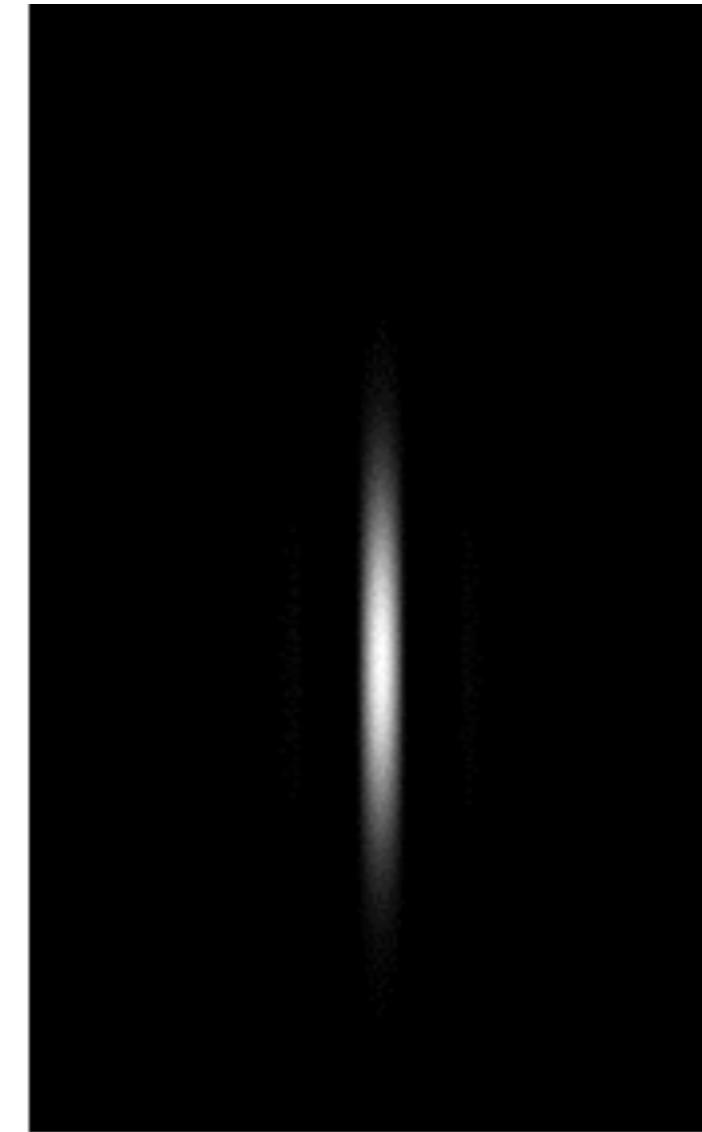


Sigma (σ) = 45

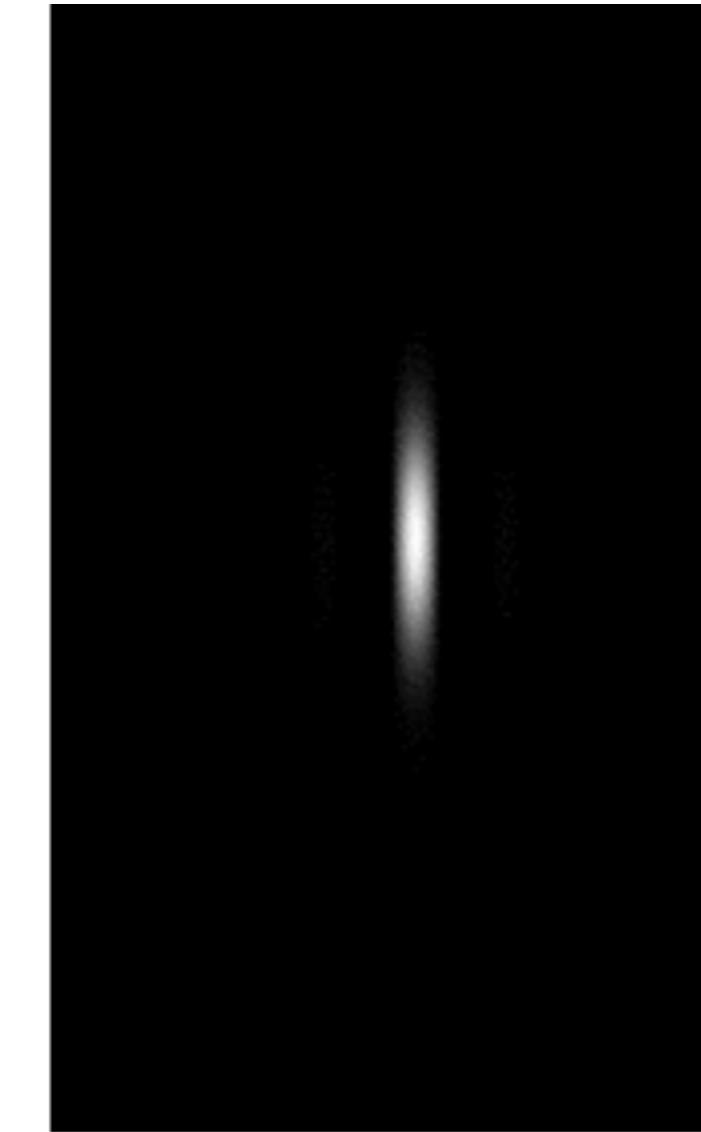
Quelle: https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97

γ : Spatial aspect ratio (of the x- and y-axis of the Gaussian ellipse)

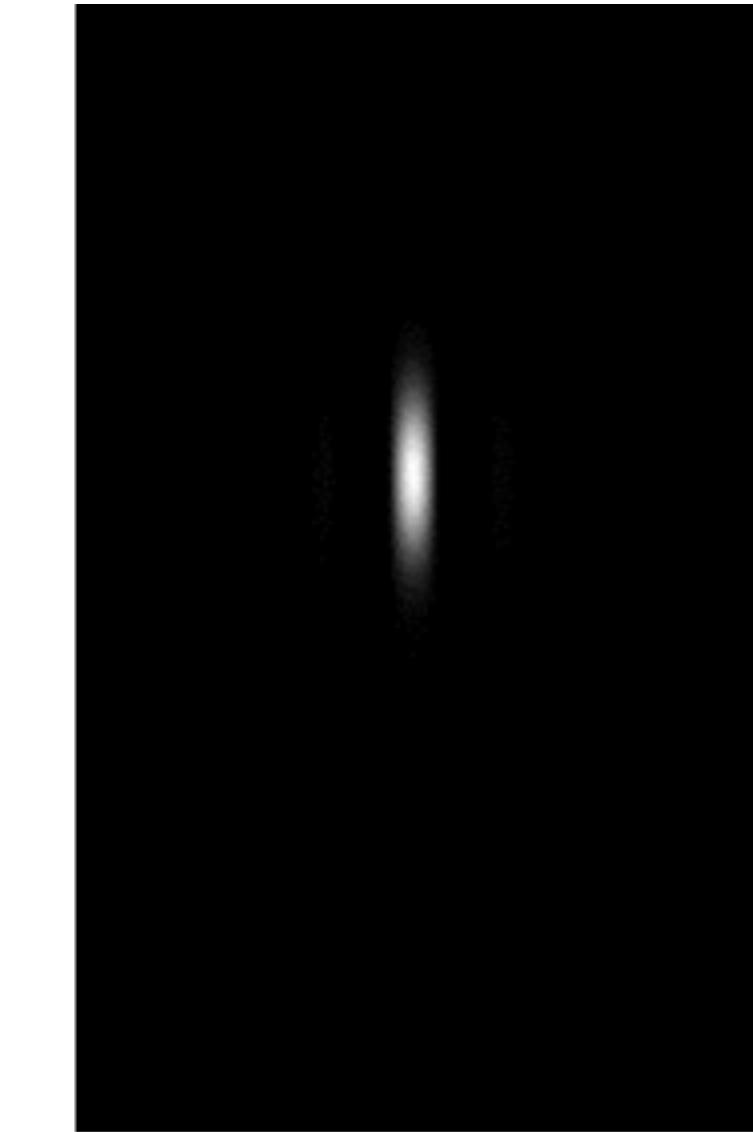
Controls the height of the Gabor filter where increasing γ , leads to a smaller height and vice versa.



Gamma (γ) = 0.25



Gamma (γ) = 0.5



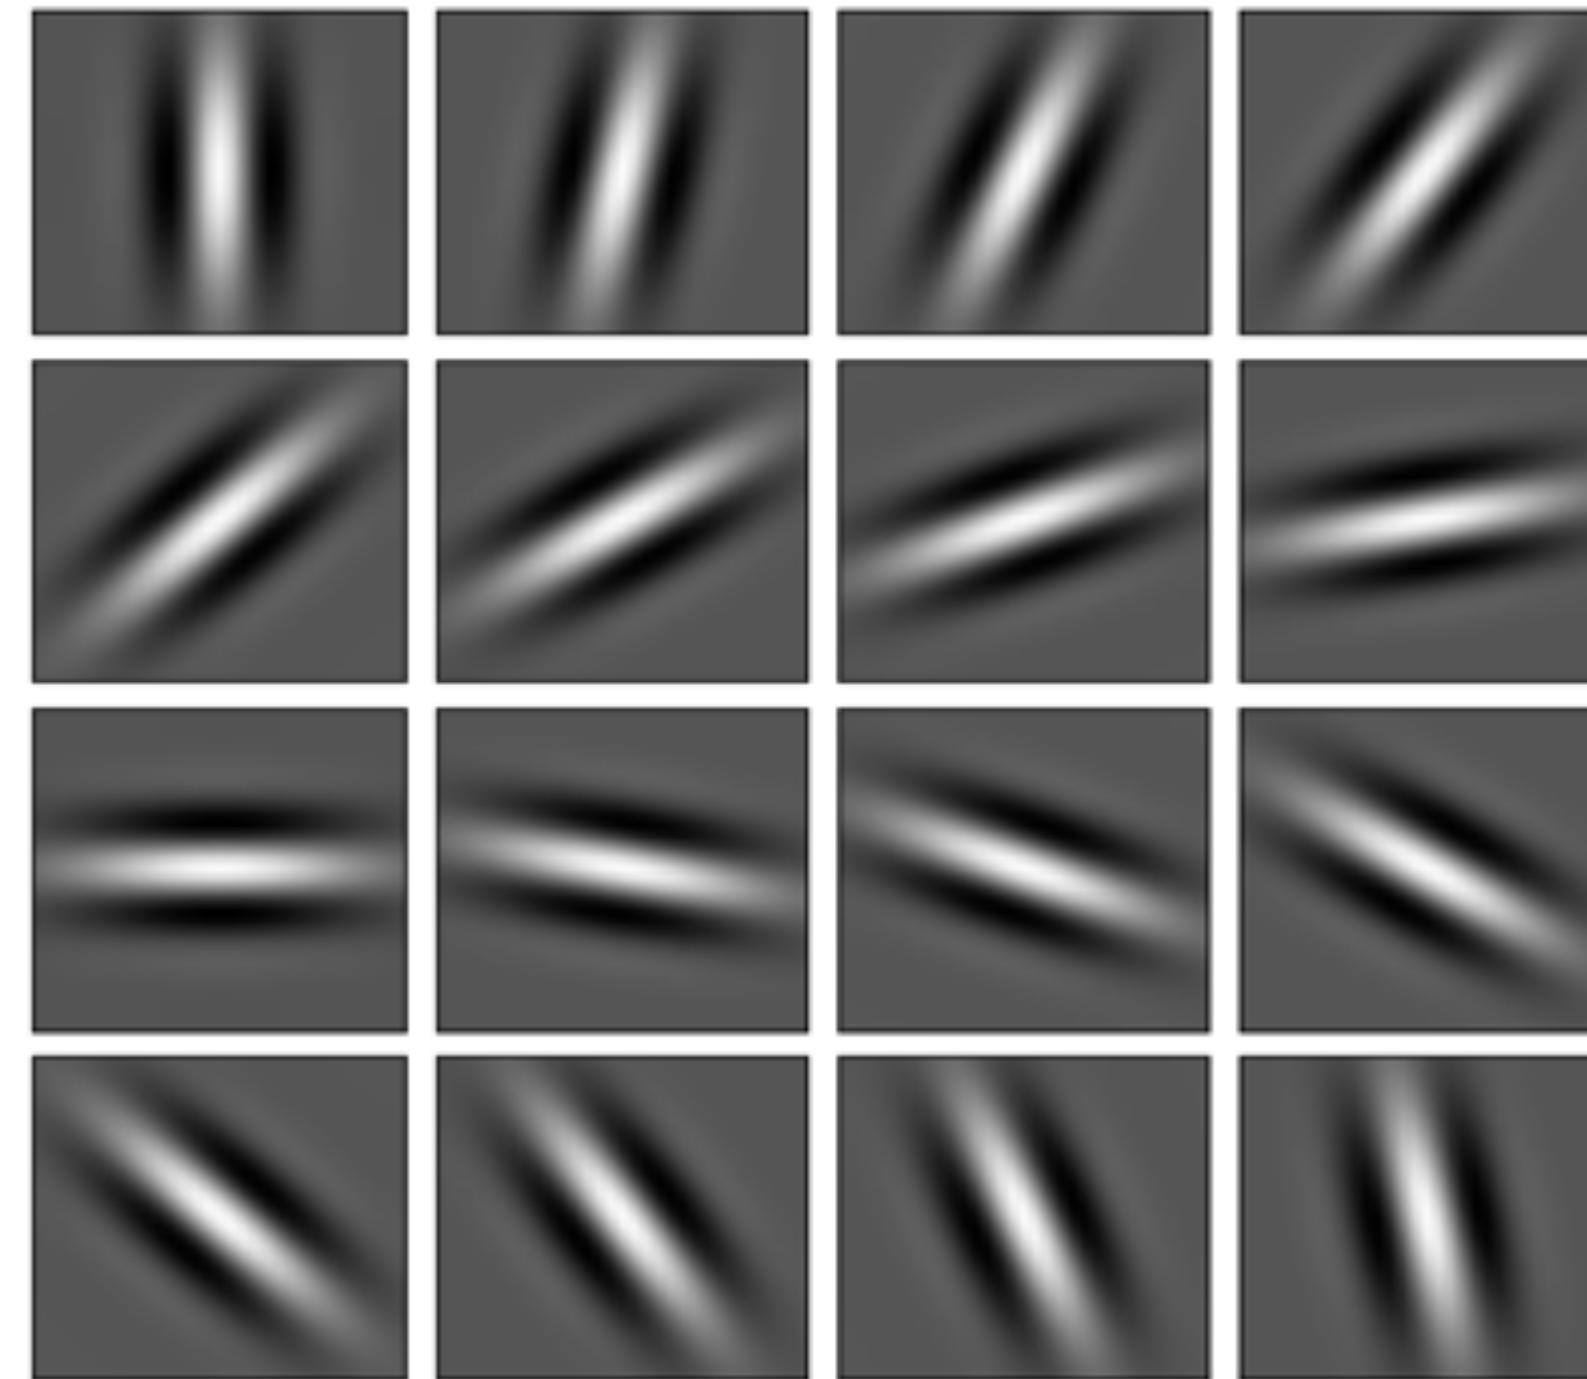
Gamma (γ) = 0.75

Quelle: https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97

Texture Kernels

Gabor Filter

- As differently parameterized Gabor filters detect different patterns in different orientation, usually a whole bank of filters is applied to the image



A bank of 16
Gabor filters
with varying
orientation

Quelle: https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97

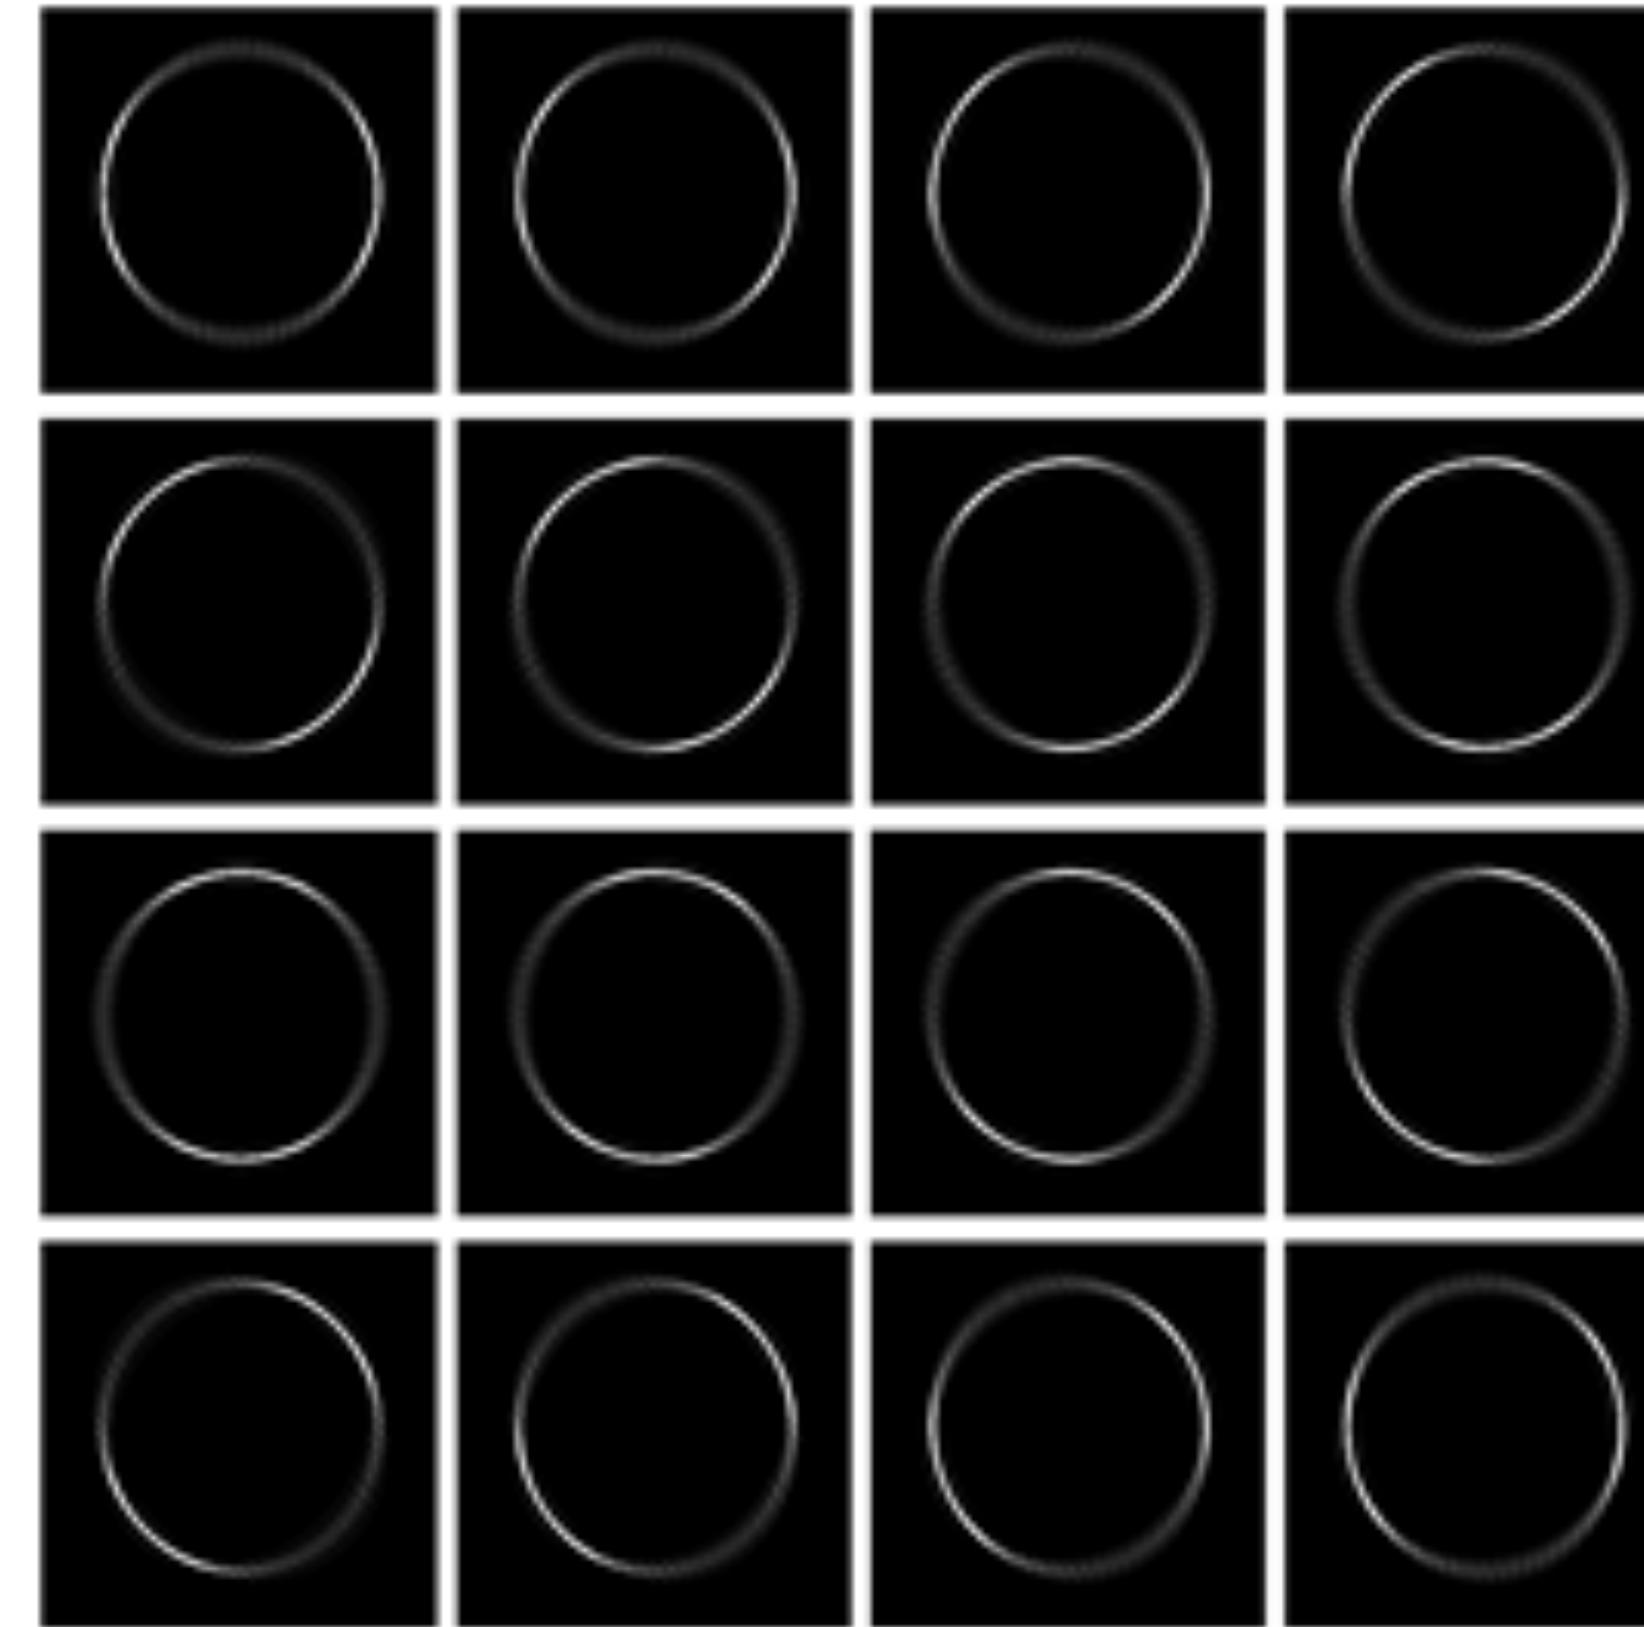
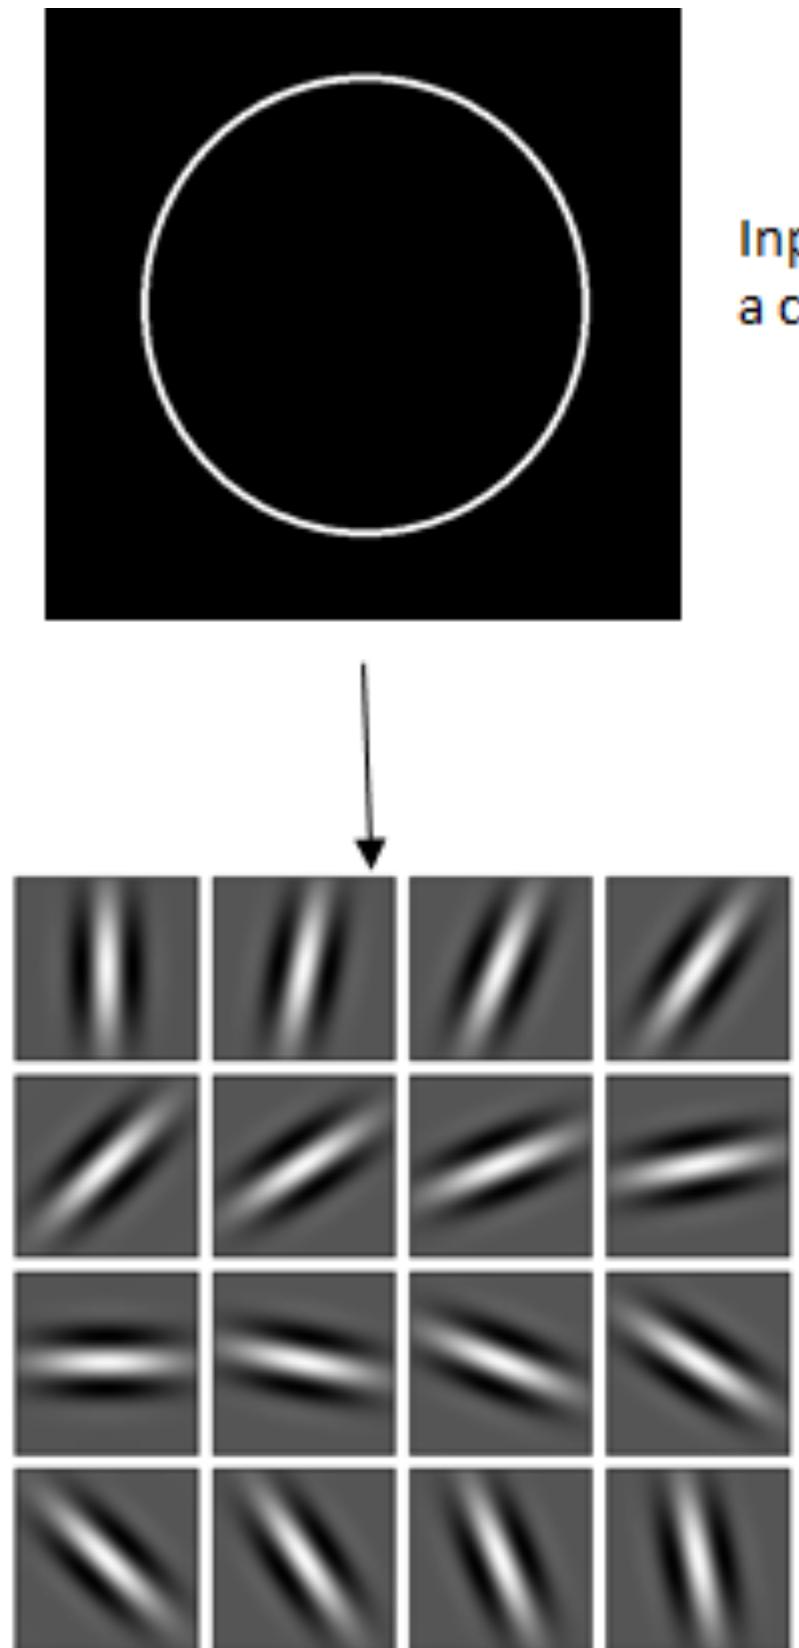
Let's apply that now!



Texture Kernels

Gabor Filter

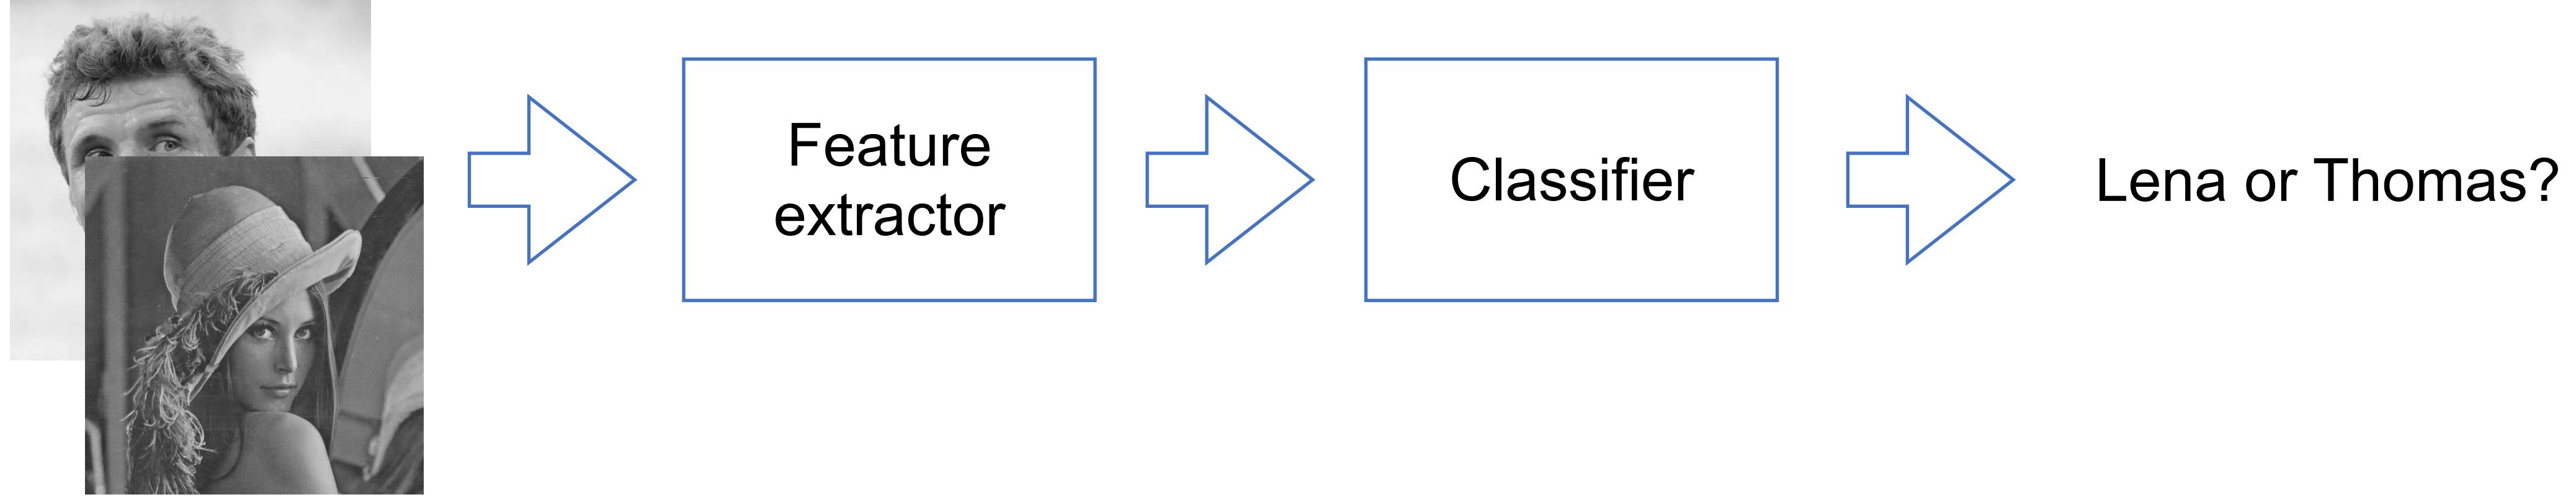
Application to a circle example



The output circle as seen when pass through individual Gabor filter

Quelle: https://medium.com/@anuj_shah/through-the-eyes-of-gabor-filter-17d1fdb3ac97

We ultimately want to recognize something from this image.

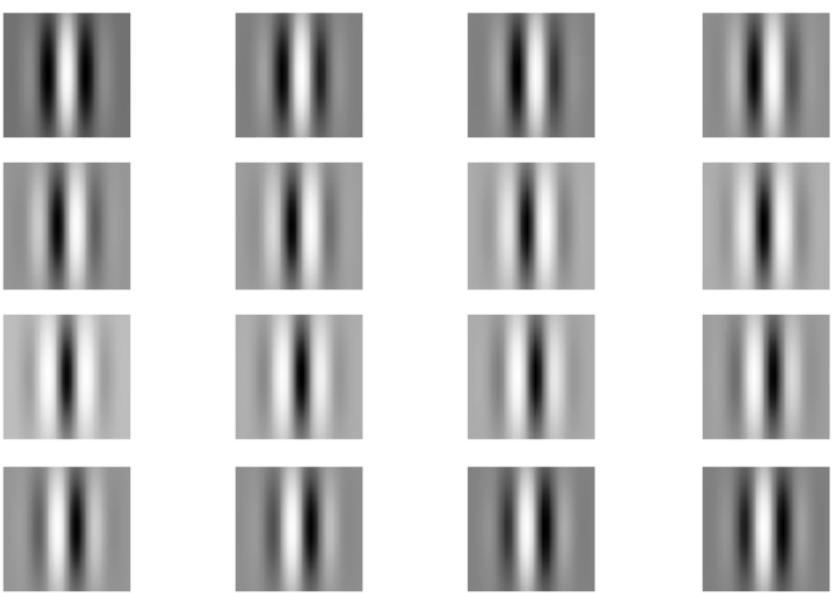


- We thus need meaningful features that explain the variance of the classes „Thomas“ and „Lena“ and allow a decision using some classifier.
- It's often being perceived that the choice of classifier is the most important in this problem.
- But in reality, the selection of proper features takes up 80% of the time.

Let's extract some features from Thomas



Gabor filters (kernel size: 31)



$$\psi \in \left[0, \frac{\pi}{8}, \dots, \frac{15\pi}{8}\right]$$

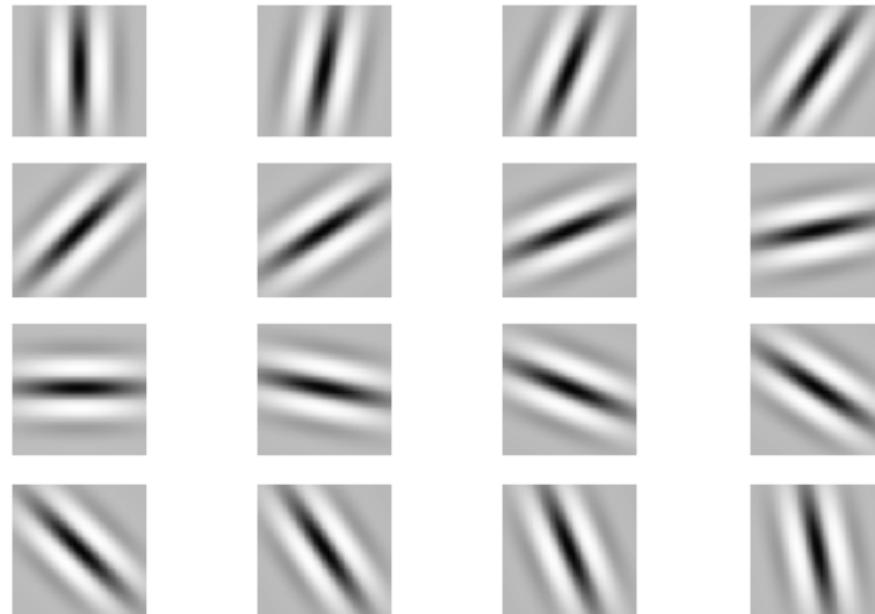
$$\theta = 0, \gamma = 0.5, \lambda = 10, \sigma = 5$$



Let's extract some features from Thomas



Gabor filters (kernel size: 31)



Eyes / lips highlighted:
Might be a
good feature

$$\theta \in \left[0, \frac{\pi}{8}, \dots, \frac{15\pi}{8}\right]$$

$$\gamma = 0.5, \lambda = 10, \psi = \pi, \sigma = 5$$



- Finding the right bank of Gabor filters is a hard and time consuming task
 - Has to be fine tuned for different kind of images and problems
 - You will get a chance to make this experience in the practical course



Wouldn't it be nice to have some tool which is creating them for us automatically?